

ID1217 Concurrent Programming
Lecture 2



ROYAL INSTITUTE
OF TECHNOLOGY

Shared Memory Programming:
Processes and Synchronization

Vladimir Vlassov
KTH/EECS



ROYAL INSTITUTE
OF TECHNOLOGY

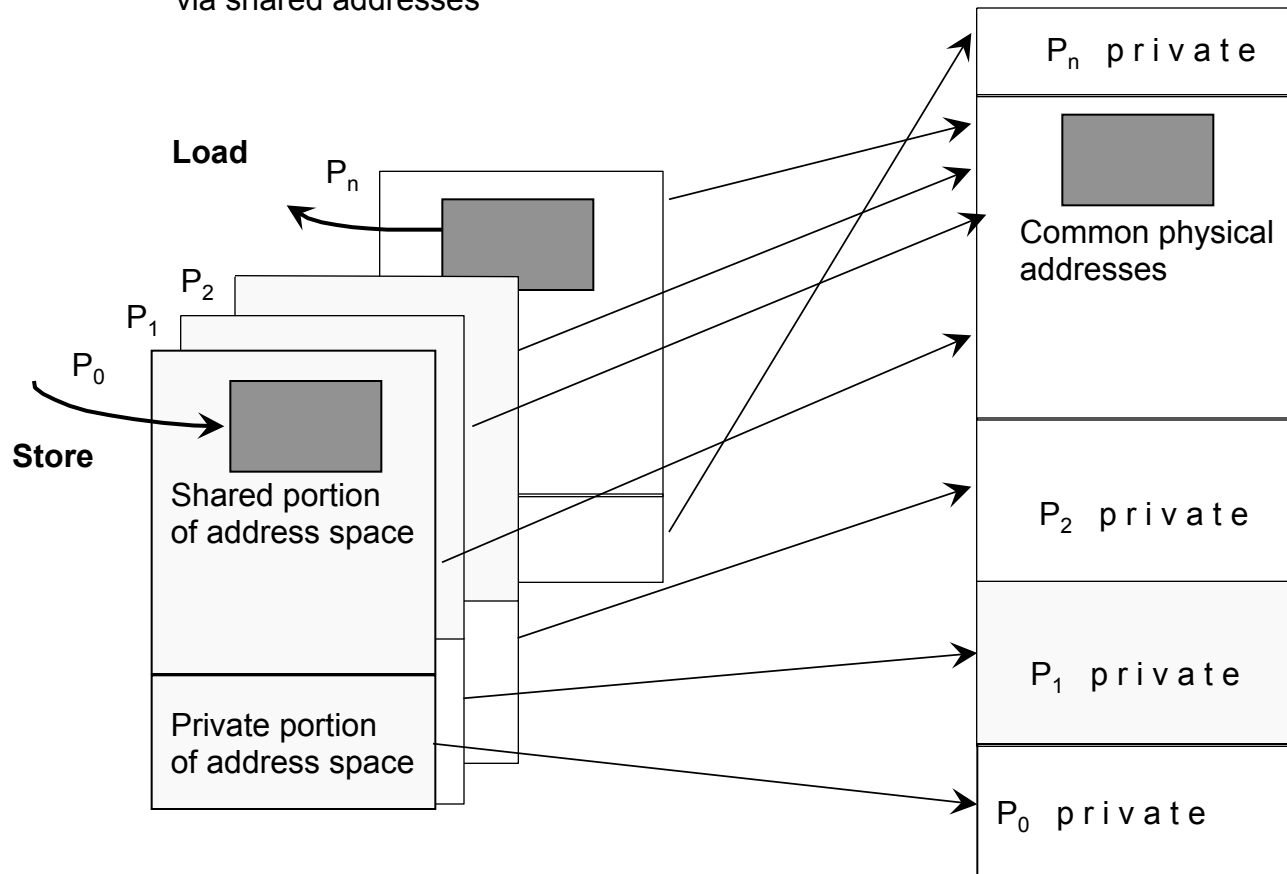
Outline

- Processes
 - Atomic actions
 - Process histories.
 - Interleaving semantics and nondeterminism of concurrent execution.
- Synchronization
 - Atomicity.
 - The await statement.
 - At-Most-Once property.
- Examples of parallelization and synchronization
 - Finding patterns in a file.
 - Find the maximum element in an array.

Shared Address Space Model

Virtual address spaces for a collection of processes communicating via shared addresses

Machine physical address space



Process State. Actions. Process History

- *Process (thread)* is an abstract entity that performs tasks assigned to processes
- *Process state* is formed of values of variables at a point in time.
- Each process executes a sequence of statements.
- Each statement consists of one or more *atomic (indivisible) actions* which transform one state into another
 - Some actions allow processes to communicate with each other
- Sequence of states makes up a process history
- The *process history* is a trace of ONE execution.

$$P: Q_0 \xrightarrow{a_1} Q_1 \xrightarrow{a_2} \dots \xrightarrow{a_m} Q_m$$

Atomic Actions

- **Atomic action** – indivisible sequence of state transitions made atomically
- Fine-grained atomic actions
 - Machine instructions (read, write, swap, etc.) – atomicity is guaranteed by HW
- Coarse-grained atomic actions
 - A sequence of fine-grained atomic actions executed indivisibly (atomically)
 - Internal state transitions are not visible “outside”
 - For example, a critical section of a code
- Notation (to be used further): **< statements >** – a list of statements to be executed atomically.

Interleaving Semantics of Concurrent Execution

- The concurrent execution of several processes can be viewed as **the interleaving of histories of the processes**
 - i.e., the interleaving of sequences of atomic actions of different processes.
- Individual histories:
 - Process 1: $s_0 \rightarrow s_1 \rightarrow s_2 \rightarrow \dots \rightarrow s_n$
 - Process 2: $p_0 \rightarrow p_1 \rightarrow p_2 \rightarrow \dots \rightarrow p_n$
 - Process 3: $q_0 \rightarrow q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_n$
- Interleaved concurrent histories:
 - Trace 1: $s_0 \rightarrow p_0 \rightarrow s_1 \rightarrow p_1 \rightarrow p_2 \rightarrow q_0 \rightarrow s_2 \rightarrow q_1 \rightarrow \dots$
 - Trace 2: $p_0 \rightarrow s_0 \rightarrow q_0 \rightarrow q_1 \rightarrow s_1 \rightarrow p_1 \rightarrow p_2 \rightarrow q_2 \rightarrow \dots$

Nondeterminism of Concurrent Execution

- **Nondeterminism**: The behavior of a concurrent program is not reproducible because different interleavings (histories) can be observed on different concurrent executions.
- Reason: **Asynchrony**
 - Each process executes at its own rate (caches misses, page faults, etc.)
 - **Concurrent means ambiguous temporal order**
- A concurrent program of n processes each of m atomic actions can produce $(n \times m)! / (m!)^n$ different histories.
 - For example, $n = 3$, $m = 2$ give 90 different histories.
- Impossible to show the correctness of a program by testing (“run the program and see what happens”).

Example of Nondeterminism

- A program:

```
int y = 0; z = 0;  
co x = y + z; || y = 1; z = 2; oc
```

- Trace 1:

```
x = y{0} + z{0}; y = 1; z = 2;      {x == 0}
```

- Trace 2:

```
y = 1; x = y{1} + z{0}; z = 2;      {x == 1}
```

- Trace 3:

```
y := 1; z := 2; x := y{1} + z{2};    {x == 3}
```

- Trace 4:

```
load y{0} to R1; y := 1; z := 2;  
add z{2} to R1{0}; store R1 to x; {x == 2}
```


Synchronization

- **Synchronization** is a mechanism to delay a process until it may proceed
 - Allows reducing the entire set of possible histories to those which are desirable (correct).
 - To preserve (true) dependences between processes
 - To avoid race conditions if any
 - To coordinate access to shared resources
 - To hide state transitions
- Two kinds of synchronization:
 - **Mutual exclusion**
 - guarantees that only one process at a time accesses a shared resource, i.e. only one process at a time executes its critical section.
 - **Condition synchronization**
 - delays a process until a certain condition is true.

Specifying Synchronization: The Await Statement

< **await**(**B**) **S**; >

- Wait for **B** to be true, then execute **S** atomically
- Combines **condition synchronization with mutual exclusion**
 - Executed as an atomic action
 - **B** is guaranteed to hold when **S** begins
 - **S** is guaranteed to terminate

< **await**(**B**); >

- Wait for a condition **B** to be true.
- Used to express **condition synchronization**

< **S**; >

- Execute a list of statements **S** atomically (indivisibly).
- Can be used to express **mutual exclusion**



ROYAL INSTITUTE
OF TECHNOLOGY

Assumptions on Machine Architecture

- RISC architecture
 - Values are manipulated by loading them into registers, operating on them, and storing them back to memory
- **Atomic memory operations (atomic registers)**
 - Values of the basic types are stored in memory elements that are read and written as atomic actions
- Process context
 - Each process has its own set of registers and stack
- All temporaries are stored in private memory
 - Any intermediate results are stored in registers or in process' private memory (on stack)

Registers (Memory Locations)***

- **Safe register**
 - If read does not overlap write, read returns the value written by the most recent write
 - If read overlaps write, it returns **any** value (from the range of valid values)
- **Regular register**
 - If read does not overlap write, the register is safe
 - If read overlaps write, it returns **either the old or the new value**
- **Atomic register (we assume)**
 - If read does not overlap write, the register is safe, i.e. read returns the value written by the most recent write
 - If read overlaps with write, it returns **either the old value or the new value but not newer than the next read**

The At-Most-Once Property

- Unnecessary explicit coarse-grained atomicity causes unnecessary overhead and may reduce parallelism.
- **Critical reference** is a reference to a variable which is (can be) changed by another process.
- A statement $\mathbf{x} = \mathbf{e}$ appears to be executed atomically when it satisfies the **At-Most-Once (AMO) property**, i.e. if
 - either (a) \mathbf{e} contains at most one critical reference and \mathbf{x} is not referenced by another process;
 - or (b) \mathbf{e} contains no critical references and \mathbf{x} may be read by other processes.



ROYAL INSTITUTE
OF TECHNOLOGY

Implication of the AMO Property

$\langle S; \rangle \equiv S;$

- as long as **S** contains at most one critical reference seeing by other processes,
- In this case, one cannot tell the difference, so **S** will appear to execute atomically w.r.t. to other processes.

$\langle \text{await}(B); \rangle \equiv \text{while (not } B) \text{ continue;}$

- if **B** contains (reads) at most one critical reference

Parallelizing a Sequential Program

- How? Identify independent and dependent tasks, assign tasks to processes, identify shared variables, synchronize processes.
- Data accessed by a process:
 - Read set (rs) – variables that are only read
 - Write set (ws) – variables that are written (and possibly read)
- **Independence:** processes are independent iff the write set of one proc is disjoint from the read and write sets of another proc:

$$ws_1 \cap (rs_2 \cup ws_2) = \emptyset \text{ AND } ws_2 \cap (rs_1 \cup ws_1) = \emptyset$$

Dependencies Between Atomic Actions w.r.t. to a Shared Variable

- Assume two atomic actions access the same variable
- Dependencies:
 - **True dependence**: read after write, e.g. $x := 1; a := x;$
 - **Anti dependence**: write after read, e.g. $a := x; x := 1;$
 - **Output dependence**: write after write, e.g. $x := 1; x := 2;$
 - **Input dependence**: read after read, e.g. $a := x; b := x;$
- If the actions are executed by different processes, all dependencies must be preserved by synchronization
 - Input dependences do not require synchronization
 - Some anti and output dependencies can be **false dependencies** caused by reuse of memory – To avoid, rename (disjoin) variables
 - An output dependence does not need to be preserved (ordered) if the order of writes is not important, e.g. OR-parallelism

Example of Parallelization and Synchronization: Finding Patterns in a File

- Find all instances of a pattern in a file: **grep pattern filename**
- Sequential program:

```
string line;  
open the file; read a line of input from the file into line;  
while (!EOF) { // EOF is end of file  
    look for pattern in line;  
    if (pattern is in line) print line;  
    read next line from the file into line;  
}
```

- Parallelization:
 - Tasks: read a line, compare with the pattern, print if match
 - Apply Producer-Consumer paradigm:
 - Producer reads a line from the file and stores it to a shared buffer
 - Consumer gets a line from the buffer, tests it against pattern, prints if match.

Example (cont'd): How to Synchronize

- True/anti dependence between P and C via the shared buffer.
- Assume a single-slot buffer
 - the buffer provides place for only one string
- Synchronization:
 - **Mutual exclusion**: The buffer is accessed by one process at a time
 - **Condition synchronization**
 - To put a line, Producer waits for the buffer to be empty
 - To get a line, Consumer waits for a buffer to become full
- Let's use two counters:
 - **p** – number of produced lines (stored to the buffer by Producer)
 - **c** – number of consumed (fetched from the buffer by Consumer)
- Synchronization requirement expressed as a predicate:

$$PC: c \leq p \leq c + 1$$

Example (cont'd): Parallel Program

```
// shared variables:
string buf;    // buffer
bool done = false; // termination
int p = 0, c = 0; // counters
Process Producer { // reads lines
    string line;
    open the file;
    while (true) {
        read a line from the file into line;
        if (EOF) {
            done = true;
            break;
        }
        // wait for buffer to be empty
        < await (p == c); >
        buf = line;
        p++; // signal that buffer is full
    }
};
```

```
Process Consumer { // finds patterns
    string line;
    while (true) {
        // wait for buffer to be full
        // or done to be true
        < await ((p > c) || done); >
        if (p > c) {
            line = buf;
            c++; // signal that buf is empty
            look for pattern in line;
            if (pattern is in line) write line;
        }
        if (done) break;
    }
}
```

Example of Synchronization:

Find the Maximum

- Given array $a[1:n]$ of positive integers. Find the maximum value m
 $(\forall j : 1 \leq j \leq n : m \geq a[j]) \wedge (\exists j : 1 \leq j \leq n : m = a[j])$
 - When program terminates, m is at least as large as every element of a
 - m is one of elements of a
- Sequential program:

```
int m = 0;
for [i = 0 to n-1]
    if (a[i] > m) m = a[i];
```

- How to parallelize? – Examining all elements in parallel
- Parallel program without synchronization:

```
int m = 0;
co [i = 0 to n-1]
    if (a[i] > m) m = a[i];
oc
```

- Program is incorrect because of the race condition
- Synchronization is needed

Example (cont'd): Concurrent Program

- First attempt to synchronize: execute cond. updates atomically

```
int m = 0;
co [i = 0 to n-1]
    < if (a[i] > m) m = a[i]; > oc
```

- Not efficient (over-constrained): all the updates are serialized just like in the sequential program but executed in an arbitrary order.
- Observations:
 - Read and test can be executed in parallel for each **i**
 - Updates of **m** require atomicity (mutual exclusion) for serialization
 - Can use second test in critical section to avoid races

Example (cont' d): Solution for Concurrent Program

- Read and test in parallel without mutual exclusion. Those who passed the test, execute conditional update atomically (with mutual exclusion)

```
int m = 0;
co [i = 0 to n-1]
    if (a[i] > m)
        < if (a[i] > m) m = a[i]; >
oc
```

Lessons Learned

- Synchronization is required whenever processes read and write shared variables (to preserve true, anti, and output dependences)
- Atomicity helps to provide mutual exclusion
- Test followed by an atomic test-and-update is a useful combination for conditional updates
 - Helps to avoid races among concurrent conditional updates that depend on the same condition

Exercise 1: Non-determinism.

Possible final values

- Consider the following fragment of a program outline:

```
var x = 1, y = 2, z = 3;  
co x = x + 1;  
  || y = y + 2;  
  || z = x + y;  
  || <await (x > 1) x = 0; y = 0; z = 0 >  
oc
```

- Assume that atomic actions in the first three arms of the **co** statement are reading and writing individual variables.
- Does the co-statement terminate? What are the possible final values of x, y and z? Explain.

Solution to Exercise 1

```
var x = 1, y = 2, z = 3;  
co x = x + 1  
|| y = y + 2  
|| z = x + y  
|| <await (x > 1) x = 0; y = 0; z = 0 >  
oc
```

- The possible final values of x, y and z are

$x == \{0\}$

$y == \{0, 2, 4\}$

$z == \{0, 1, 2, 3, 4, 5, 6\}$

Exercise 2: Await versus If

- Consider the following program fragment that spawns three threads in the co-statement:

```
int x = 0;
co <await (x != 0) x = x - 2;>
|| <await (x != 0) x = x - 3;>
|| <await (x == 0) x = x + 5;>
oc
```

- Does the program terminate? If so, what are the possible final values of **x**? If not, why not?
- Suppose the **await** statements are changed to **if** statements – namely, that **await** is replaced by **if** and the angle brackets are deleted. Now the program is sure to terminate. What are the possible final values of **x**?

Solution to Exercise 2

(a) Yes, it does terminate. The final value of x is zero. The last statement executes first, then the other two in either order.

(b)

```
int x = 0;
co if (x != 0) x = x - 2; // S1
|| if (x != 0) x = x - 3; // S2
|| if (x == 0) x = x + 5; // S3
oc
```

Possible histories and corresponding final values are the following:

S1 and S2 tests fail then S3	$x == 5$
S2 test fails; S3; S1	$x == 3$
S1 test fails; S3; S2	$x == 2$
S3; S1 S2	$x == \{0, 2, 3\}$

Exercise 3: Intersection of two arrays

- Given integer arrays **a[1:m]** and **b[1:n]**, assume that each array is sorted in ascending order.
 - (a) Develop an outline of a sequential program to compute the number of values that appear in both **a** and **b**, i.e. the intersection of **a** and **b**.
 - (b) Develop an outline of a parallel program based on your sequential program.

Use the **co** statement for concurrent execution and the **await** statement for synchronization (if needed).

Possible Solution to Exercise 3 (a)

(a) We can do a linear search in the arrays to compare the elements and find the common ones.

- The following sequential solution is inefficient:

```
int i = 0, j = 0, count = 0;
for (i = 0; i < m; i++)
    for (j = 0; j < n; j++)
        if (a[i] == b[j]) count++;
```

- The outline of a more efficient sequential program:

```
int i = 0, j = 0, count = 0;
while (i < m && j < n) {
    if (a[i] < b[j]) i++;
    else if (a[i] > b[j]) j++;
    else { count++; i++; j++; }
}
```

Possible Solution to Exercise 3 (b)

(a) The outline of a parallel program:

```
int i = 0, j = 0, count = 0, A, B;
while (i < m && j < n) {
    A = b[i]; B = b[j];
    co if (A < B) i++;
    || if (A > B) j++;
    || if (A == B) { count++; i++; j++; }
    oc
}
```

- There is no need of any synchronization because the conditions are disjoint. Only one of them in every iteration will increment the shared variable(s).

Properties of a Program

- A property of a program is an attribute of ALL histories of a program, e.g., correctness.
- Two kinds of properties:
 - **Safety properties state that nothing bad ever happens**
 - A safety property is one in which the program never enters a “bad” state
 - **Liveness properties state that something good eventually happens**
 - A liveness property is one in which the program eventually enters a “good” (desirable) state



ROYAL INSTITUTE
OF TECHNOLOGY

Properties of a Program (cont'd)

- **A property can be represented as a conjunction of a safety property and a liveness property**
- For example:
 - Property: A message sent is always delivered (exactly once)
 - Safety: A message sent is delivered at most once
 - Liveness: A message sent is delivered at least once



ROYAL INSTITUTE
OF TECHNOLOGY

Examples of Safety and Liveness Properties

- **Safety properties:**
 - Partial correctness
 - The final state is correct, assuming that the program terminates
 - Mutual exclusion of critical sections
 - “Bad” when more than one proc. are in critical sessions
 - Absence of deadlocks (deadlock-freedom)
 - A set of processes is deadlocked when each process in the set is waiting for an event which can only be caused by another process in that set.
- **Liveness properties:**
 - Termination
 - Every history (trace) is finite.
 - Eventually enter a critical section (if any)
- **Total correctness** combines termination and partial correctness
 - A program always terminates with a correct answer

Proving Properties

Three main approaches

1. Testing or debugging

- Run and see what happens
- Limited to considered cases

2. Operational reasoning

- “Exhaustive case analysis”
- Considers enormous number of histories: $n \cdot m! / (m!)^n$
- Helps in development

3. **Assertional reasoning**

- Based on axiomatic semantics: axioms, inference rules, assertions
- Work is proportional to the number of atomic actions