

ID1217 Concurrent Programming
Lecture 6



Tutorial: Introduction to Pthreads

Vladimir Vlassov
KTH/ICT/EECS



ROYAL INSTITUTE
OF TECHNOLOGY

Additional reading

- Tutorial: POSIX Threads Programming at <https://computing.llnl.gov/tutorials/pthreads/>
- Ch 4 in An Introduction to Parallel Programming, by Peter Pacheco, Morgan Kaufmann (available online via KTH library)



ROYAL INSTITUTE
OF TECHNOLOGY

Outline

- Pthreads – IEEE POSIX threads API
 - Pthreads types and functions
 - Thread creation and termination
 - Joinable and detached threads
 - Synchronization primitives: Mutex locks, Condition variables, Semaphores
 - Examples



ROYAL INSTITUTE
OF TECHNOLOGY

Pthreads: POSIX Threads

- **Pthreads** is a standard set of C library routines for multithreaded programming with shared variables
 - IEEE Portable Operating System Interface, POSIX, section 1003.1 standard, 1995
- Allows to create and synchronize multiple threads in a heavy-weight process
 - Threads share a common address space (thru common variables)
 - Each thread has a private stack for local variables
- Goal in developing the Pthreads API:
 - “To give programmers the ability to write concurrent applications that run on both uniprocessor and multiprocessor machines transparently, taking advantage of the additional processors if any.”



ROYAL INSTITUTE
OF TECHNOLOGY

Pthreads API

- The Pthreads API includes a library of functions that contains in total
 - over 60 functions excluding extensions;
 - about 100 subroutines including extensions.
- Thread management (29 functions): create, exit, detach, join, get/set attributes, ...
- Thread cancellation (9): cancel, test for cancellation, ...
- Mutex locks (19): init, destroy, lock, unlock, try lock, get/set attributes;
- Condition variables (11): init, destroy, wait, timed wait, signal, broadcast, get/set attributes;
- Read/Write locks (13): init, destroy, write lock/unlock, read lock/unlock, get/set attributes;
- Thread specific storage (4)
- Signals (3): send a signal to thread, signal mask
- Unsupported extension (16): get/set scheduling policy, ...
- Extension: semaphores (semaphore.h)



Header Files, Compiling and Linking

- Header files:
`#include <pthread.h>`
`#include <sched.h>`
`#include <semaphore.h>` to use semaphores
- Compile using **gcc** (or **g++**) and link to **pthreadlib**
 - For example:
`g++ -c -O3 common.cpp`
`g++ -O3 -o myprog myprog.cpp common.o -lm -lpthread`
 - To link with POSIX 4 extension (e.g. semaphores): **-lposix4**
- Run as an ordinary executable code:
`.\myprog <command line parameters>`








ROYAL INSTITUTE
OF TECHNOLOGY

Pthreads Naming Convention

- Types: **pthread[_object][_np]_t**
- Functions: **pthread[_object]_action[_np]**
- Constants and Macros: **PTHREAD_PURPOSE[_NP]**
- If the type of an object is not a thread, then:
 - *object* represents the type of object, e.g., **mutex**
 - **action** is an operation to be performed on the object, e.g., **lock**
 - **np** or **NP** indicates that the name or symbol is a non-portable extension to the API set,
 - *PURPOSE* indicates the use or purpose of the symbol, e.g. **PTHREAD_CREATE_DETACHED**

Primitive Data Types



type	Description
pthread_attr_t	Thread creation attribute
pthread_cleanup_entry_np_t	Cancellation cleanup handler entry
pthread_condattr_t	Condition variable creation attribute
pthread_cond_t	Condition Variable synchronization primitive
pthread_joinoption_np_t	Options structure for extensions to pthread_join()
pthread_key_t	Thread local storage key
pthread_mutexattr_t	Mutex creation attribute
pthread_mutex_t	Mutex (Mutual exclusion) synchronization primitive
pthread_once_t	Once time initialization control variable
pthread_option_np_t	Pthread run-time options structure
pthread_rwlockattr_t	Read/Write lock attribute
pthread_rwlock_t	Read/Write synchronization primitive
pthread_t	Pthread handle
pthread_id_np_t	Thread ID. For use as an integral type.
struct sched_param	Scheduling parameters (priority and policy)

Thread Declaration and Creation

Declaration:

```
pthread_t tid; /* thread handle */
```

Creation:

```
int pthread_create(pthread_t* tid,  
                  const pthread_attr_t* attr,  
                  void* (*start_routine)(void *),  
                  void *arg);
```

- Analogous to a combined **fork** and **exec** routine
- The routine creates a thread with the specified thread attributes.
- When **attr** is **NULL** the default thread attributes are used.
- Returns a thread id in **tid**.
- The new thread begins execution by calling **start_routine** with a single argument **arg**
- For example:

```
pthread_create (&tid, &tattr, (void *) tf, (void *) &x);
```

Thread Attributes

- Thread attributes can be set before a thread is created
 - If not set, the attributes are set to default values
- Typical usage:
 1. Declare a pthread attribute variable of the `pthread_attr_t` data type
 2. Initialize the attribute variable:
`pthread_attr_init(&attr)`
 3. Set an attribute value:
`pthread_attr_setattribute(&attr, value)`
 4. Destroy the attribute var if not longer needed:
`pthread_attr_destroy(&attr)`



ROYAL INSTITUTE
OF TECHNOLOGY

Thread Attributes (cont'd)

- Detachable or joinable thread
 - **detachstate** – indicates whether a thread is detached or joinable.
- Stack information
 - **stackaddr** – specifies the stack address (void pointer) of a thread created with this attributes object.
 - **stacksize** – specifies the minimum stack size, in bytes, of a thread created with this attributes object.
- Thread scheduling attributes
 - **schedparam** – specifies the scheduling parameters such as priority.
 - **schedpolicy** – specifies the scheduling policy (FIFO, round-robin, other).
 - **inheritsched** – specifies the inheritance of scheduling properties (either inherits or not).
 - **scope** – the contention-scope attribute specifies the contention scope (local or global) of a thread



Passing Arguments to a Thread

- **pthread_create** permits to pass one argument (a void pointer) to the thread start routine.
 - The argument must be passed by reference and cast to **void ***.
 - To pass multiple arguments, collect them into a structure and pass a pointer to that structure.
 - Important: the argument data structure is located in the parent thread's memory space and it must not be corrupted/modified until the thread has finished accessing it.



Thread Identifiers

- When created, a thread is assigned a unique thread identifier (handle) that is used to reference the thread

```
pthread_t tid;
```

```
...
```

```
pthread_create(&tid, ...)
```

```
...
```

```
pthread_join(tid, NULL);
```

- A thread can get its own handle:

```
pthread_t mytid = pthread_self()
```

- returns the handle of the calling thread

- To compare two threads, use

```
pthread_equal(tid1, tid2)
```

- returns 0 if different
- Should not use == to compare threads.

Thread Termination

A thread may terminate its execution in two ways

- Explicitly by executing

void pthread_exit(void* return_value);

- **return_value** – a single return value or **NULL**
- Analogous to **exit**
- If the current thread is the last thread then the process terminates
- The **exit** routine kills all threads and exits the process

- Implicitly by returning from the thread routine

return (status)

- Return from the thread routine is equivalent to calling pthread_exit
- Return from the initial thread main is the equivalent to calling exit

Join a Thread


- A thread can wait for its child thread to terminate, i.e. it joins the child thread by executing:

```
int pthread_join(pthread_t tid, void** value);
```

- **tid** – child's descriptor
 - **value** – the address of the return value
 - Analogous to the UNIX **wait**
 - Must specify thread. There is no wait any.
 - Current thread blocks until the child thread terminates
 - The return value of thread is returned in **value**
- All threads can be either detached or joinable.

Detached Threads

- Threads that are not joinable are called **detached threads**.
 - A parent does not need to join a child.
 - When a detached thread terminates, it is destroyed.



- A thread can be created as "detached"


```
pthread_attr_setdetachstate(&tattr, PTHREAD_CREATE_DETACHED)
pthread_create(&tid, &tattr, ...)
```
- A thread can be “detached” explicitly by calling


```
int pthread_detach(pthread_t tid);
```




ROYAL INSTITUTE
OF TECHNOLOGY

Example 1: Independent Threads

- Creates a number of threads and then joins them.
- Each thread estimates its life time.
- The main thread waits for the child threads to terminate and then prints the total life time.

ind.c on the course website under “Lectures”

```
#ifndef _REENTRANT
#define _REENTRANT
#endif
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <sys/times.h>
#define MAXTHREADS 8
void *tf (void *);
int n = MAXTHREADS, results[MAXTHREADS], total = 0;
int main (int argc, char *argv[]) {
    pthread_t tid[MAXTHREADS];
    long i;
    n = (argc > 1)? atoi(argv[1]) : MAXTHREADS;
    for (i = 0; i < n; i++) pthread_create (&tid[i], NULL, tf, (void *) i);
    for (i = 0; i < n; i++) pthread_join (tid[i], NULL); total +=
        results[i];
    printf ("Total: %d\n", total);
    return 1;
}
void *tf (void *num) {
    int i;
    long mynum = (long)num;
    struct tms buffer;
    clock_t time = times (&buffer);
    for (i = 0; i < 5; i++)
        printf ("I am thread %d of %d\n", mynum, n); sleep (rand() % 2);
    time = times (&buffer) - time;
    results[mynum] = (int) time;
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

Synchronization Mechanisms in Pthreads

- **Mutexs** – mutual exclusion locks
- **Condition variables** – queues of thread (used for condition synchronization)
- **Reader/Writer locks** – shared read and exclusive write locks
- **Semaphores** – an extension (POSIX 1b) to Pthreads API implemented on top of mutexes and condition variables:

`#include <semaphore.h>`

Link to `-lpthread`

Mutexs

- **Mutex** is an object of the `pthread_mutex_t` type used to protect critical sections of code, i.e. for mutual exclusion
 - Acts like a lock protecting access to shared data within a critical section: Lock, unlock and trylock operations
 - Before to be used, must be initialized.
 - A thread can set/get mutex attributes.
 - Should be destroyed when no longer needed.
- Declaration:
`pthread_mutex_t mutex;`



ROYAL INSTITUTE
OF TECHNOLOGY

Functions on Mutex

- **`pthread_mutex_init(&mutex, &attr)`**
 - Create and initialize a new mutex object, set its attributes.
 - The mutex is initially unlocked.
- **`pthread_mutex_destroy(mutex)`**
 - Destroy the mutex when no longer needed
- **`pthread_mutex_lock(&mutex)`**
 - Lock the mutex. If it is already locked, the call blocks the calling thread until the mutex is unlocked.
- **`pthread_mutex_trylock(&mutex)`**
 - Attempt to lock a mutex. Returns non-zero value if the mutex is already locked, otherwise returns 0.
- **`pthread_mutex_unlock(&mutex)`**
 - Unlock the mutex (if called by the owning thread)

Typical Usage

- Create and initialize mutex, ..., spawn threads, ..., lock the mutex, execute critical section, unlock the mutex, ... destroy the mutex
- For example:

```
pthread_mutex_t mutex;  
pthread_mutex_init(&mutex, NULL);  
...  
pthread_mutex_lock(&mutex);  
critical section;  
pthread_mutex_unlock(&mutex);  
non-critical section;
```

Example 2: Counters

counters.c on the course website under “Lectures”

```
#ifndef _REENTRANT
#define _REENTRANT
#endif
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5                /* default number of threads */
/* shared variables */
double total;
pthread_mutex_t lock;
void *Counter (void *null) {
    int i;
    double *result = (double *) calloc (1, sizeof (double));
    for (i = 0; i < 1000000; i++) *result = *result + (double) (random ()%100);
    pthread_mutex_lock (&lock);
    total += *result;
    pthread_mutex_unlock (&lock);
    pthread_exit ((void *) result);
}
int main (int argc, char *argv[]) {
    int n = NUM_THREADS;
    double *result;
    result = (double*)calloc(1, sizeof(double));
    pthread_t thread[NUM_THREADS];
    int t;
    if (argc > 1) n = atoi (argv[1]);
    if (n > NUM_THREADS || n < 1) n = NUM_THREADS;
    for (t = 0; t < n; t++) pthread_create (&thread[t], NULL, Counter, NULL);
    for (t = 0; t < n; t++) {
        pthread_join (thread[t], (void*) &result);
        printf ("Completed join with thread %d. Result =%f\n", t, *result);
    }
    printf ("The total = %f\n", total);
}
```



Condition Variables

- **A condition variable (a.k.a. queue variable)** is an object of the **pthread_cond_t** data type used for blocking and resuming threads holding mutexes
 - Used for condition synchronization.
 - A condition variable is always used in conjunction with a mutex lock.
- Declaration
pthread_cond_t cond;
- Operations on condition variables:
 - wait, signal, broadcast
 - Signal-and-continue signaling discipline

Functions on Condition Variables

- **`pthread_cond_init(&cond, &attr)`**
 - Create and initialize a new condition variable.
- **`pthread_cond_destroy(&cond)`**
 - Free the condition variable that is no longer needed.
- **`pthread_cond_wait(&cond, &mutex)`**
 - Wait on the condition variable, i.e. release the mutex (if owner) and place the calling thread to the tail on the cond var queue.
- **`pthread_cond_signal(&cond)`**
 - Signal the condition variable, i.e. move a waiting thread (if any) from the head of the cond var queue to the mutex queue.
- **`pthread_cond_broadcast(&cond)`**
 - Signal all: Awaken all threads waiting on the condition variable.

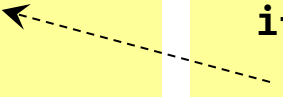
Typical Usage of A Condition Variable

- Take an action when $x == 0$

<await ($x == 0$) take_action();>

```
action() {  
    ...  
    pthread_mutex_lock(&m);  
    while (x != 0)  
        pthread_cond_wait(&cv, &m);  
    take_action();  
    pthread_mutex_unlock(&m);  
    ...  
}
```

```
counter() {  
    ...  
    pthread_mutex_lock(&m);  
    x--;  
    if (x == 0)  
        pthread_cond_signal(&cv);  
    pthread_mutex_unlock(&m);  
    ...  
}
```

A dashed arrow points from the pthread_cond_signal(&cv); line in the counter() function to the pthread_cond_wait(&cv, &m); line in the action() function, illustrating the signal-wait relationship.

Example 3: Summing Matrix Elements

```
#define _REENTRANT
#endif
#include <pthread.h>
#include <stdio.h>
#define SHARED 0
#define MAXSIZE 10000 /* maximum matrix size */
#define MAXWORKERS 8 /* maximum number of workers */
pthread_mutex_t barrier; /* mutex lock for the barrier */
pthread_cond_t go; /* condition variable for leaving */
int numWorkers;
int numArrived = 0; /* number who have arrived */
void Barrier() { /* a reusable counter barrier */
    pthread_mutex_lock(&barrier);
    numArrived++;
    if (numArrived == numWorkers) {
        numArrived = 0;
        pthread_cond_broadcast(&go);
    } else pthread_cond_wait(&go, &barrier);
    pthread_mutex_unlock(&barrier);
}
int size, stripSize; /* assume size is multiple of numWorkers */
int sums[MAXWORKERS];
int matrix[MAXSIZE][MAXSIZE];
void *Worker(void *);
```

matrixSum.c on the course website under “Lectures”



ROYAL INSTITUTE
OF TECHNOLOGY

Summing Matrix Elements (cont'd)

```
/* read command line, initialize, and create threads */
int main( int argc, char *argv[] ) {
    int i, j;
    pthread_attr_t attr;
    pthread_t workerid[MAXWORKERS];
    /* set global thread attributes */
    pthread_attr_init(&attr);
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
    /* initialize mutex and condition variable */
    pthread_mutex_init(&barrier, NULL);
    pthread_cond_init(&go, NULL);
    /* read command line args if any */
    size = (argc > 1)? atoi(argv[1]) : MAXSIZE;
    numWorkers = (argc > 2)? atoi(argv[2]) : MAXWORKERS;
    if (size > MAXSIZE) size = MAXSIZE;
    if (numWorkers > MAXWORKERS) numWorkers = MAXWORKERS;
    stripSize = size/numWorkers;
    /* initialize the matrix */
    for (i = 0; i < size; i++) {
        printf("[ ");
        for (j = 0; j < size; j++) {
            matrix[i][j] = rand()%99;
            printf(" %d", matrix[i][j]);
        }
        printf(" ]\n");
    }
    /* create the workers, then exit */
    for (i = 0; i < numWorkers; i++)
        pthread_create(&workerid[i], &attr, Worker, (void *) i);
    pthread_exit(NULL);
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

Summing Matrix Elements (cont'd)

```
/* Each worker sums the values in one strip of the matrix.
   After a barrier, worker(0) computes and prints the total */
void *Worker(void *arg) {
    int myid = (int) arg;
    int total, i, j, first, last;

    printf("worker %d (pthread id %d) has started\n",
           myid, pthread_self());

    /* determine first and last rows of my strip */
    first = myid * stripSize;
    last = (myid == numWorkers - 1) ? (size - 1) :
                                              (first + stripSize - 1);

    /* sum values in my strip */
    total = 0;
    for (i = first; i <= last; i++)
        for (j = 0; j < size; j++)
            total += matrix[i][j];
    sums[myid] = total;
    Barrier();
    if (myid == 0) {
        total = 0;
        for (i = 0; i < numWorkers; i++)
            total += sums[i];
        printf("the total is %d\n", total);
    }
}
```

Thread-Safe (Reentrant) Functions

- A multithreaded program must use **thread-safe (reentrant)** versions of standard functions.
 - A **re-entrant function** can have multiple simultaneous, interleaved, or nested invocations which will not interfere with each other.
 - A **thread-safe function** is either re-entrant or protected from multiple simultaneous execution by some form of mutual exclusion.
- To get declared reentrant (thread safe) versions of standard functions, define the **_REENTRANT** macro before any **include**:

```
#ifndef _REENTRANT
#define _REENTRANT
#endif
#include <stdio.h>
#include <pthread.h>
...
```



ROYAL INSTITUTE
OF TECHNOLOGY

Semaphore Extension

- A pthread **semaphore** is a special kind of integer object of the **sem_t** type that can take any nonnegative value and can be altered by **sem_wait** (P, decrement) and **sem_post** (V, increment) operations
 - If the semaphore is zero, the P operation blocks the calling thread until the semaphore is positive.
- **#include <semaphore.h>**
 - Defines semaphore types and functions

Semaphores (cont'd)

- Declaration: **sem_t sem;**
- Functions
 - **sem_wait(&sem)** – P(sem) – decrements the semaphore if it's positive, otherwise blocks the process until the semaphore is positive
 - **sem_post(sem)** – V(sem) – increments the semaphore
- For example:

```
sem_init(&sem, SHARED, 1); /* sem = 1; */  
...  
sem_wait(&sem);           /* P(sem) */  
Critical section;  
sem_post(&sem);           /* V(sem) */
```



ROYAL INSTITUTE
OF TECHNOLOGY

Example 4: Simple Producer/Consumer Using Semaphores

```
#ifndef _REENTRANT
#define _REENTRANT
#endif
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#include <pthread.h>
```

```
#include <semaphore.h>
```

```
#define SHARED 1
```

```
void *Producer(void *); /* the two threads */
```

```
void *Consumer(void *);
```

```
sem_t empty, full; /* the global semaphores */
```

```
int data; /* shared buffer */
```

```
int numIters;
```

```
/* main() -- read command line and create threads, then  
print result when the threads have quit */
```

```
int main( int argc, char *argv[] ) {
```

```
/* thread ids and attributes */
```

```
pthread_t pid, cid;
```

```
pthread_attr_t attr;
```

```
pthread_attr_init(&attr);
```

```
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
```

```
numIters = atoi(argv[1]);
```

```
sem_init(&empty, SHARED, 1); /* sem empty = 1 */
```

```
sem_init(&full, SHARED, 0); /* sem full = 0 */
```

```
printf("main started\n");
```

```
pthread_create(&pid, &attr, Producer, NULL);
```

```
pthread_create(&cid, &attr, Consumer, NULL);
```

```
pthread_join(pid, NULL);
```

```
pthread_join(cid, NULL);
```

```
printf("main done\n");
```

```
}
```

pc.sems.c on the course website under "Lectures"

Example4: Simple Producer/Consumer Using Semaphores (cont' d)

```
/* deposit 1, ..., numIters into the data buffer */
void *Producer( void *arg ) {
    int produced;
    printf("Producer created\n");
    for (produced = 0; produced < numIters; produced++) {
        sem_wait(&empty);
        data = produced;
        sem_post(&full);
    }
}

/* fetch numIters items from the buffer and sum them */
void *Consumer( void *arg ) {
    int total = 0, consumed;
    printf("Consumer created\n");
    for (consumed = 0; consumed < numIters; consumed++) {
        sem_wait(&full);
        total = total+data;
        sem_post(&empty);
    }
    printf("for %d iterations, the total is %d\n", numIters, total);
}
```