

ID1217 Concurrent Programming  
Lecture 8



ROYAL INSTITUTE  
OF TECHNOLOGY

Programming Models for Multicores  
and their Architectures

Vladimir Vlassov  
KTH/ICT/EECS




# Outline


- Towards multicore CPUs
- Programming model design space
- Multicore Architecture: The main ideas, cache coherency

# Moore's Law; CPU Improvements

- **Moore's law: 2X transistors / 18 months**
  - Smaller transistors => faster transistors (MHz -> GHz)
  - Many transistors => more on-chip state (caches)
  - Many transistors => advanced architectures (ILP+ +)
- 
- Processors have historically had the **performance improvement rate of 50-60% per year.**
  - Most of that due to technology increases improvements like **increasing clock frequency.**



# The truth...



The updated chart from "Data Processing in Exascale-Class Computing Systems" by Chuck Moore, presented at the 2011 Salishan Conference on High-speed Computing. Original data by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten, dotted line extrapolations by C. Moore


# The truth...

- Since the beginning of 2000, clock rate started to level off.
- **The clock rate of CPUs** from main vendors (Intel, AMD, IBM,...) **is no longer increasing, it's even declining.**
  - The top clock rate of any Intel processor was around 4 GHz, and now it's around 3-3.5 GHz.
  - It's definitely not improving.
- On the other hand **the Moore's law** that says that the number of transistors is doubling every 18 month, **still continues.**
  - Smaller transistors are faster.
  - **But we cannot increase the clock rate because the power consumption goes way up, and it's impossible to cool the chip.**



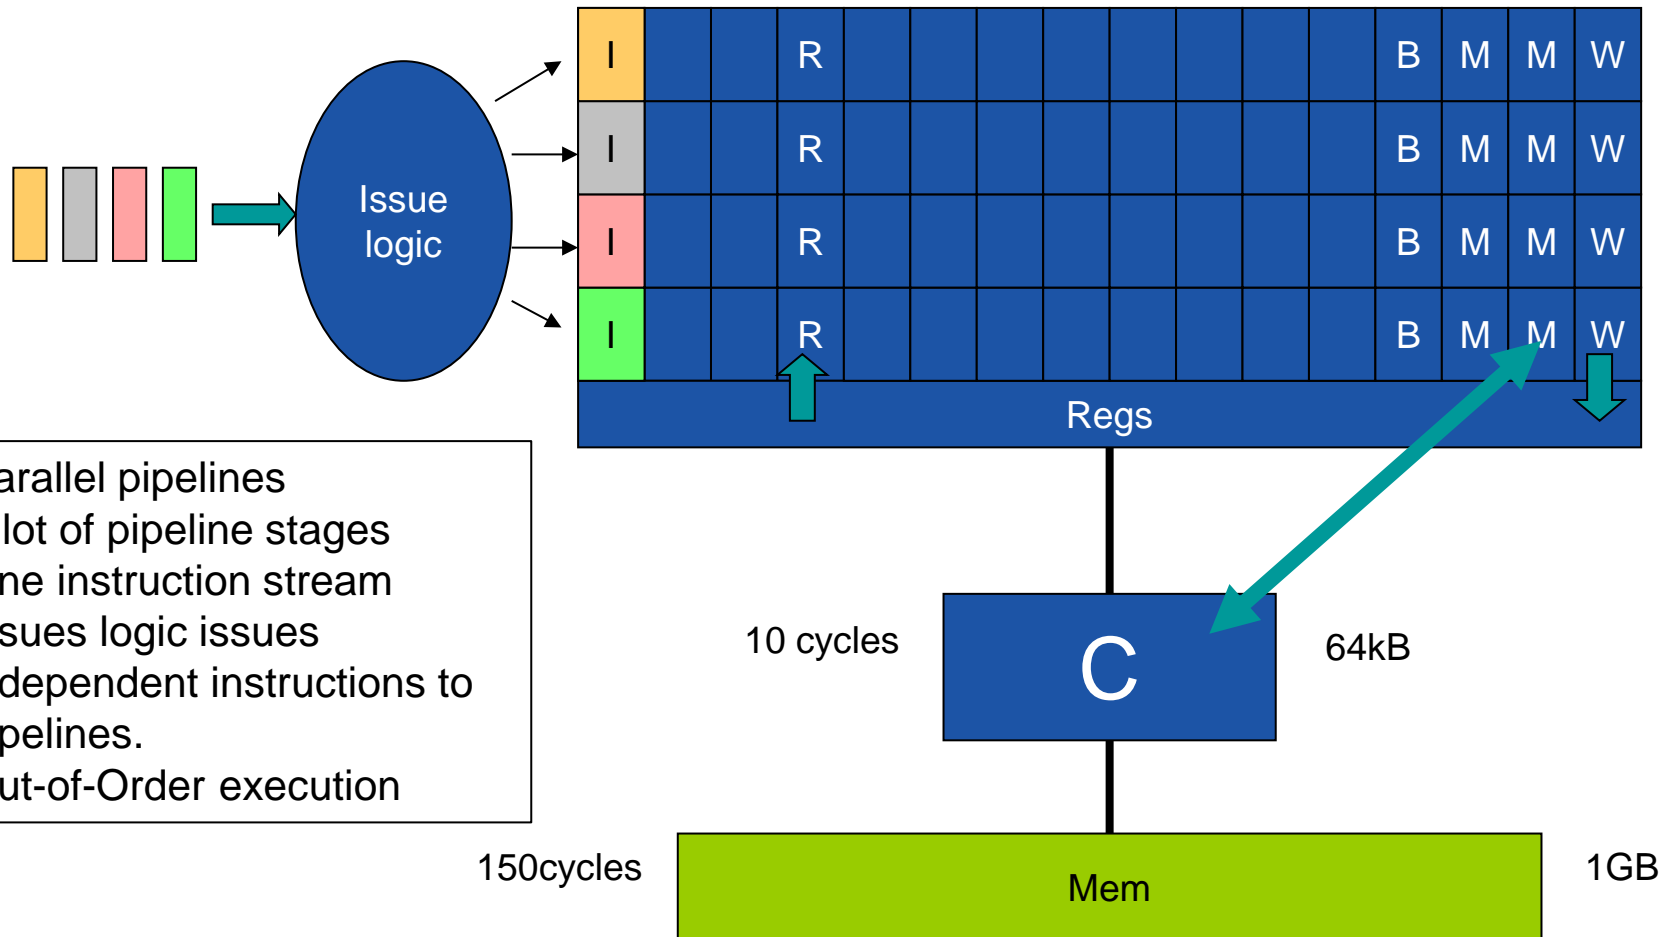
# Main steps in CPU development

- Single-chip, in-order issue pipelined CPU
  - Superscalar processors
  - Super-pipelined processors
  - Out-of-order (ooo) execution
  - Speculative computing
- 
- All in the quest for high ***Instruction-Level Parallelism (ILP)***





# A High-Performance Processor Architecture






# Could we continue to make large and complicated CPUs?...

- Diminishing return on..
  - pipeline stages
  - parallel pipelines
- Smaller reach →
  - Hard to add complexity, smaller caches
- Memory is often the bottleneck
- We cannot keep the power budget
- The IPL Wall
  - Instruction level parallelism requires independent instructions
  - Short basic blocks lead to little opportunity to find parallelism
- Industry-wide trend: ***Thread-Level Parallelism (TLP)***



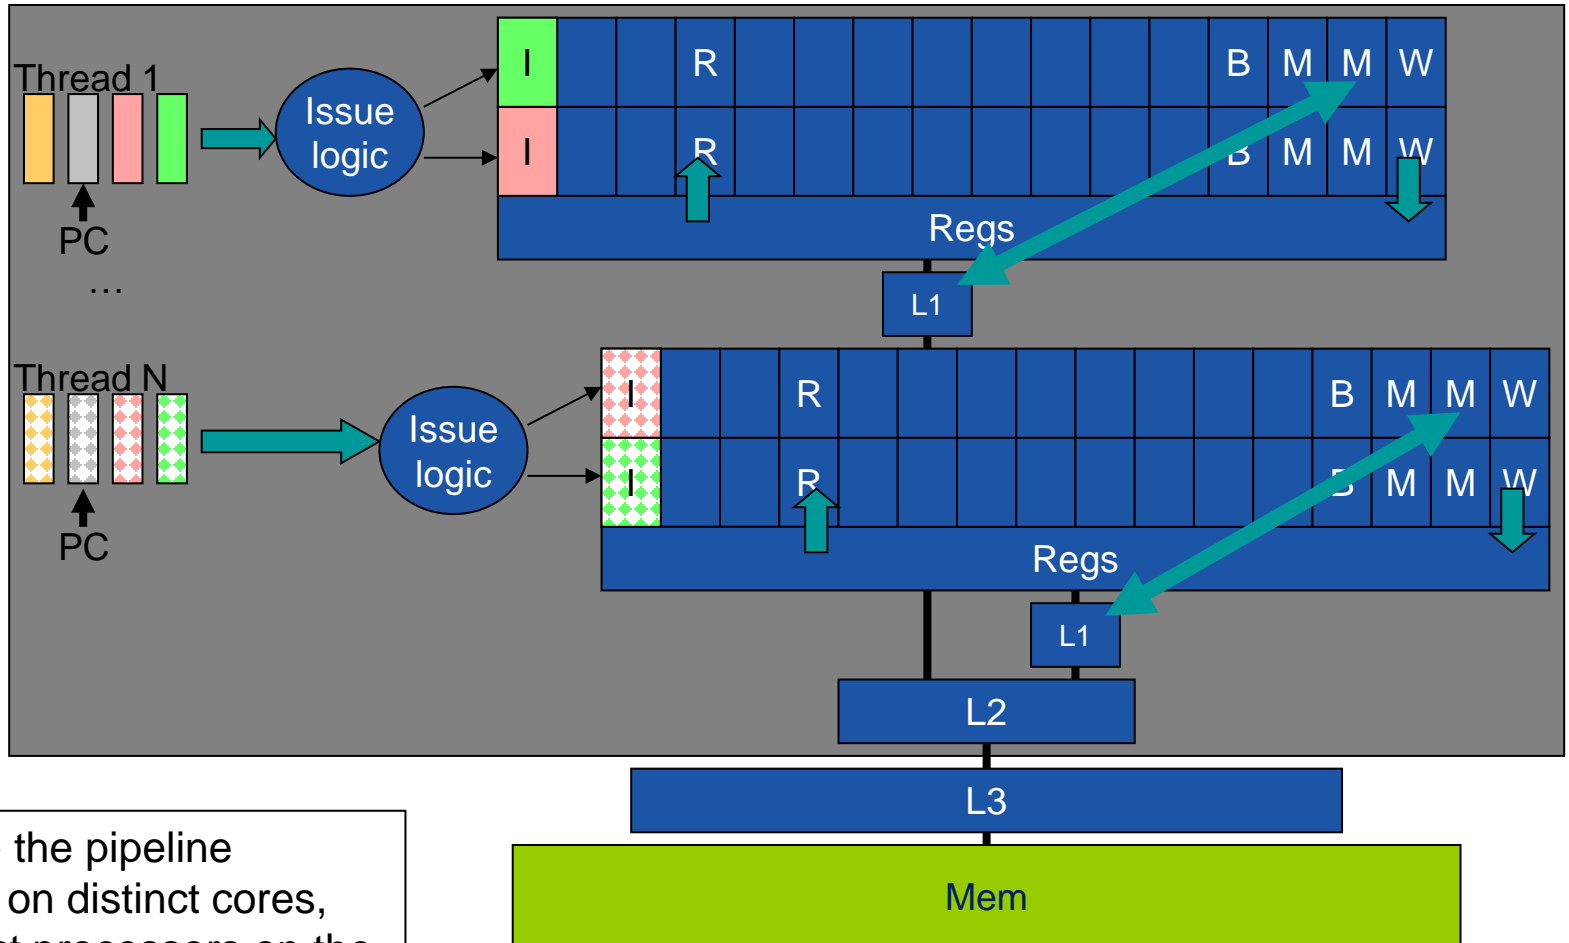


# SMT: Simultaneous Multithreading



- Several thread contexts in HW in the same time;
- Multiple threads sharing the pipeline structure

# CMP: Chip Multiprocessor (multi-core processor)




- Separate the pipeline structure on distinct cores, on distinct processors on the same chip, each executing a single or multiple threads.




# Why not earlier?


Old Mainframes



Super Minis:




Microprocessor:



Paradigm shift

Chip Multiprocessor (CMP):



Sequential execution

Parallel execution



# Many years experience

- We had a working horse with many years experience, and it worked well, performed better and better every year. Sort of “horse on steroids”, but now...





# New challenges...



+



= ?

- Now we are facing small little rabbits that are still going to carry around the same logs as the horse before and we need to coordinate them in some way.
  - So, that is, sort of, the reason for not doing this earlier.
- A *paradigm shift in programming*.
- And every software developer has to deal with it now.

# Why Multicore (1/2)

- Software developers were used to ever increasing performance, for free
  - Increasing clock frequencies
  - Increasingly advanced architectures exploiting Instruction level parallelism (ILP)
  - Larger cache memories
- However,
  - Clock frequencies no longer go up!
  - We have exhausted the available ILP in applications
- ***Parallelism is the only way increasing performance significantly***

## Why Multicore (2/2)

- Because these are the only chips we know how to build with the increasing number of transistors within the power envelope
- Because of performance reasons: potential of linear performance improvements with growing number of cores
  - If we could only program them...
- Because of power reasons



# Multicore Processors


## Huge design space:

- Homogeneous vs heterogeneous
- Low count, fat cores vs High count, thin cores
- Bus-based interconnect vs Network on chip
- Shared memory vs Message passing
- ...



# Example 1: AMD Opteron 6172 "Magny-Cours"


- 12 x x86 cores
- Private L1/L2 caches
- Shared L3 cache between 6 cores
- Software compatibility





## Example 2: Intel® Xeon X5650 "Gulftown"

- 6 cores (12 threads)
- 2-way hyperthreading (2 hw threads / core)
- Private L2 caches
- Shared L3 cache



See: <http://www.bit-tech.net/hardware/cpus/2010/03/31/amd-opteron-6174-vs-intel-xeon-x5650-review/>




## Example 3: Intel® Xeon Phi™



- 60 x86-64 cores with 4 threads each
- Cache coherence
- Ring interconnect
  - 512-bit
- Vector processing
  - 32 512-bit vector registers



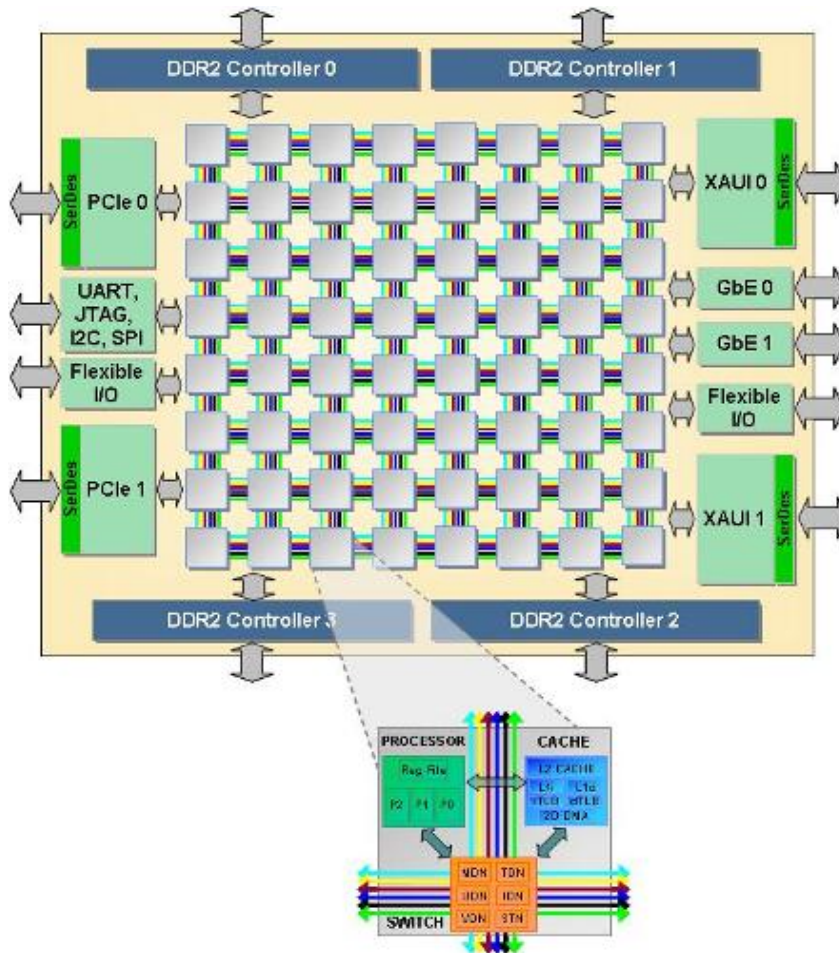
## Example 4: Oracle's UltraSPARC T5



- 16 cores, 128 threads (8 threads per core)
- A three-level cache architecture
  - L1 and L2 are private
  - L3 is shared
- Cross-bar interconnect
- Huge software support
  - Linux/Solaris
  - All kinds of languages



## Example 5: Tileria Tile-Gx



- 100 cores
- Cache coherent
- Standard OS (Linux)
- C/C++
- Pthreads
- Highly configurable
  - Programmable interconnect

# AMD Ryzen vs Intel Core i7

Make	Intel	AMD
Model	Core i7-6900K	Ryzen
Process	14nm	14nm
Cores	8	8
Threads	16	16
Base frequency	3.2GHz	3.4GHz
Turbo Boost	3.7GHz	Unknown
Turbo Boost Max	4GHz	Unknown
Cache	20MB	20MB
TDP	140 W	Unknown
Memory Type	DDR4	DDR4
Memory Channels	4	2
Socket	LGA-2011 V3	AM4
TDP	140 watts	95 watts
Price	\$1,100	Unknown

<http://www.pcadvisor.co.uk/new-product/pc-components/amd-zen-processor-release-date-price-specs-features-3643552/>



# What to expect from the future...

## By 2020 we will have...

- 200-2000 cores per chip\*
  - 35 Billion transistors
  - 400 mm<sup>2</sup> chip area
- Point-to-point on-chip interconnection networks
  - For design productivity
- Shared distributed memory
  - For software compatibility
- Non-reliable hardware
  - Software makes it reliable
- Highly dynamic workload
  - Workload adapting to architecture and
  - Architecture adapting to workload

\* International Technology Roadmap for Semiconductors



# The era of the “Supercomputers”

- The one with the most blinking lights wins
- The one with the strangest languages wins
- The niftier the better!



**Connection Machine 5  
by Thinking Machines  
Corporation**



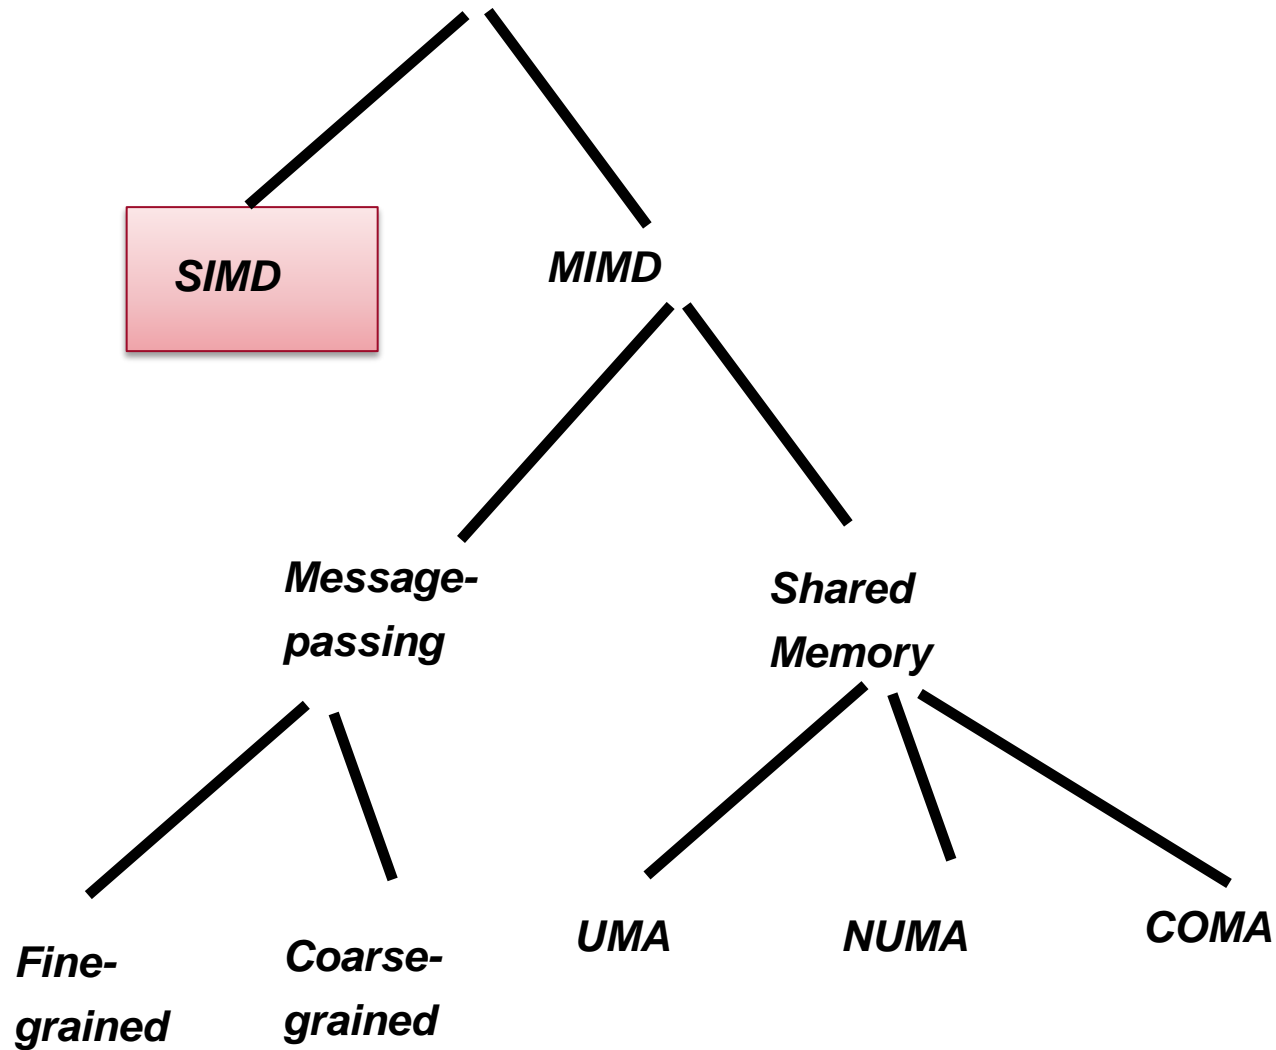
# Flynn's Taxonomy

{ Single, Multiple } Instruction + { Single, Multiple } Data

- SISD - Our good old simple CPUs
- SIMD – Vectors, "SSE", DSPs, CM-2,...
- MIMD – TLP, cluster, shared-mem MP,...
- MISD – Can't think of any...





# MP Taxonomy



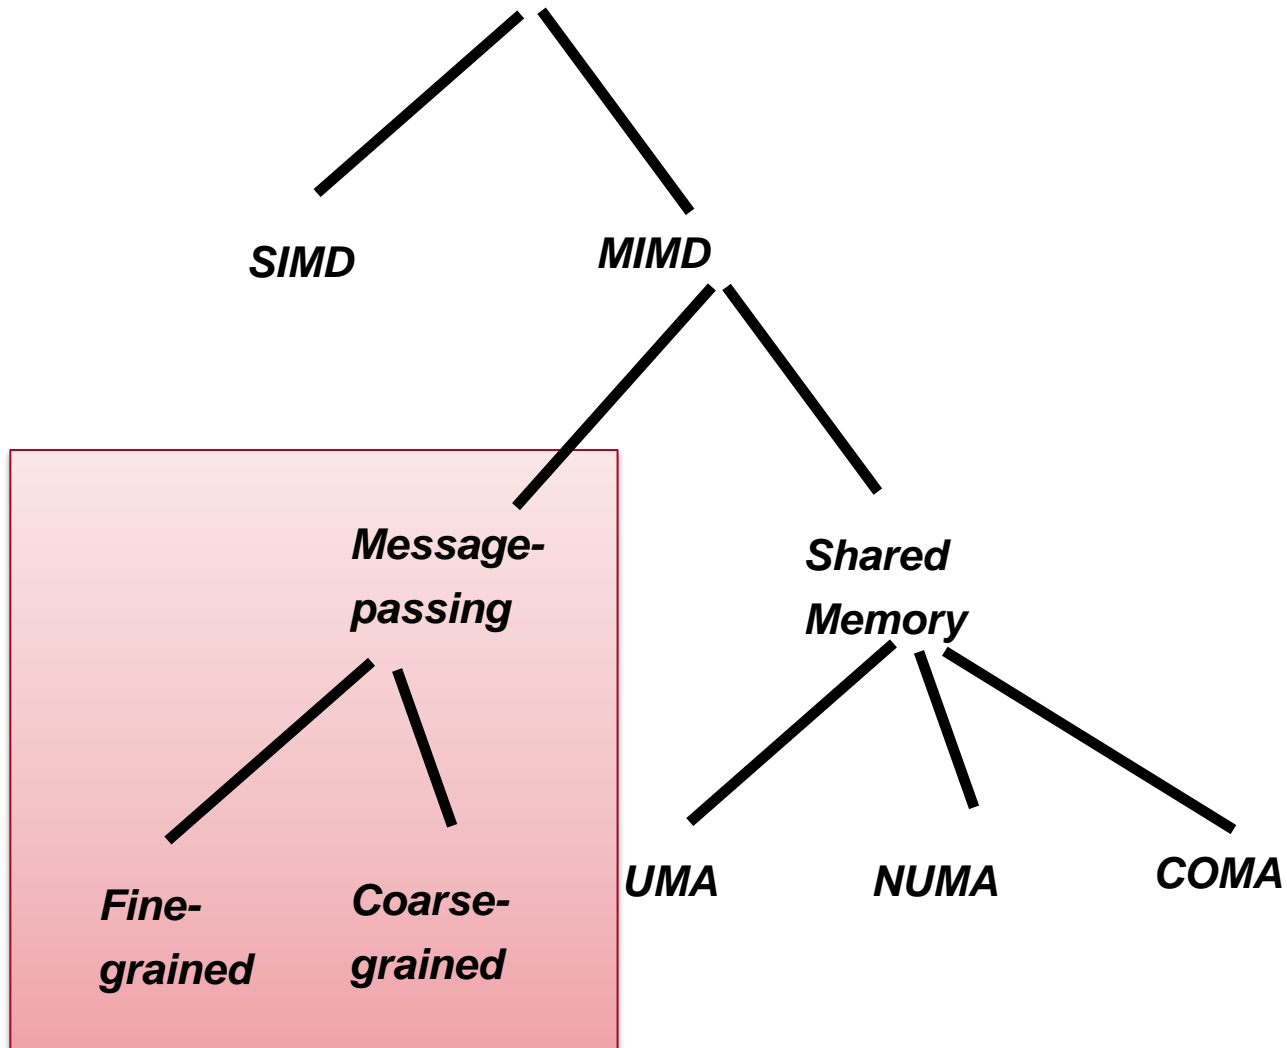


# *SIMD*






# Message-passing





# Message-passing Arch

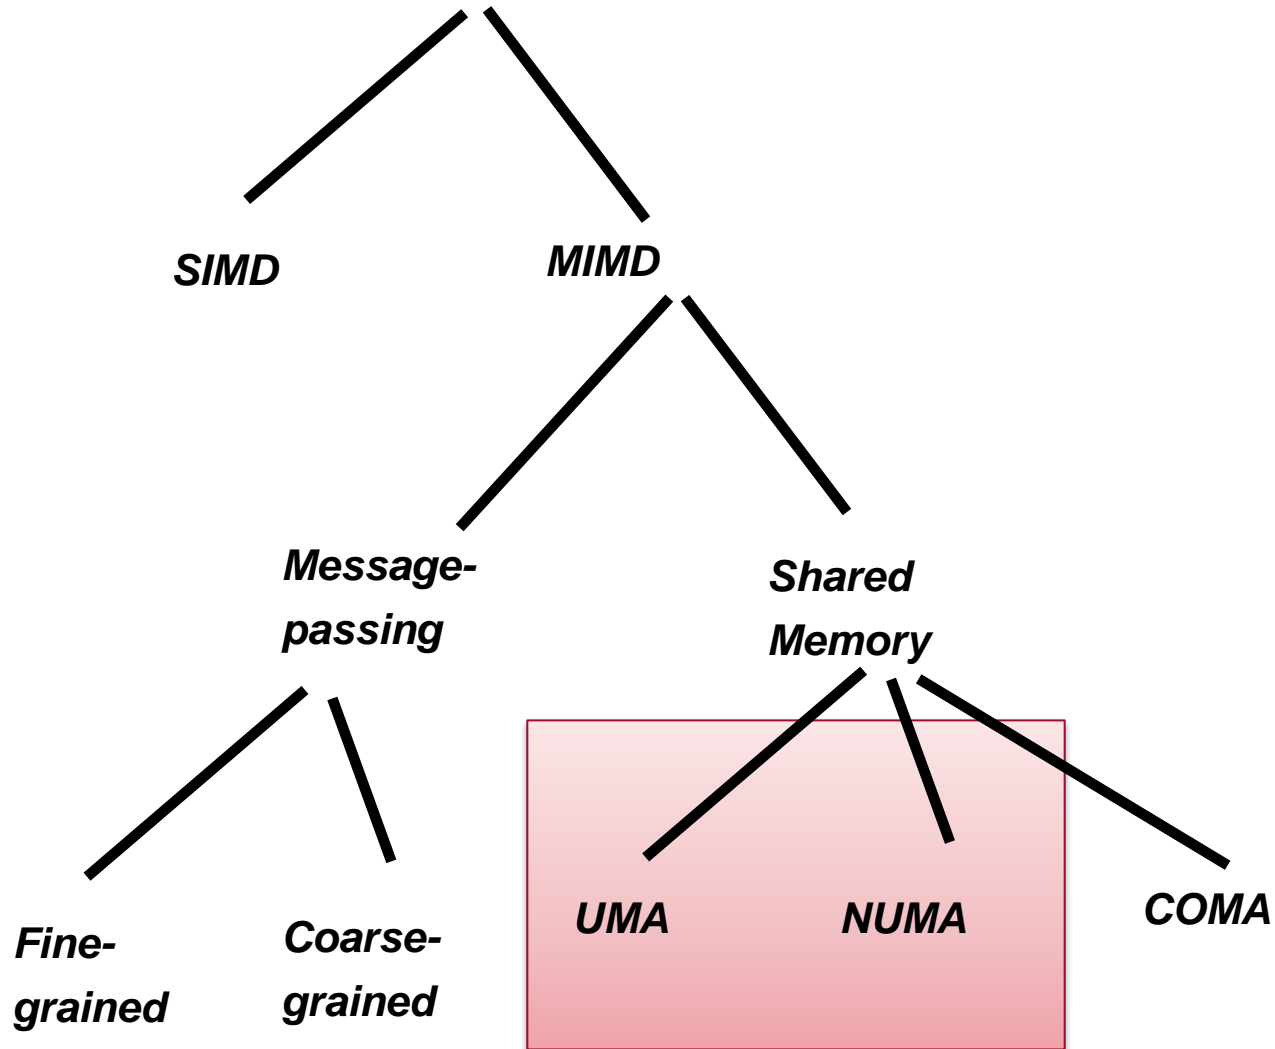
## MIMD



*Explicit  
Messages*




# Shared Memory MIMD






# The Shared Memory Model





# Adding Caches: More Concurrency



- Two coherence options
  - Snoop-based
  - Directory-based






# The Multicore Software Triad

## Desirable properties

- Application performance: Scalable up and down
- Software reliability: Do not introduce new bugs with parallelism
- Development time: Easy to write and maintain




# Algorithmic Complexity Measures

$T_P$  = execution time on  $P$  processors

$$T_1 = \text{work}$$

$$T_\infty = \text{span}^*$$



## LOWER BOUNDS

- $T_P \geq T_1/P$
- $T_P \geq T_\infty$

\* Also called *critical-path length* or *computational depth*.



# Parallelism

Defined as:  $T_1/T_\infty$

1.  $T_1/T_\infty$  is the average amount of work along each step of the critical path
2.  $T_1/T_\infty$  is **the maximum possible speedup** on any number of processors
3. Perfect linear speedup ( $T_1/T_P = P$ ) cannot be obtained for any  $P > T_1/T_\infty$



# Programming model design space

- General-purpose languages (GPL)
  - Native threads/processes
  - Data parallelism
  - SPMD
  - Loop-based parallelism
  - Task-based parallelism
- Domain-specific languages (DSL)
  - Erlang/SDL/Matlab/Labview
  - GPU-languages, e.g., CUDA
- Automatic parallelization
- Tool support

# A coding example

```
#include <stdlib.h>
#include <stdio.h>
long fib(long n) {
    long x, y;
    if (n < 2)
        return n;
    else {
        x = fib(n - 1);
        y = fib(n - 2);
    }
    return x + y;
}
int main(int argc, char *argv[]) {
    long n = (argc > 1) ? atoi(argv[1]) : 30;
    long result;
    result = fib(n);
    printf("fib(%ld) = %ld\n", n, result);
}
```

- Serial code
- Fibonacci numbers
- Stupid algorithm
- Good didactic example

# Explicit threading: pthreads

```
#ifndef _REENTRANT
#define _REENTRANT
#endif
#include <pthread.h>
#include <stdlib.h>
#include <stdio.h>
long fib(long n){
    long x, y;
    if (n < 2) return n;
    else {
        x = fib(n-1);
        y = fib(n-2);
    }
    return x + y;
}


typedef struct {
    long input;
    long output;
} thread_args;

void *thread_func(void *ptr){
    long i = ((thread_args *)ptr)->input;
    ((thread_args *)ptr)->output = fib(i);
    return NULL;
}
```

```
int main(int argc, char *argv[]){
    pthread_t thread;
    thread_args args;
    long n = (argc > 1)? atoi(argv[1]) : 30;
    long result;
    if (n < 30) result = fib(n);
    else {
        args.input = n-1;
        pthread_create(&thread, NULL, thread_func,
            (void*)&args);
        result = fib(n-2);
        pthread_join(thread, NULL);
        result += args.output;
    }
    printf("fib(%ld) = %ld\n", n, result);
}
```



# Consequences of pthreads



## Desirable properties:

- Scalable up and down
  - Code is fixed to two threads
- Do not introduce new bugs with parallelism
  - Code no longer modular
- Easy to write and maintain
  - 50 year leap backwards!



# Same example using OpenMP

```
#include <stdlib.h>
#include <stdio.h>
#include <omp.h>
long fib(long n){
    long x, y;
    if (n < 2) return n;
    else {
        #pragma omp task shared(x)
        x = fib(n-1);
        #pragma omp task shared(y)
        y = fib(n-2);
        #pragma omp taskwait
    }
    return x + y;
}
```

```
int main(int argc, char *argv[]){
    long n = (argc > 1)? atoi(argv[1]) : 30;
    long result;
    #pragma omp parallel
    #pragma omp single
    result = fib(n);
    printf("fib(%ld) = %ld\n", n, result);
}
```





# (some) Concurrency platforms

- Task-based languages:
  - OpenMP
    - also loop parallelism
  - Cilk(++)
- Data parallel languages
  - Ct
  - HPF
- Message passing libraries:
  - MPI
  - MCAPI
- Task-parallelism libraries:
  - Intel Threading Building Blocks
  - Microsoft TPL
  - Wool
  - Apple GCD

# Another coding example

```
#include <stdlib.h>
#include <stdio.h>
#define T 10
#define N 100
double X1[N+2], X2[N+2]; // initially zeros
int main(int argc, char *argv[]) {
    int t, i;
    int size = (argc > 1)? atoi(argv[1]) : N;
    if (size > N) size = N;
    int numIterations = (argc>2)? atoi(argv[2]) : T;
    X1[0] = X2[0] = X1[size+1] = X2[size+1] = 100;
    for (t = 0; t < numIterations; t++) {
        if ((t % 2) == 0) {
            for (i = 1; i < size+1; i++)
                X2[i] = (X1[i-1]+2*X1[i]+X1[i+1]) / 4;
        } else {
            for (i = 1; i < size+1; i++)
                X1[i] = (X2[i-1]+2*X2[i]+X2[i+1]) / 4;
        }
    }
}
```

$$X_i^{(t+1)} = \frac{X_{i-1}^{(t)} + 2X_i^{(t)} + X_{i+1}^{(t)}}{4}, \quad 0 < i \leq N, 0 \leq t < T$$

```
/* print the vector */
printf("[ ");
if (numIterations % 2 == 0)
    for (i = 0; i < size+2; i++)
        printf(" %f", X1[i]);
else
    for (i = 0; i < size+2; i++)
        printf(" %f", X2[i]);
printf(" ]\n");
}
```

# Same example using OpenMP

```
#include <stdlib.h>
#include <stdio.h>
#define T 10
#define N 100
double X1[N+2], X2[N+2]; // initially zeros
int main(int argc, char *argv[]) {
    int t, i;
    int size = (argc > 1)? atoi(argv[1]) : N;
    if (size > N) size = N;
    int numIterations = (argc>2)? atoi(argv[2]) : T;
    X1[0] = X2[0] = X1[size+1] = X2[size+1] = 100;
    #pragma omp parallel private(t,i) num_threads(4)
        for (t = 0; t < numIterations; t++) {
            if ((t % 2) == 0) {
                #pragma omp for
                    for (i = 1; i < size+1; i++)
                        X2[i] = (X1[i-1] + 2*X1[i] + X1[i+1]) / 4;
            } else {
                #pragma omp for
                    for (i = 1; i < size+1; i++)
                        X1[i] = (X2[i-1] + 2*X2[i] + X2[i+1]) / 4;
            }
        }
}
```

```
/* print the vector */
printf("[ ");
if (numIterations %2 == 0)
    for (i = 0; i < size+2; i++)
        printf(" %f", X1[i]);
else
    for (i = 0; i < size+2; i++)
        printf(" %f", X2[i]);
printf(" ]\n");
}
```

# Same example in "MPI"

```
MPI_Comm_size(MPI_COMM_WORLD, &p);
MPI_Comm_rank(MPI_COMM_WORLD, &id);

int mystart = id * N / p + 1;
int myend = (id + 1) * N / p + 1;

for (t = 0; t < T; t++) {
    if ((t % 2) == 0) {
        for (i = mystart; i < myend; i++)
            X2[i] = (X1[i-1]+2*X1[i]+X1[i+1])/4;
        if (id != (p-1)) {
            send(X2[myend-1], id+1, sizeof(double));
            receive(X2[mystart-1], id-1, sizeof(double));
            /* blocking send and receive */
        }
    }
}
```

```
if (id != 0) {
    send(X2[mystart], id-1, sizeof(double));
    receive(X2[myend], id+1, sizeof(double));
    /* blocking send and receive */
}
} else {
    for (i = mystart; i < myend; i++)
        X1[i] = (X2[i-1]+2*X2[i]+X2[i+1])/4;
    if (id != (p-1)) {
        send(X1[myend-1], id+1, sizeof(double));
        receives(X1[mystart-1], id-1, sizeof(double));
        /* blocking send and receive */
    }
    if (id != 0) {
        send(X1[mystart], id-1, sizeof(double));
        receive(X1[myend], id+1, sizeof(double));
        /* blocking send and receive */
    }
}
}
```

# Sending a message

```
MPI_Send(&N, 1, MPI_INT, i, tag,  
MPI_COMM_WORLD);
```

- *&N* – address to send buffer
  - *1* – the number of elements to send
  - *MPI\_INT* – the datatype of message
  - *i* – the receiving process
  - *tag* – used to label this message
  - *MPI\_COMM\_WORLD* – Communicator
- 
- Buffered and non-blocking

# Receiving a message

```
MPI_Recv(&tmp, 1, MPI_INT, i,  
        tag, MPI_COMM_WORLD, &status);
```

- *&tmp* – address to receive buffer
- *1* – the *maximum* number of elements to receive
- *MPI\_INT* – the datatype of message
- *i* – the sending process
- *tag* – used to label this message. This must match the corresponding tag in *MPI\_Send*
- *status* – information about the message received
- Blocking




ROYAL INSTITUTE  
OF TECHNOLOGY

# Multicore Architecture \* \* \*


The main ideas and future trends

# The Shared Memory Model

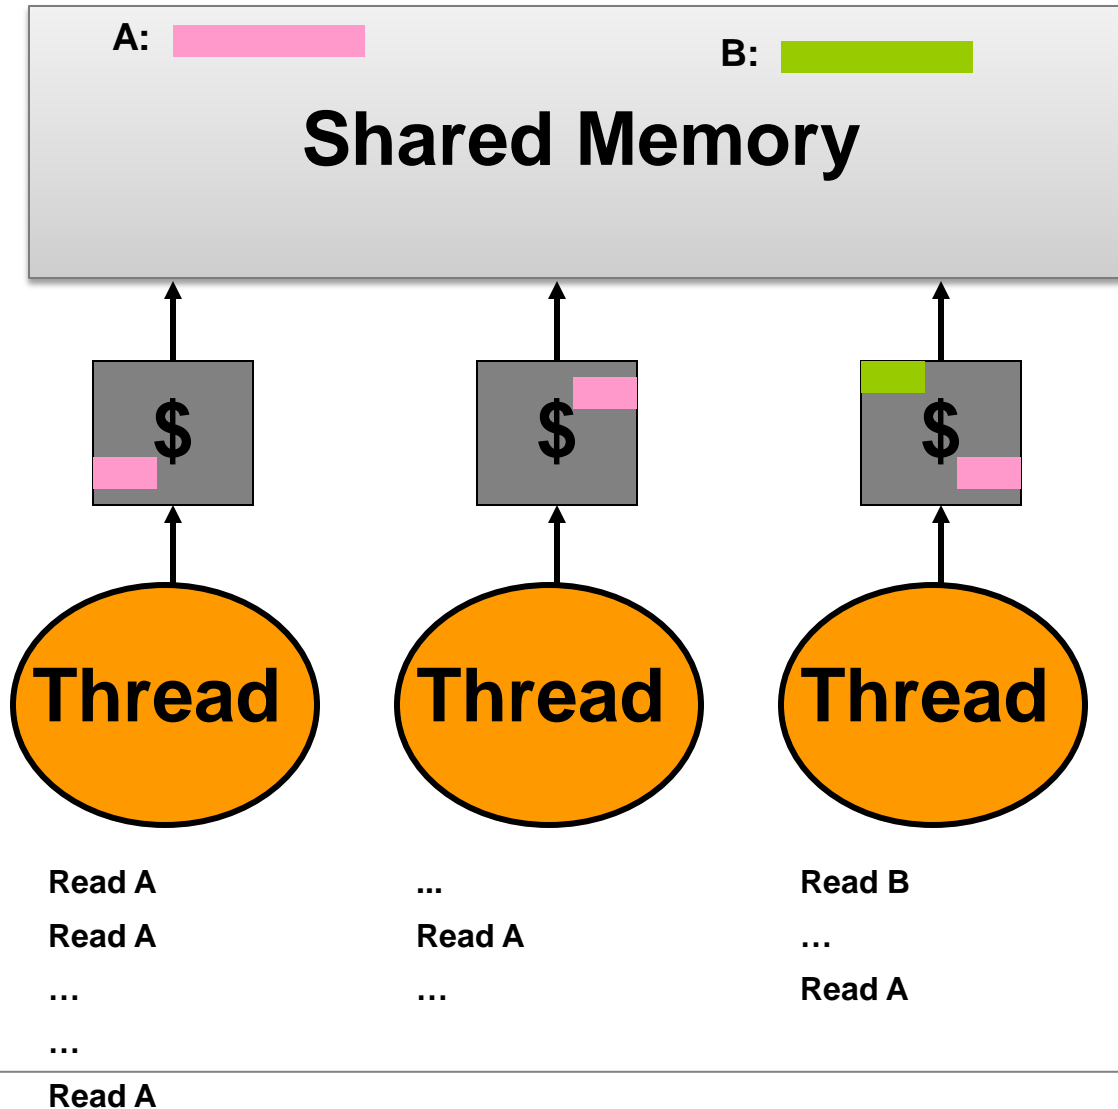




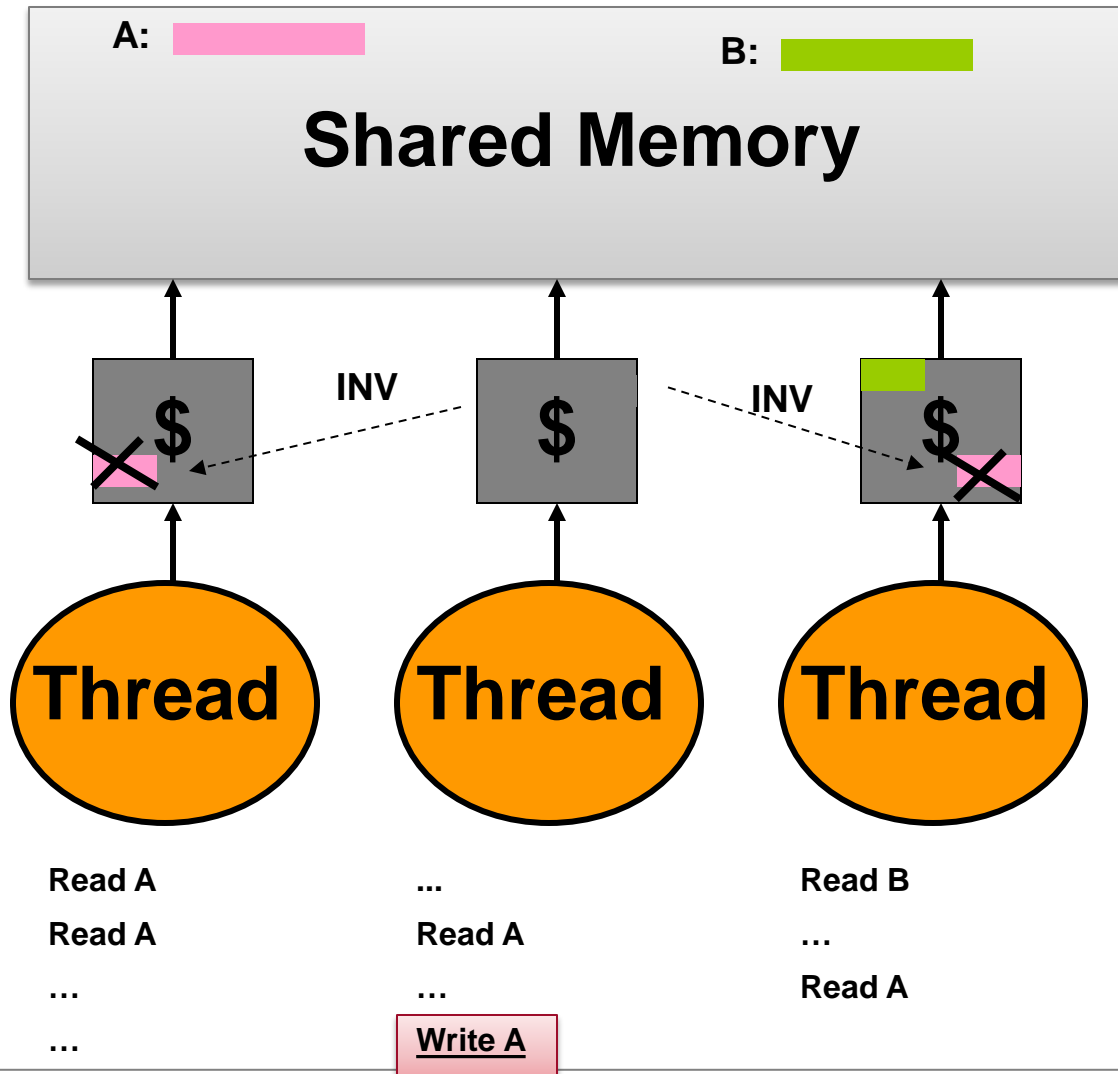
# Adding Caches: More Concurrency



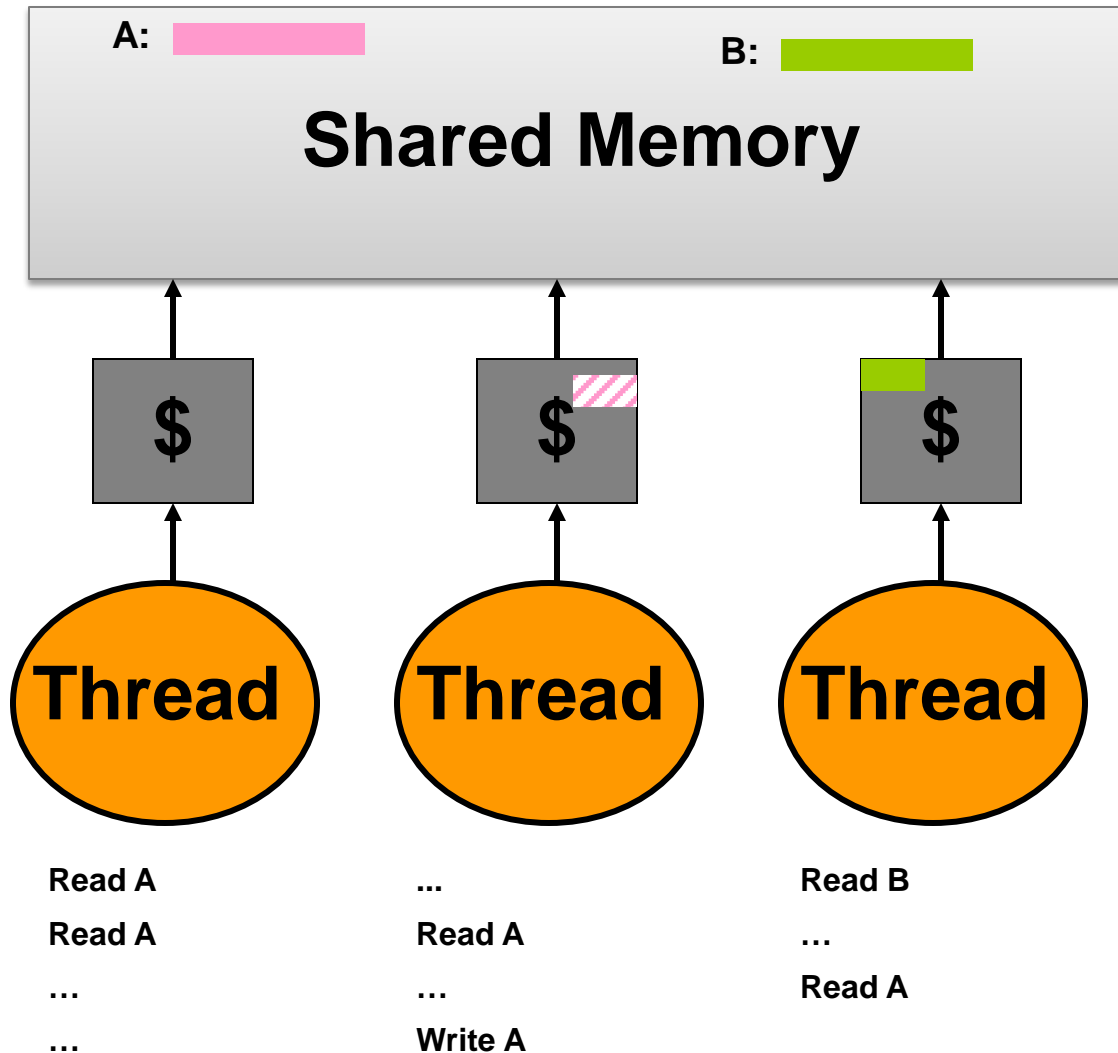
# Automatic Replication of Data



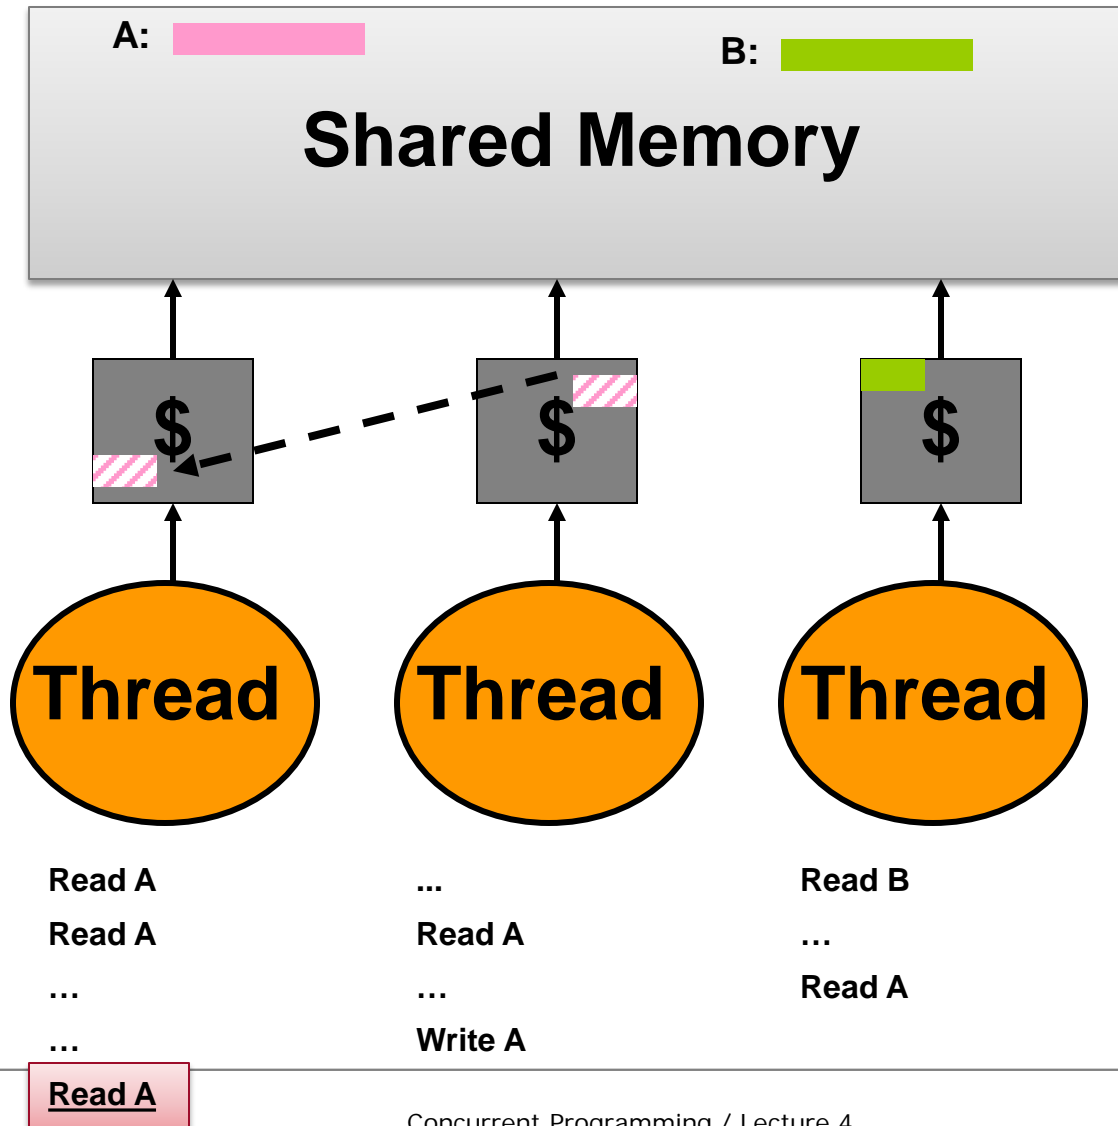
# The Cache Coherent Memory System



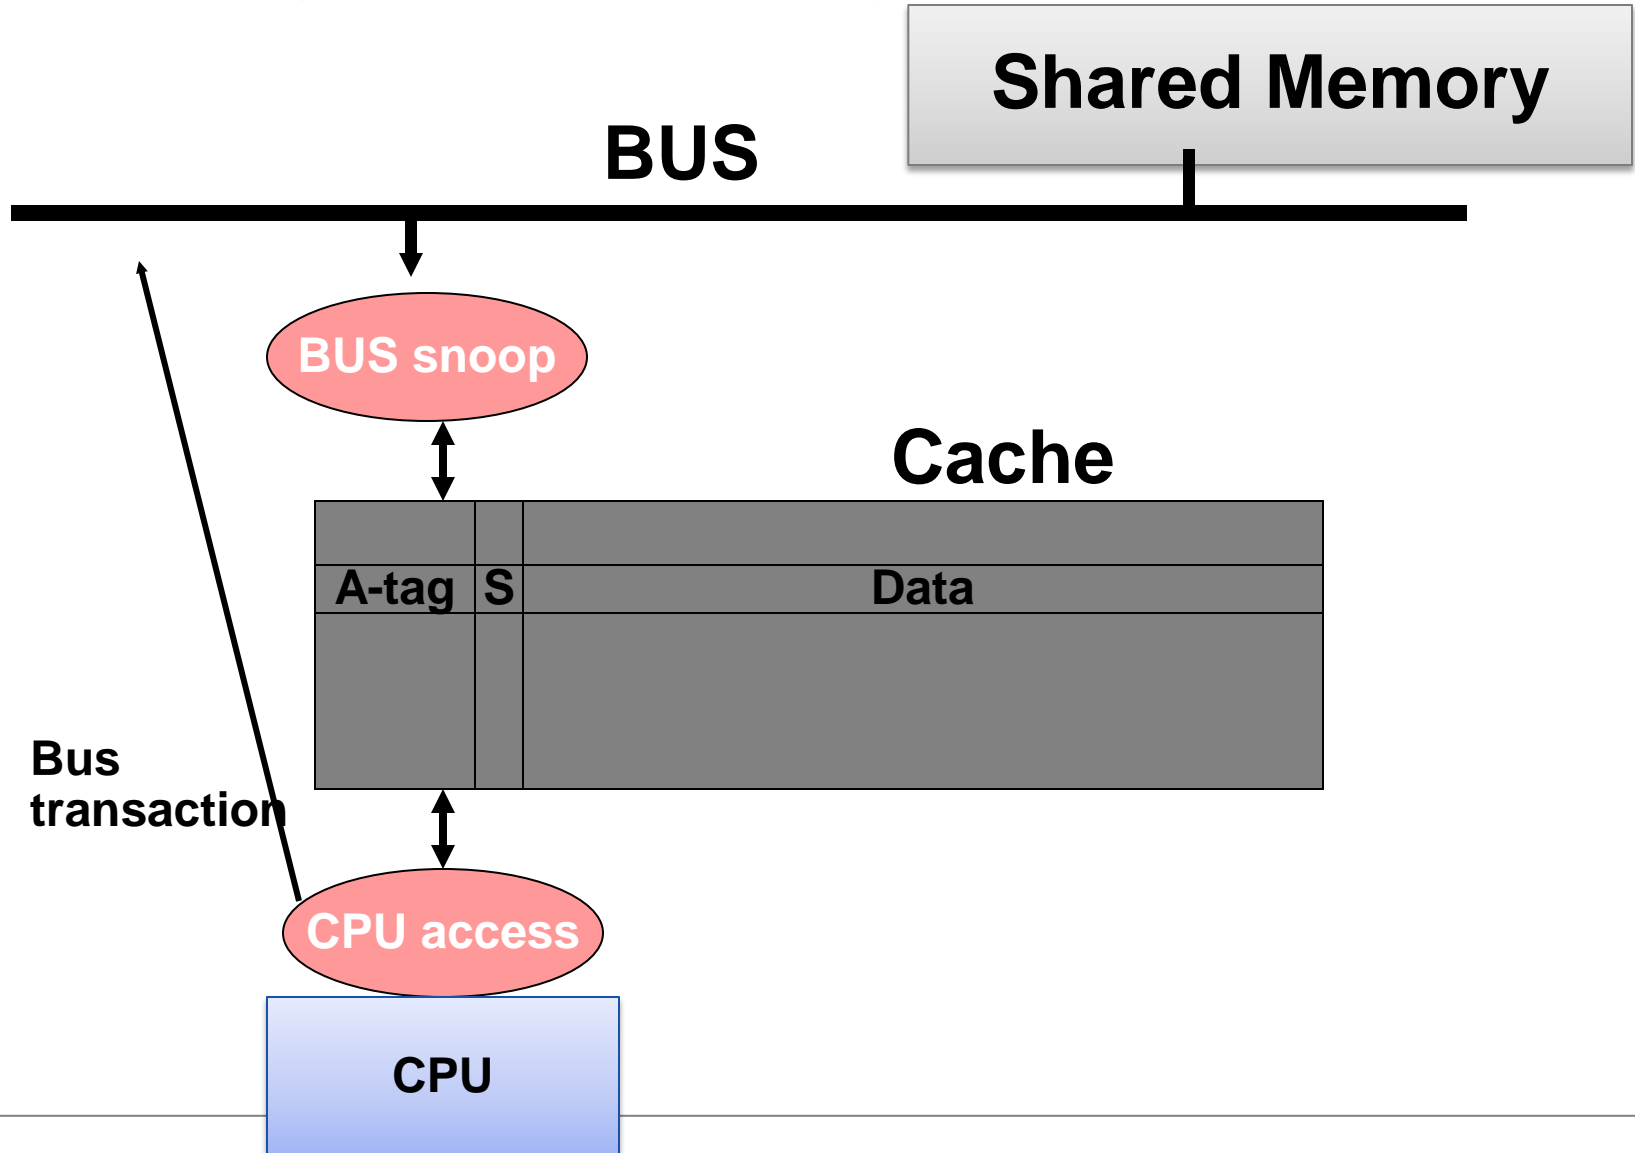
# The Cache Coherent Memory System



# The Cache Coherent Memory System



# Snoop-based Protocol Implementation



# MOSI Protocol, 1(2)

## Bus Transactions

**BUSrts: ReadtoShare** (reading the data with the intention to read it)

**BUSrtw, ReadToWrite** (reading the data with the intention to modify it)

**BUSwb**: Writing data back to memory

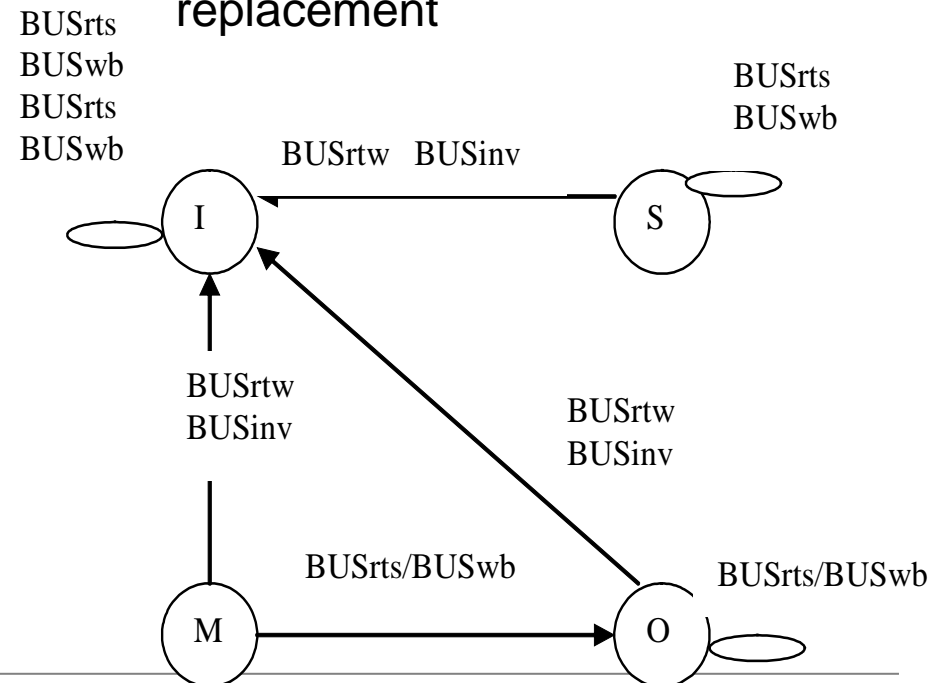
**BUSinv**: Invalidating other caches copies

## CPU-Initiated Events

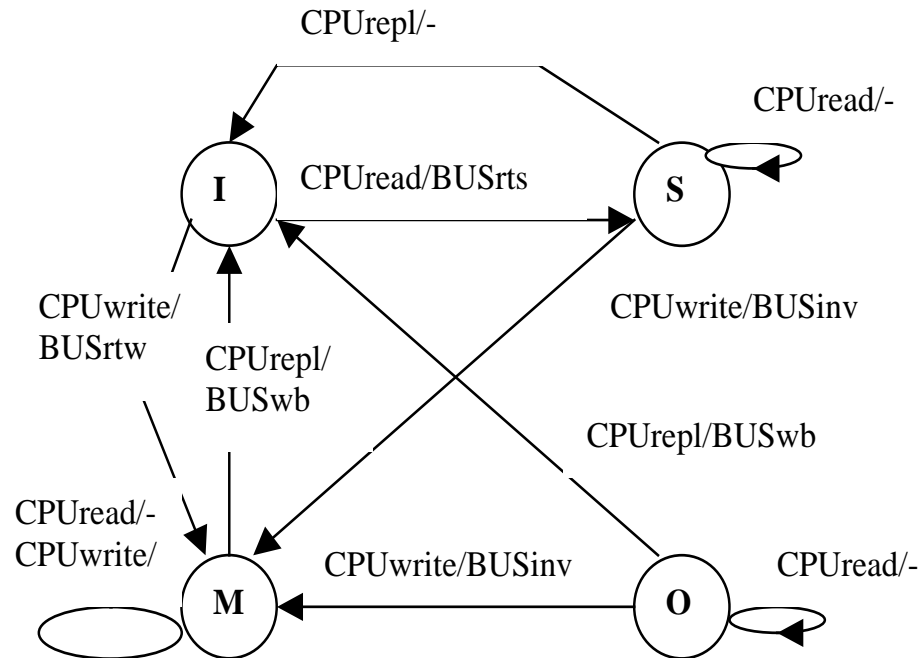
**CPUwrite**: Caused by a store miss

**CPUread** Caused by a loadmiss

**CPUrepl**: Caused by a replacement

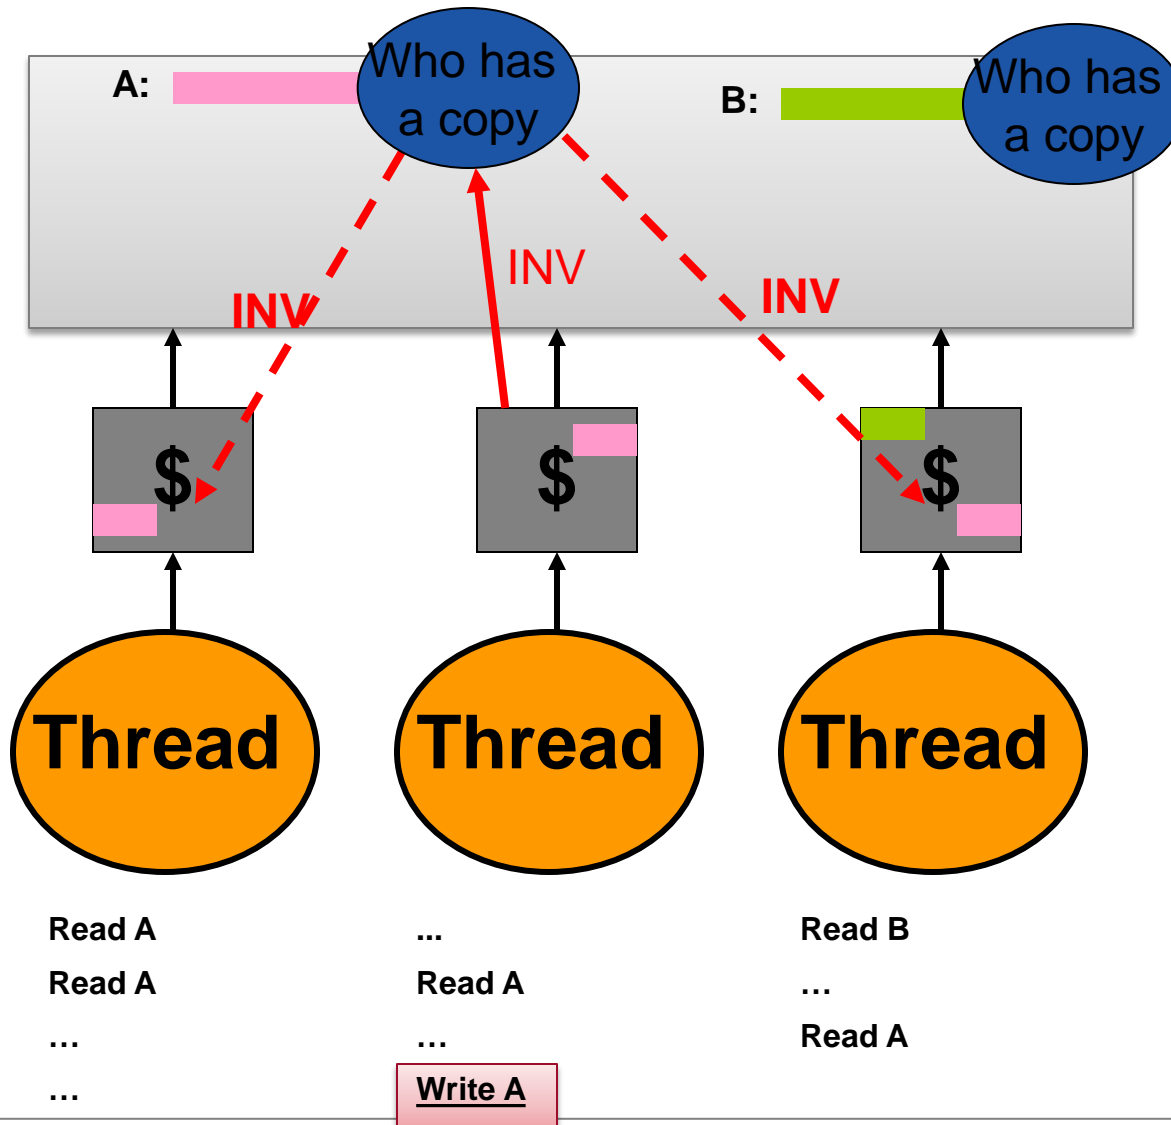


## MOSI Protocol 2 (2)





# Directory-based snooping



# Why do you miss in a cache?

- Capacity misses – *the cache is too small*
- Conflict misses – *the cache organization is not perfect*
- Compulsory misses – *touching the data for the first time*
- Coherence misses – *caused by communication*

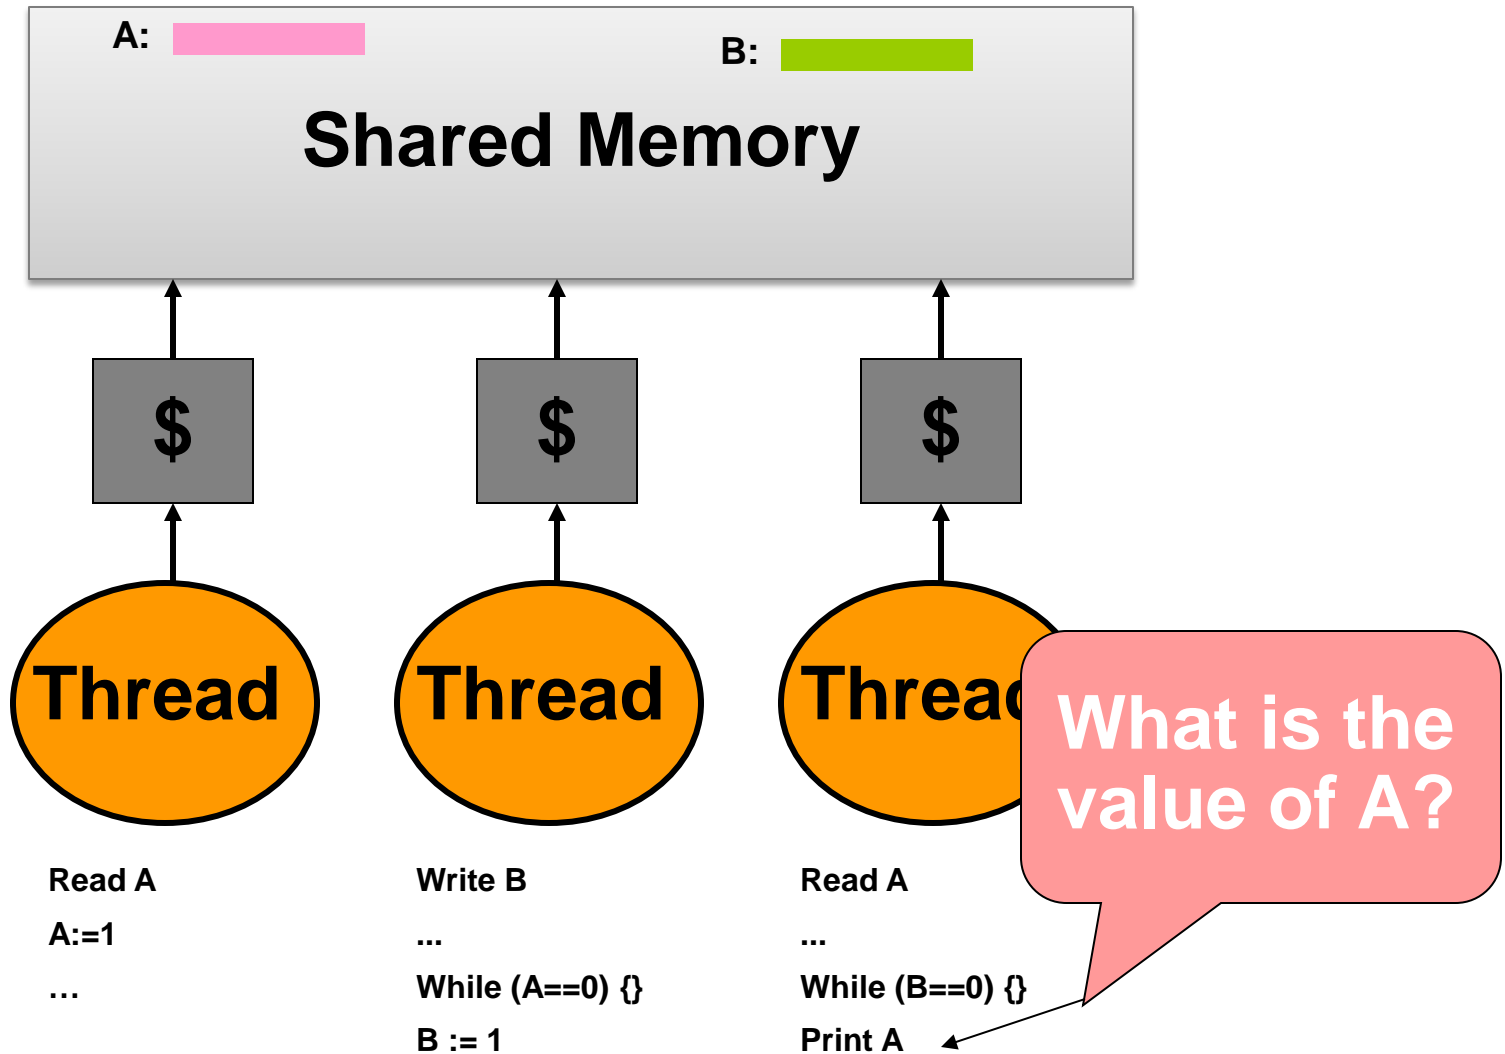
# Avoiding cache misses?

- **Capacity misses** – the cache is too small
  - larger caches, selective cacheing, prefetch cache, ...
- **Conflict misses** – the cache organization is not perfect
  - more associativity, skewed caches, victim caches, ...
- **Compulsory misses** – touching the data for the first time
  - larger cache lines, SW prefetching, HW prefetching, ...
- **Coherence misses** – caused by communication
  - update-based coherence, migr. optimization,...

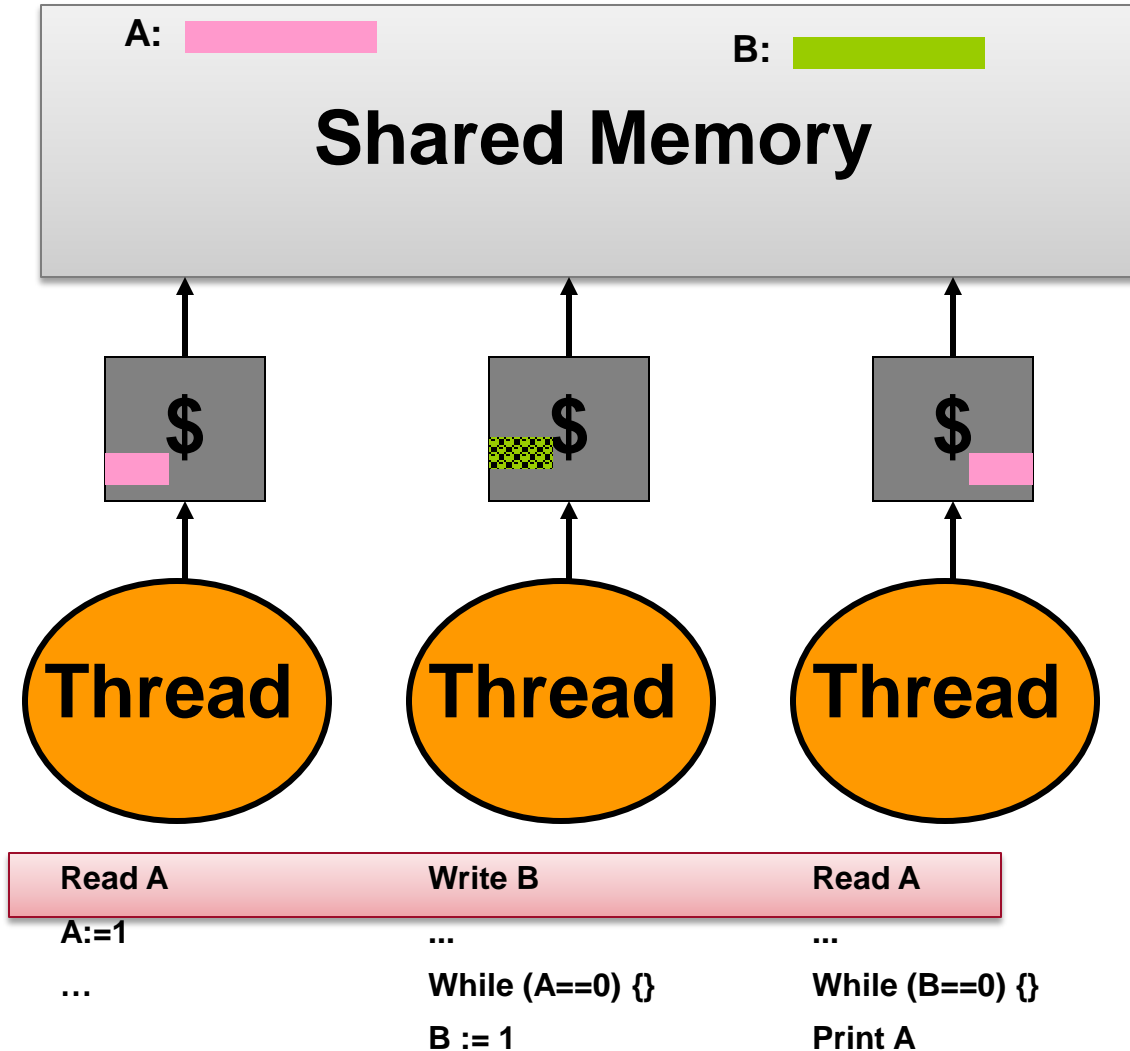
## So Far...

- Coherent shared memory
- “Snoopy-based coherence protocol”
  - All global accesses are broadcasted to all caches
  - Cache lines are automatically invalidated/fetched
- Ensure ordering and serialization for a single cache line

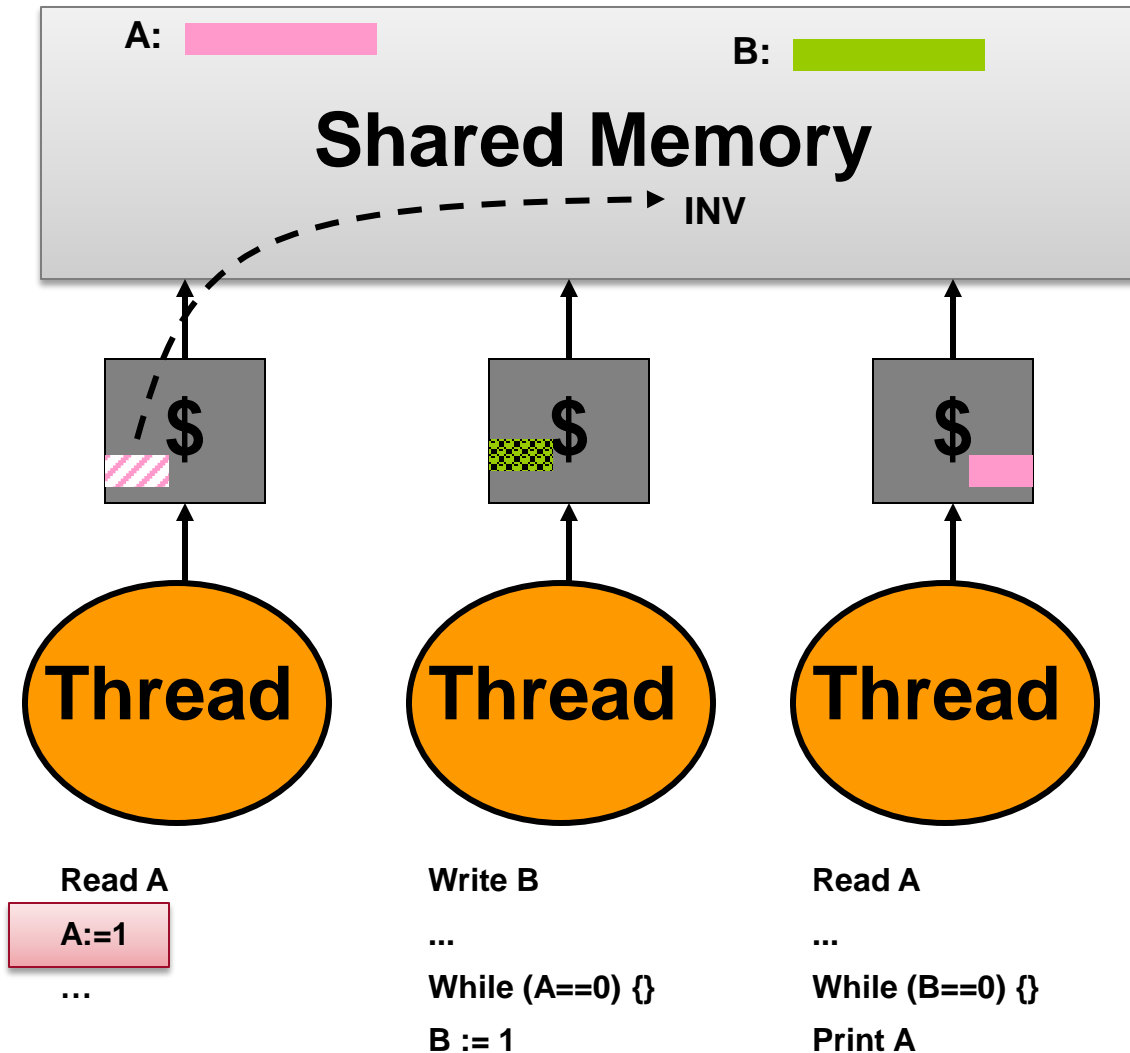
# Example1: Causal Correctness Issues



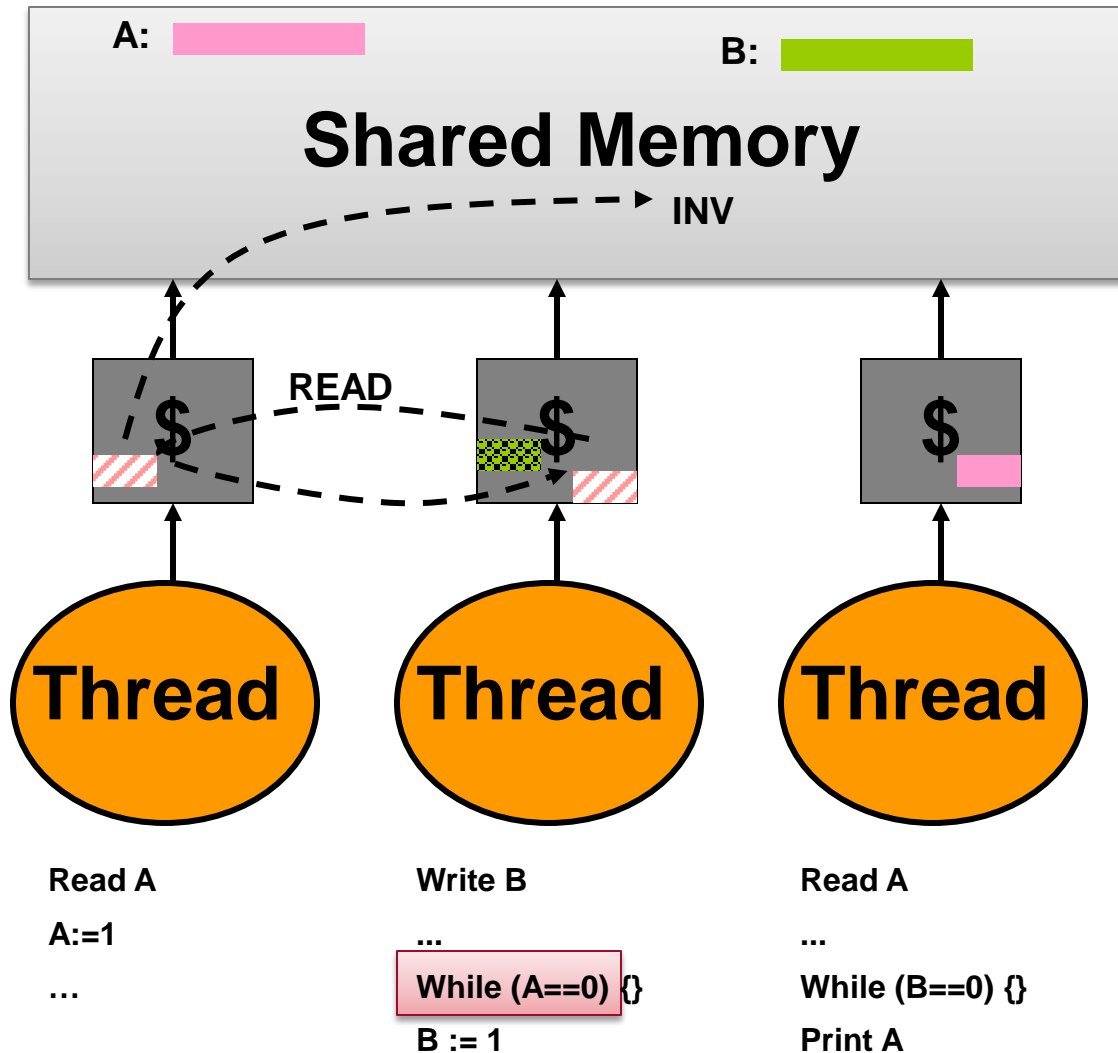
# Example1: Causal Correctness Issues



# Example1: Causal Correctness Issues

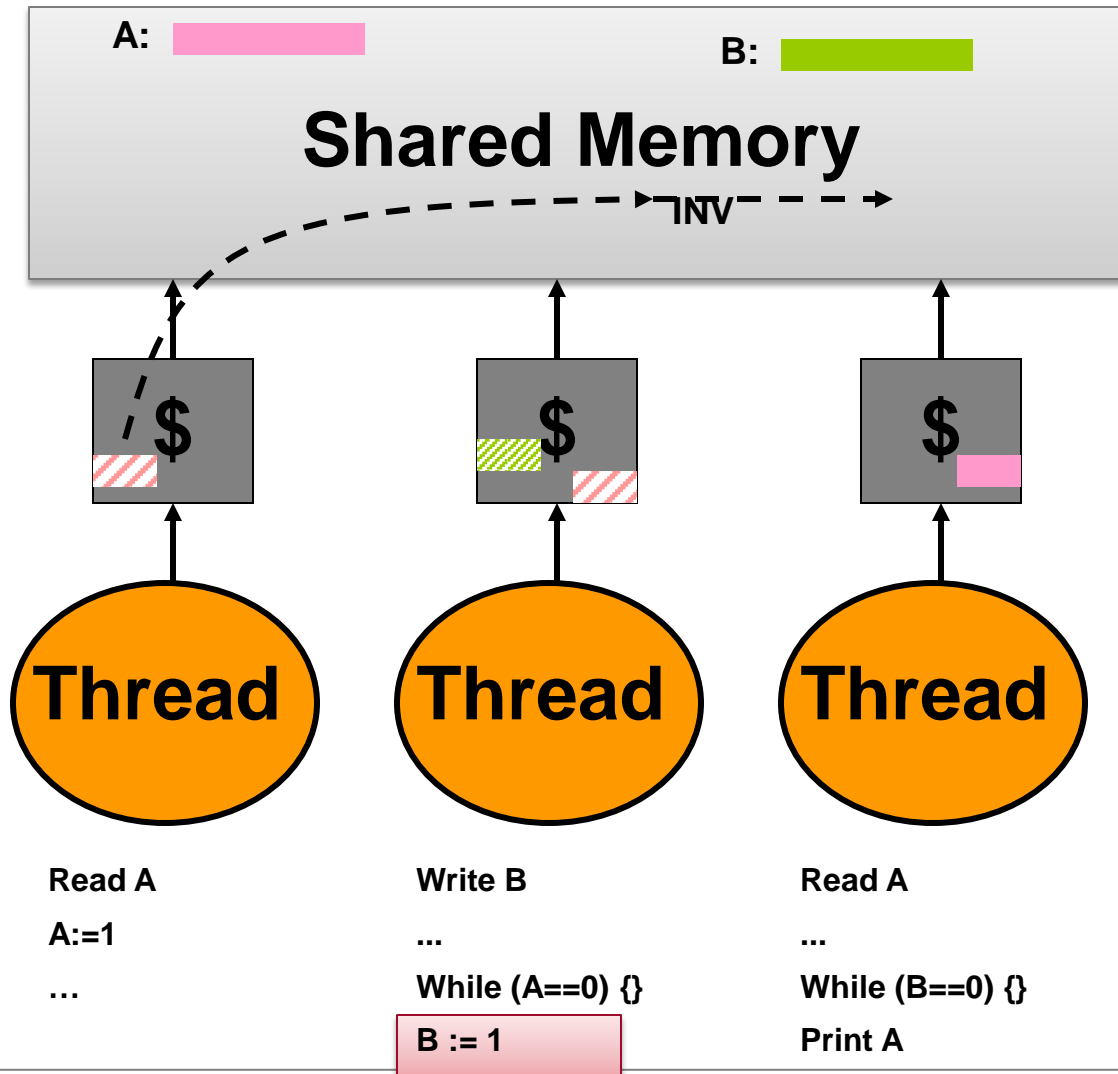


# Example1: Causal Correctness Issues

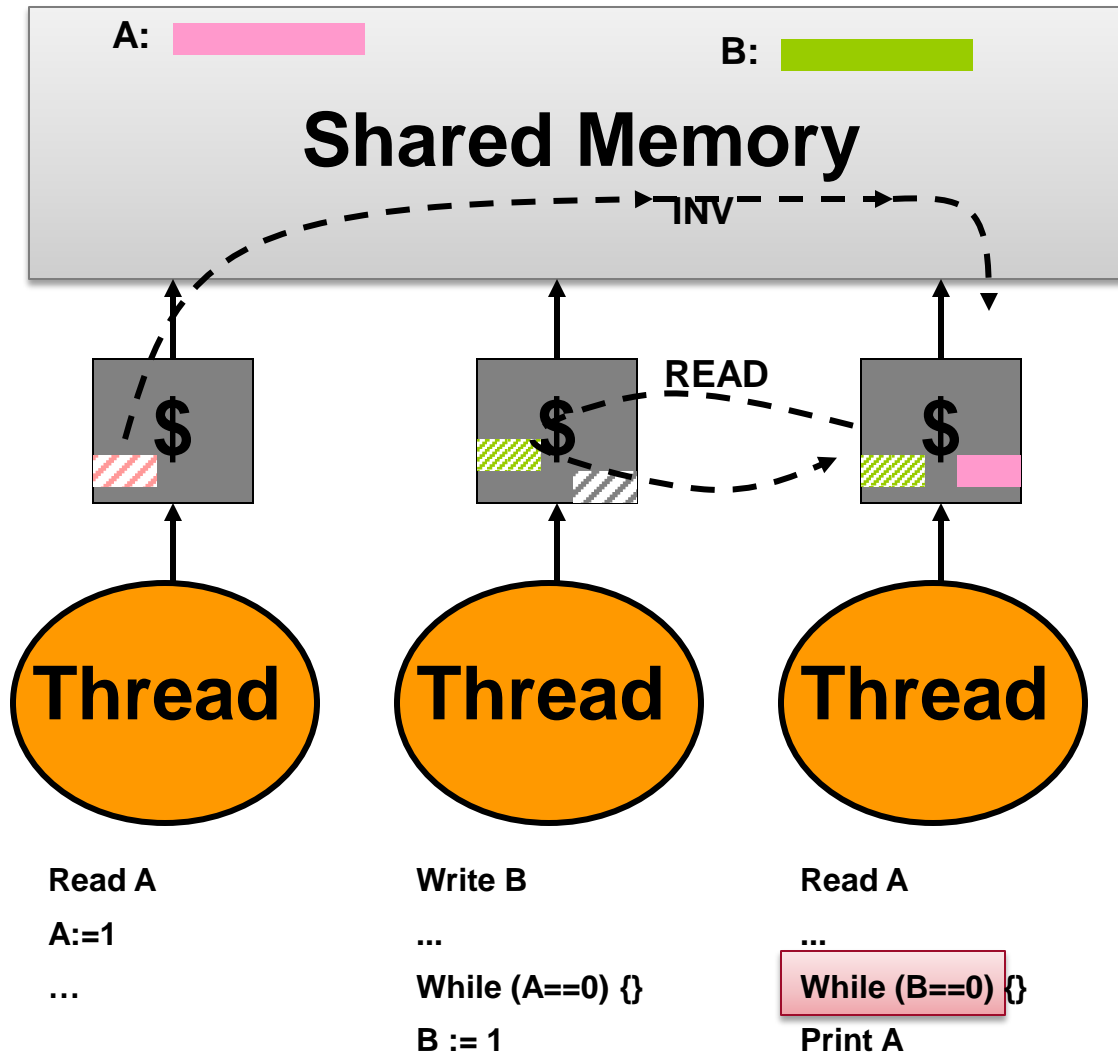




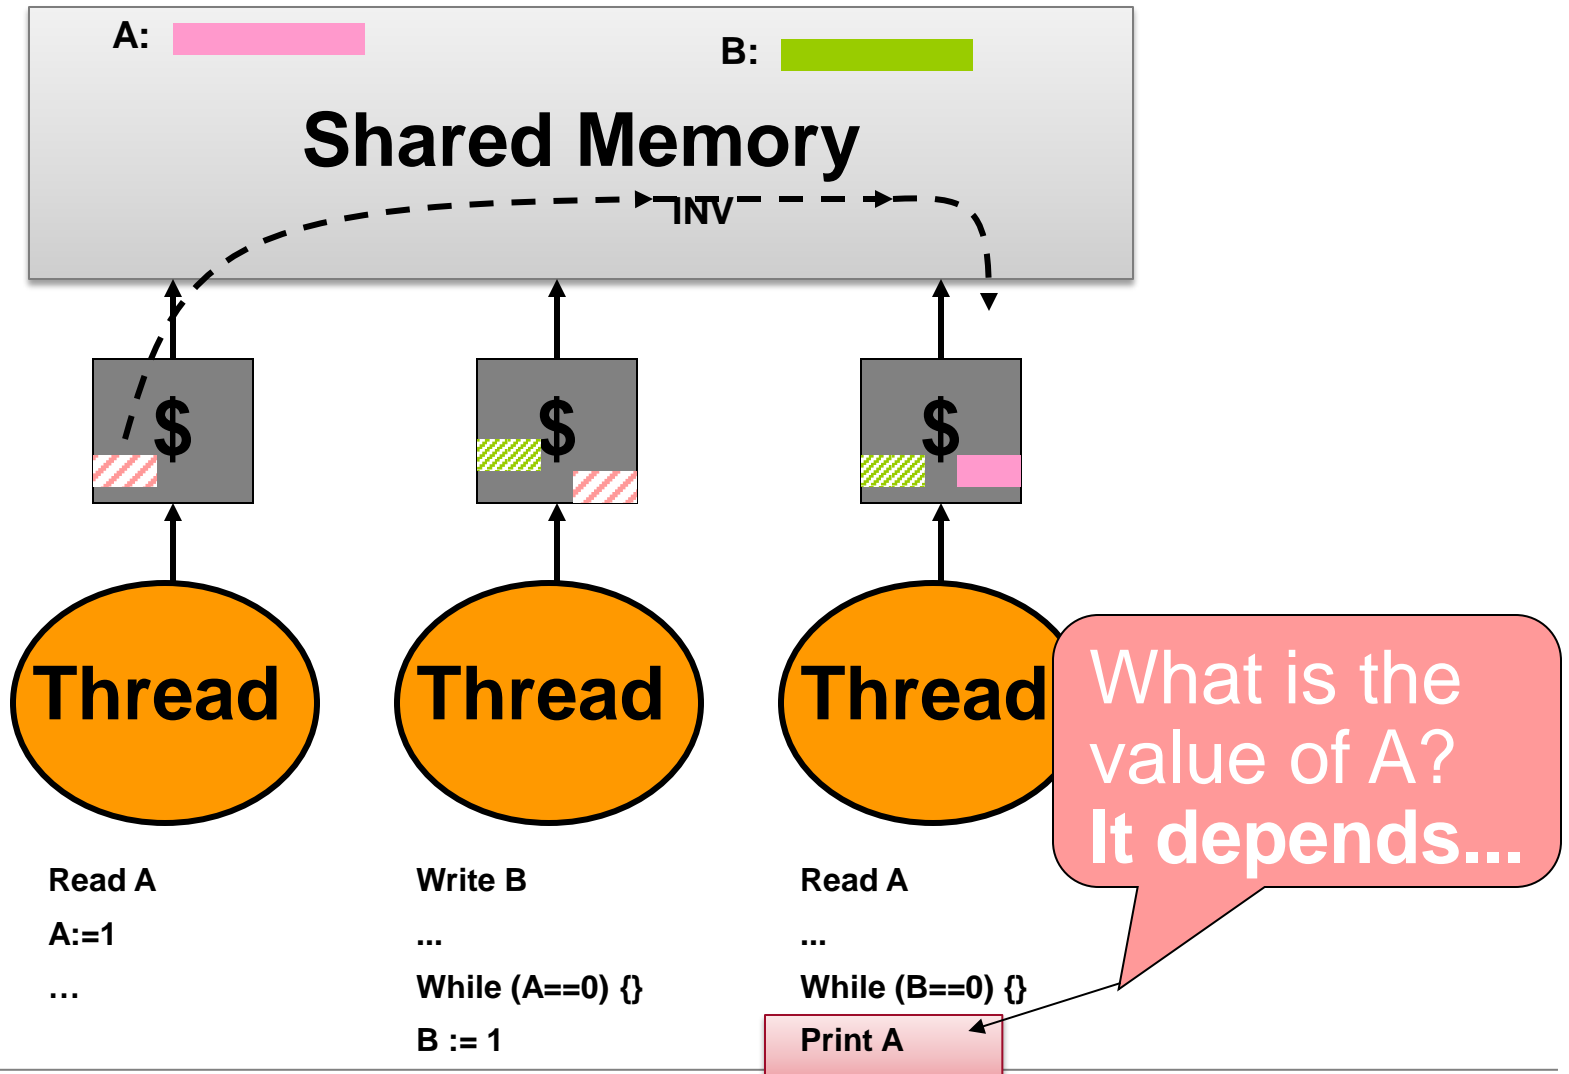
# Example1: Causal Correctness Issues



# Example1: Causal Correctness Issues



# Example1: Causal Correctness Issues




# Memory [Consistency] Model

- Specifies constraints on the order in which memory operations (from any process) can appear to execute with respect to one another
- Without it, can't tell much about a program's execution
- A contract between the programmer and the system designer
  - Programmer uses it to reason about correctness and possible results
  - System designer can use to constrain how much accesses can be reordered by compiler or hardware

# Sequential Consistency


- "A multiprocessor is sequentially consistent if the result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program." '
- [Lamport, 1979]

# Sequential Consistency



- Global order achieved by *interleaving* all memory accesses from different processes
- "Programmer's intuition is maintained"
  - Store causality? Yes
  - Does Dekker work? Yes
- -- Unnecessarily restrictive ==> performance penalty

# Weak/release Consistency



- Most accesses are unordered
- "Programmer's intuition is not maintained"
  - Store causality? No
  - Does Dekker work? No
- Global order only established when the programmer explicitly inserts memory barrier instructions
- ++ Great performance!!
- --- Interesting bugs!!
- But programming models such as OpenMP can help!