

ID1217 Programming with Processes  
Lecture 13


Concurrent Objects

Vladimir Vlassov  
KTH/ICT/EECS

# Outline

- Concurrent objects
  - Design space
  - Linearization
  - Patterns
  - Case study: Concurrent set based on ordered linked list
- Based on "*The Art of Multiprocessor Programming*" by Maurice Herlihy & Nir Shavit
  - Chapter 3 Concurrent objects
  - Chapter 9 Concurrent Lists: The Role of Locking
- *Acknowledgement*: slides are adapted from slides provided by the authors M. Herlihy and Nir Shavit

# Concurrent Computation with Shared Objects



# Shared Object Design Space (1/3)

- **Lock-based** vs **Lock-free** objects
  - It's about mutual exclusion synchronization on shared objects: whether and how accesses to shared objects are synchronized
  - **Lock-based (synchronized)**
    - *Monitors*: A lock per object; One caller (thread) at a time with mutual exclusion;
  - **Lock-free (unsynchronized)**
    - *Concurrent objects*: No locks; Allow concurrent access by multiple callers (threads)

## Shared Object Design Space (2/3)

- **Coarse-grained** vs **Fine-grained** synchronization
  - It's about the number of locks protecting a shared object and its components
  - **Coarse-grained synchronization**: one lock for the entire object (Monitor model)
  - **Fine-grained synchronization**: multiple locks: each object component has a lock
    - Split object into independently-synchronized components
    - E.g. each element of a set has its own lock.

# Shared Object Design Space (3/3)

- **Blocking** vs **Non-blocking (wait-free)** objects
  - It's about condition synchronization: await or not on a condition in an object
  - **Blocking**
    - Caller waits until state changes, e.g. a dequeuer awaits until a queue is not empty;
    - Spinning (busy waiting) vs blocking (using wait/signal)
  - **Non-Blocking**
    - No wait: Method throws exception, e.g. EmptyException
    - Methods return null
  - **Wait-free**
    - Every method call completes in a finite number of steps
    - Implies no mutual exclusion

# Objectivism (1/2)

- What is a **concurrent object**?
  - How do we **describe** one?
  - How do we **implement** one?
  - How do we **tell if we're right**?
    - Safety (correct behavior) and liveness (progress) properties
- First we focus on
  - how to describe, and
  - how to reason about safety (correctness) and liveness (progress) properties of an implementation algorithm.

## Objectivism (2/2)


- Each object has a **state**
  - Usually given by a set of **fields**
  - Queue example: sequence of items
- Each object has a set of **methods**
  - Only way to manipulate (inspect, alter) state
  - Queue example: **enq** and **deq** methods
- Methods take time: **A method call is a time interval**
- Concurrent method calls may overlap, may be nested
  - How do we describe and reason about concurrent objects?
  - How do we know that the implementation is correct and threads make progress?




# FIFO Queue Example

- Queue
  - pool of items
  - **enq()** & **deq()** methods
  - First-In-First-Out (FIFO) order

# FIFO Queue: Enqueue Method



# FIFO Queue: Dequeue Method



# Consider Two Implementations of a Bounded Queue

1. Lock-based (synchronized) bounded FIFO queue
  2. Lock-free (wait-free) bounded FIFO queue
- How to check whether implementations are correct?

## Lock-based Bounded FIFO Queue


- Lock-based (synchronized)
  - The Monitor model: Lock to enqueue/dequeue; unlock when done or on exception
  - Modifications are mutually exclusive
- Bounded
  - Fixed capacity
  - Good when resources are an issue
- Non-blocking (no busy-waiting)
  - **enq()** throws **FullException** if the queue is full
  - **deq()** throws **EmptyException** if the queue is empty
  - Versus methods that wait ...

# Lock-based Bounded Queue

```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```

# Lock-based Bounded Queue

```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```



Queue fields  
protected by single  
shared lock

# Lock-based Bounded Queue

```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```

Initially head = tail



## Implementation: Deq

```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```

## Implementation: Deq

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Method calls  
mutually exclusive

## Implementation: Deq

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

If queue empty  
throw exception

## Implementation: Deq

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Queue not empty:  
remove item and update  
head

## Implementation: Deq

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Return result

## Implementation: Deq

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Release lock no matter  
what!

## Implementation: Deq

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Should be correct (i.e. FIFO)  
because modifications are  
mutually exclusive, i.e. one at  
a time

# Lock-free Wait-free Bounded FIFO Queue

- Lock-free (unsynchronized)
  - The same thing without mutual exclusion
- For simplicity, only **two threads**
  - One thread **enq only**
  - The other **deq only**



# Wait-free 2-Thread Queue

```
public class waitFreeQueue {  
  
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];  
  
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }  
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }  
}
```

# Wait-free 2-Thread Queue


```
public class waitFreeQueue {
```

```
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];
```

```
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }
```

```
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }
```

```
}}
```



# Wait-free 2-Thread Queue

```
public class waitFreeQueue {
```

```
    int head = 0, tail = 0;  
    T[] items = (T[]) new Object[capacity];
```

```
    public void enq(Item x) {  
        if (tail - head == capacity) throw  
            new FullException();
```

```
        items[tail % capacity] = x; tail++;
```

```
    }  
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();
```

```
        Item item = items[head];  
        return item;
```

```
    }  
}
```

How do we define "correct"  
when modifications are not  
mutually exclusive?

Queue is updated without a lock!

# Defining Concurrent Queue Implementations

- We saw two types of queue implementations,
  - one that used locks and intuitively felt like a queue,
  - and one that did not use locks, it was in fact, wait-free.
- What properties are they actually providing, are they both implementations of a concurrent queue?
- In what ways are they similar and in what ways do they differ?

## Defining Concurrent Queue Implementations (cont' d)

- Need a way to **specify** a concurrent queue object
- Need a way to prove that an algorithm **implements** the object's specification, i.e. verify correctness
- Lets talk about object specifications ...

# Correctness and Progress

- In a concurrent setting, we need to specify both the **safety** and the **liveness** properties of an object
  - Safety: Properties that state that nothing bad ever happens
  - Liveness: Properties that state that something good eventually happens
- Need a way to define
  - when an implementation is **correct** (safety)
  - the conditions under which it guarantees **progress** (liveness)

**Lets begin with correctness**

# Sequential Objects

- Each object has a *state*
  - Usually given by a set of *fields*
  - Queue example: sequence of items
- Each object has a set of *methods*
  - Only way to manipulate state
  - Queue example: **enq** and **deq** methods

# Sequential Specifications

- If (precondition)
  - the object is in such-and-such a state
  - before you call the method,
- Then (postcondition)
  - the method will return a particular value
  - or throw a particular exception.
- and (postcondition, con't)
  - the object will be in some other state
  - when the method returns,



# Pre and PostConditions for Dequeue

- **Precondition:**
    - Queue is non-empty
  - **Postcondition:**
    - Returns first item in queue
  - **Postcondition:**
    - Removes first item in queue
- **Precondition:**
    - Queue is empty
  - **Postcondition:**
    - Throws Empty exception
  - **Postcondition:**
    - Queue state unchanged

# Why Sequential Specifications Totally Rock

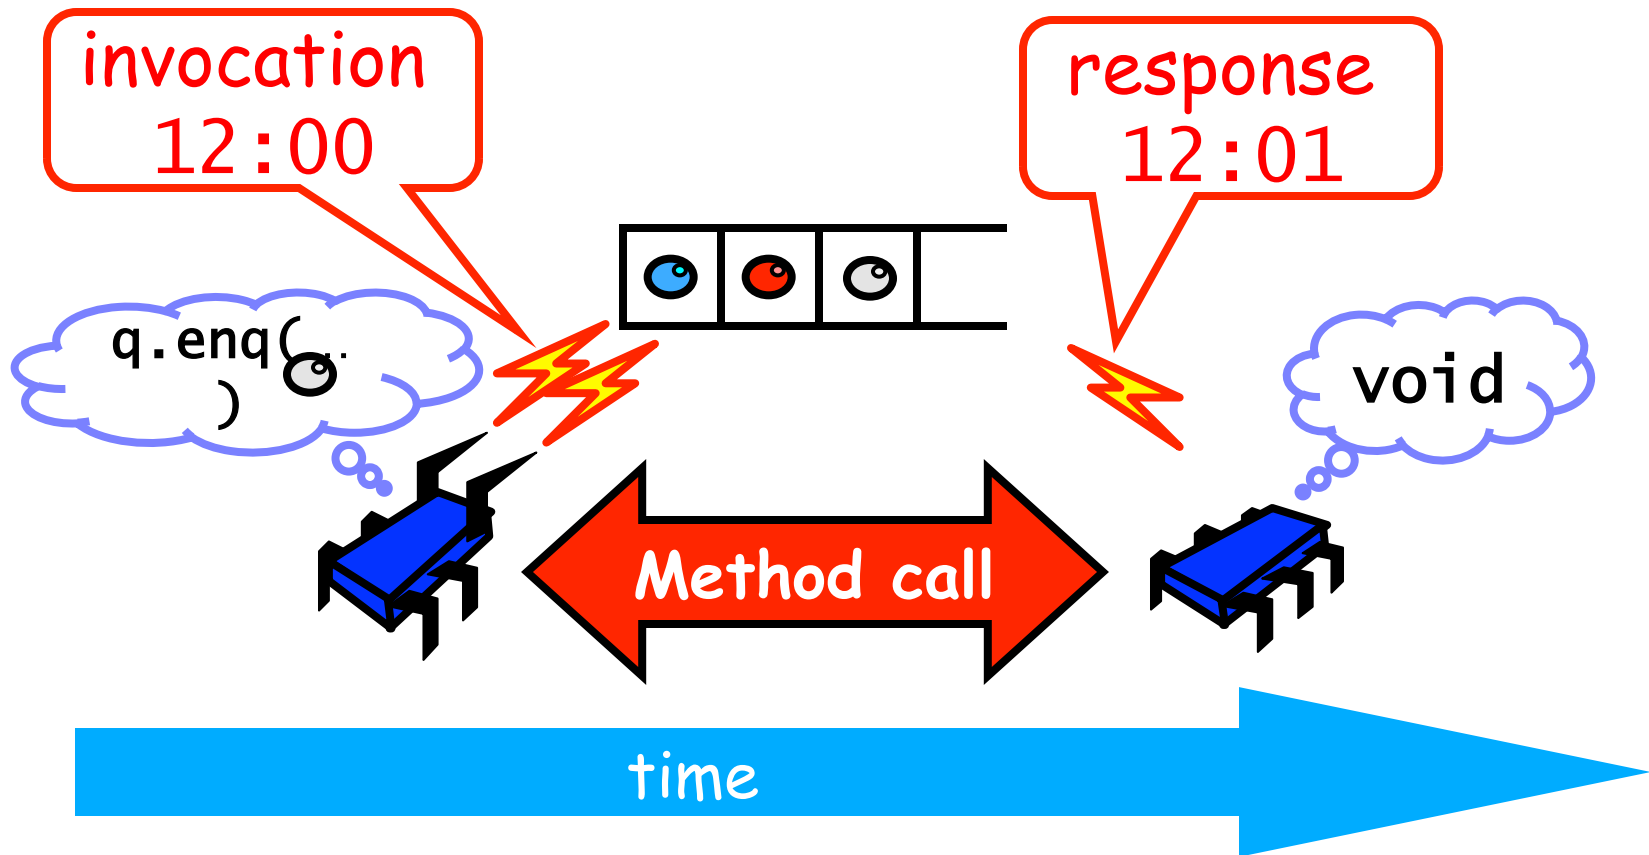
- Interactions among methods captured by side-effects on object state
  - State meaningful between method calls
  - E.g. after enq, an empty queue becomes non-empty; deq from the non-empty queue can make it empty
- Documentation size linear in number of methods
  - Each method described in isolation
- Can add new methods
  - Without changing descriptions of old methods

# What About Concurrent Specifications?

- Methods?
- Documentation?
- Adding new methods?

# Methods Take Time


In the sequential world, we treat method calls as if they were instantaneous, but in fact, they take time.



# Sequential vs Concurrent

- Sequential
  - Methods take time? Who knew?
- Concurrent
  - Method call is not an **event**
  - Method call is an **interval**.
    - Invocation and response events
    - Method takes effect in between the invocation and response events.

# Concurrent Methods Take Overlapping Time



# Sequential vs Concurrent

Sequential	Concurrent
<ul style="list-style-type: none"><li>Object needs meaningful state only <i>between</i> method calls</li></ul>	<ul style="list-style-type: none"><li>Because method calls overlap, object might <i>never</i> be between method calls</li></ul>
<ul style="list-style-type: none"><li>Each method described in isolation</li></ul>	<ul style="list-style-type: none"><li>Must characterize <i>all</i> possible interactions with concurrent calls<ul style="list-style-type: none"><li>– What if two <b>enqs</b> overlap?</li><li>– Two <b>deqs</b>? <b>enq</b> and <b>deq</b>? ...</li></ul></li></ul>
<ul style="list-style-type: none"><li>Can add new methods without affecting older methods</li></ul>	<ul style="list-style-type: none"><li>Everything can potentially interact with everything else</li></ul>

# Sequential vs Concurrent

Sequential	Concurrent
<ul style="list-style-type: none"><li>Object needs meaningful state only <i>between</i> method calls</li></ul>	<ul style="list-style-type: none"><li>Because method calls overlap, object might <i>never</i> be between method calls</li></ul>
<ul style="list-style-type: none"><li>Each method described in isolation</li></ul>	<ul style="list-style-type: none"><li>Must characterize <i>all</i> possible interactions with concurrent calls<ul style="list-style-type: none"><li>– What if two <b>enqs</b> overlap?</li><li>– Two <b>deqs</b>? <b>enq</b> and <b>d</b> <b>deq</b>? ...</li></ul></li></ul>
<ul style="list-style-type: none"><li>Can add new methods without affecting older methods</li></ul>	<ul style="list-style-type: none"><li>Everything can potentially interact with everything else</li></ul>

**Panic!**



# The Big Question

- What does it **mean** for a *concurrent* object to be correct?
  - What *is* a concurrent FIFO queue?
  - FIFO means *strict temporal order*
  - Concurrent means *ambiguous temporal order*

## Intuitively...

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

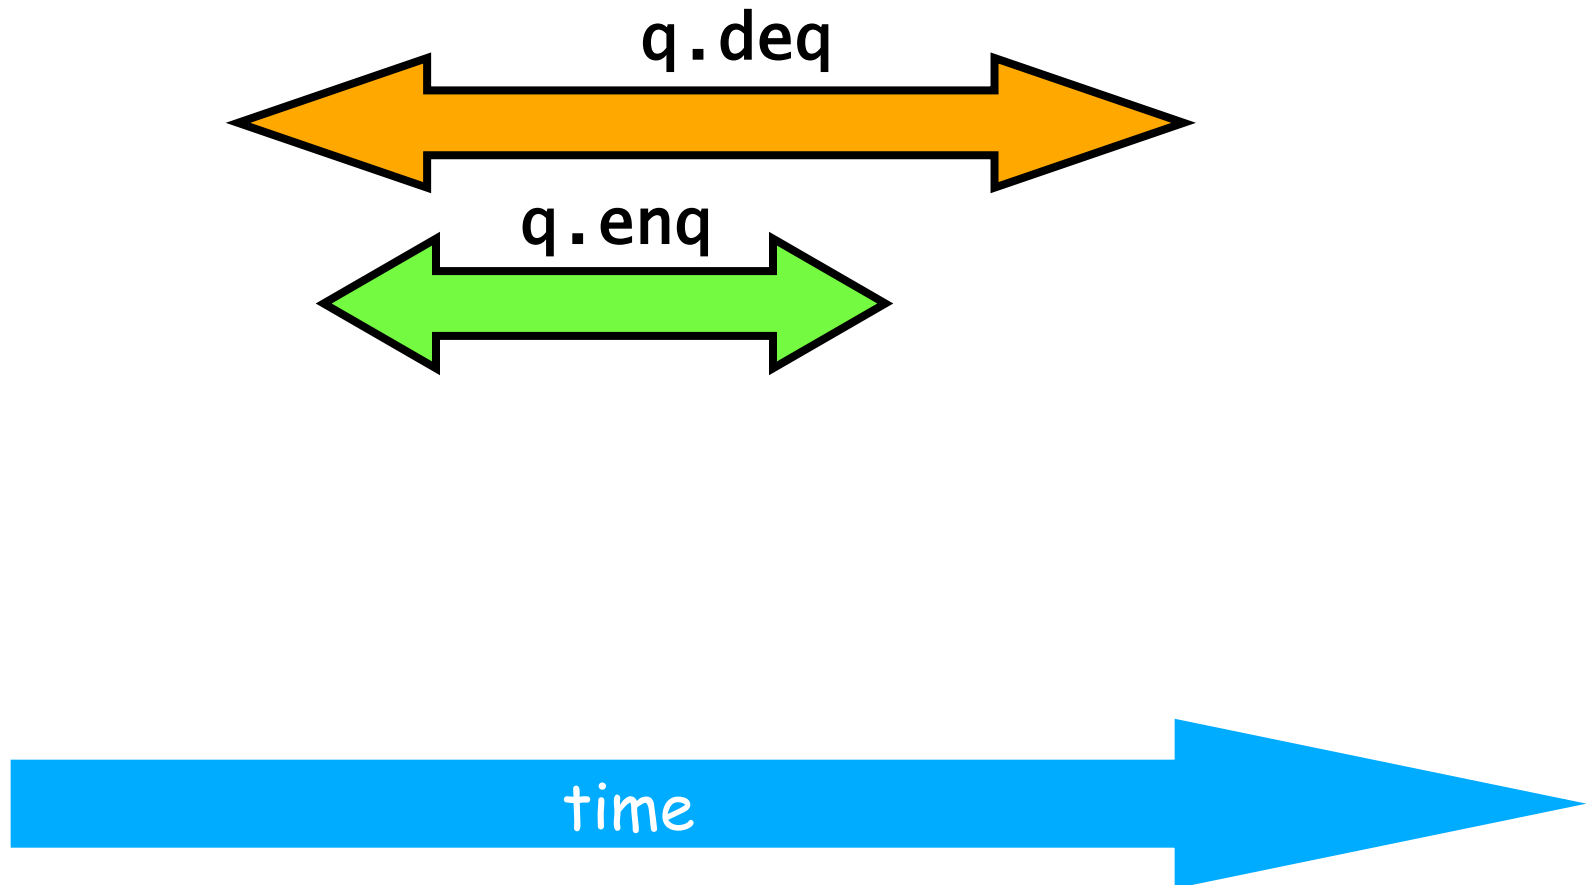
## Intuitively...

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```


All modifications  
of queue are done  
mutually exclusive

# Intuitively

Assume method calls overlap.



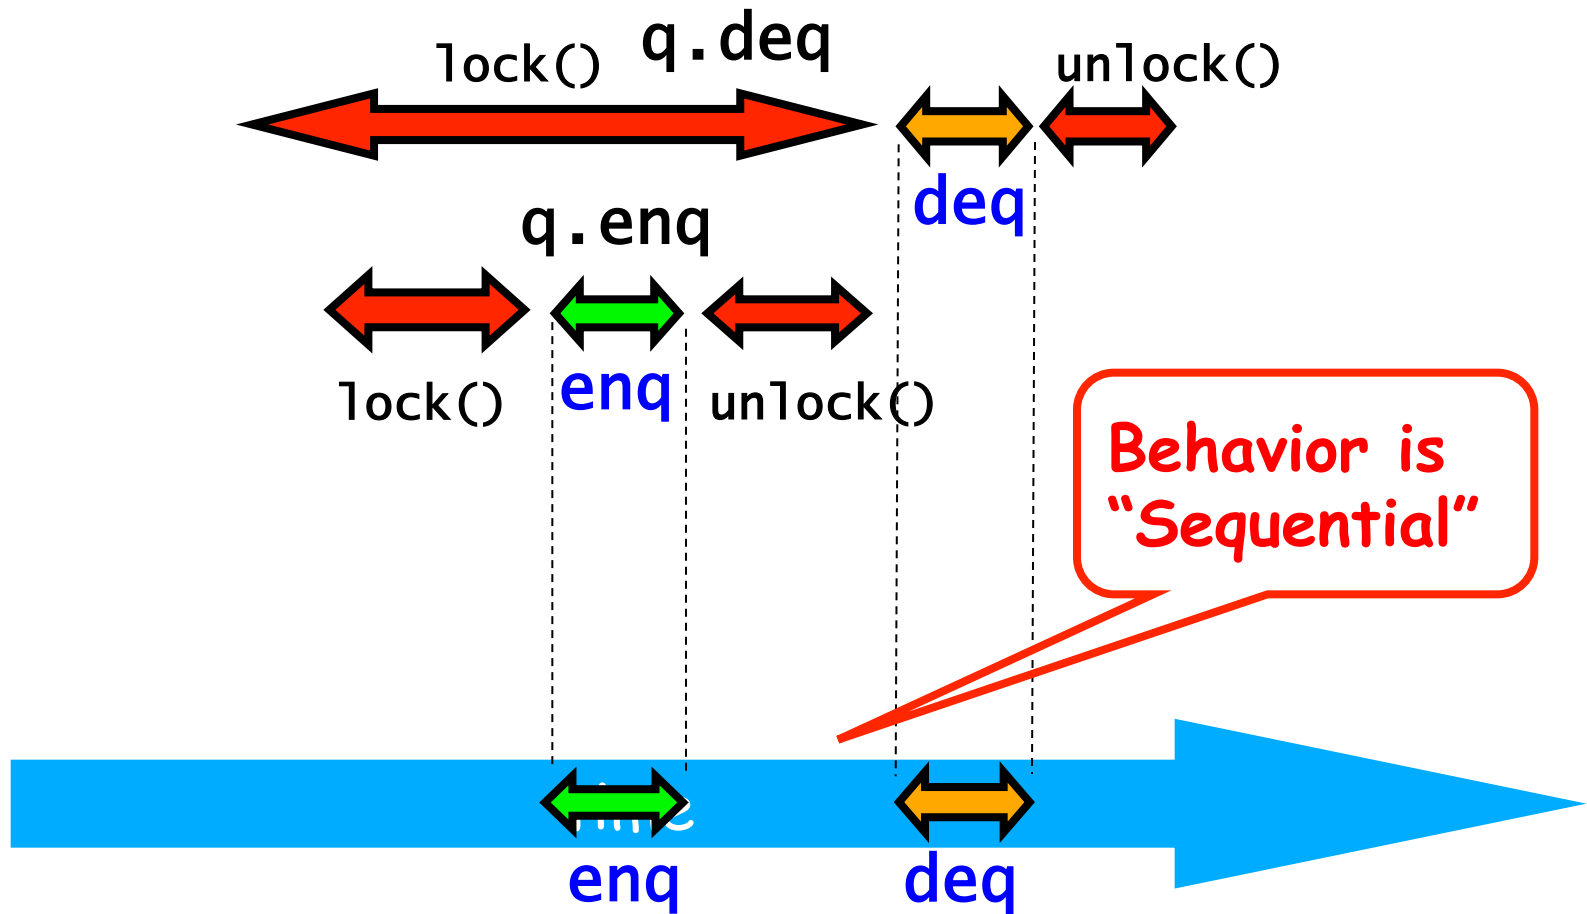
# Intuitively



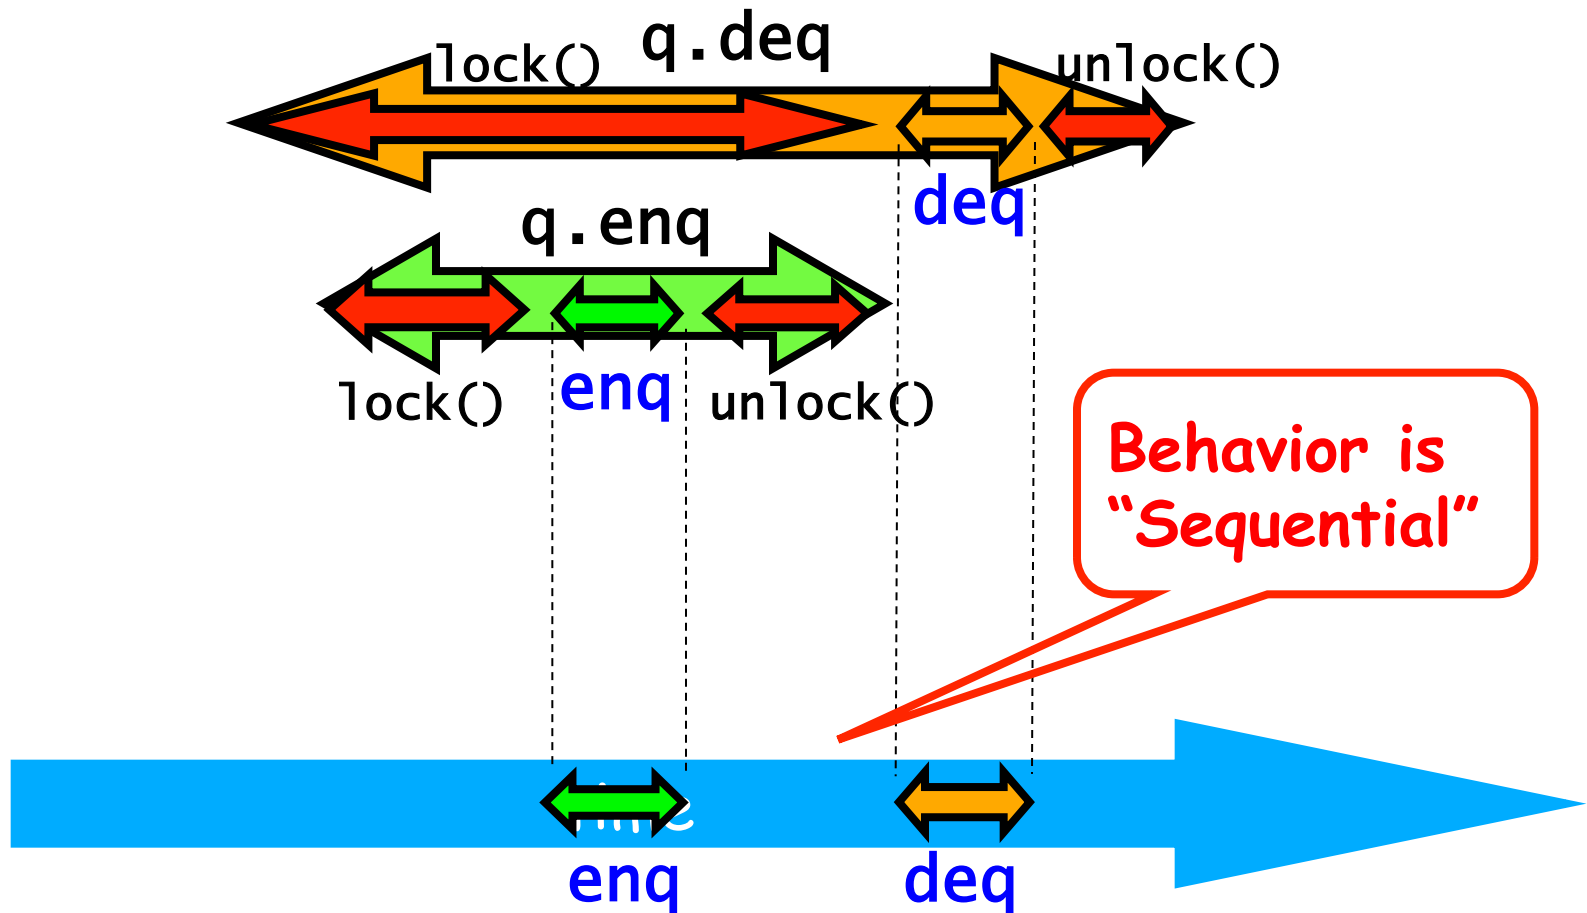
Since we use mutual exclusion locks, then each of the actual updates happens in a non-overlapping interval so the behavior as a whole looks sequential.



# Intuitively



Lets capture the idea of describing  
the concurrent via the sequential



# Linearizability (1/4)

- Each method should
  - **“take effect”** (changes become visible to other threads) **instantaneously between invocation and response events** (within the invocation interval)
  - **Linearization points** at which the method “takes effect”
  - A method may have several linearization points
- A concurrent object execution can be viewed as an **interleaving of linearization points** of method invocations on the object by multiple threads



## Linearizability (2/4)

- **Every concurrent execution history is equivalent to some sequential execution history** (which can be either legal or not)
  - If one method call precedes another, then the earlier call must have taken effect before the later call.
  - By contrast, if two method calls overlap, then their order is ambiguous, and we are free to order them in any convenient way.
- **A history is legal (correct) if it conforms to the sequential specification for the object**
- **We check the linearized execution (the sequential history) for correctness**

## Linearizability (3/4)

- Map the concurrent execution to the sequential one and check the latter for correctness
- **Object is correct if its “sequential” behavior (i.e. its linearized execution) is correct**
- Any such concurrent object is **Linearizable<sup>TM</sup>**
- **Linearizability is a correctness condition for concurrent objects.**

## Linearizability (4/4)

- Is it really about the object?
  - Each method should
  - “take effect”
  - instantaneously
  - between invocation and response events.
- Sounds like a property of **an execution...**
- **A linearizable object is an object all of whose possible executions are linearizable**

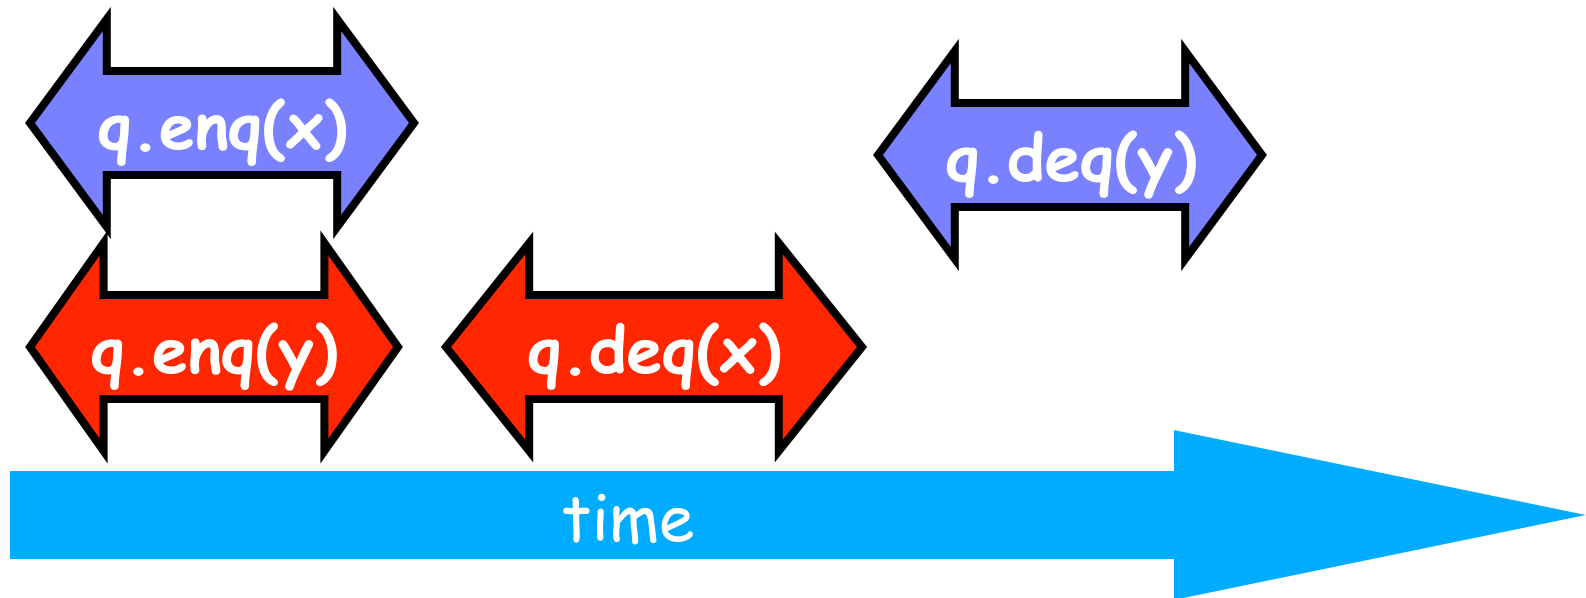

# Critical Sections

- Easy way to implement linearizability
  - Take sequential object
  - Make each method a critical section
- Problems
  - Blocking
  - No concurrency

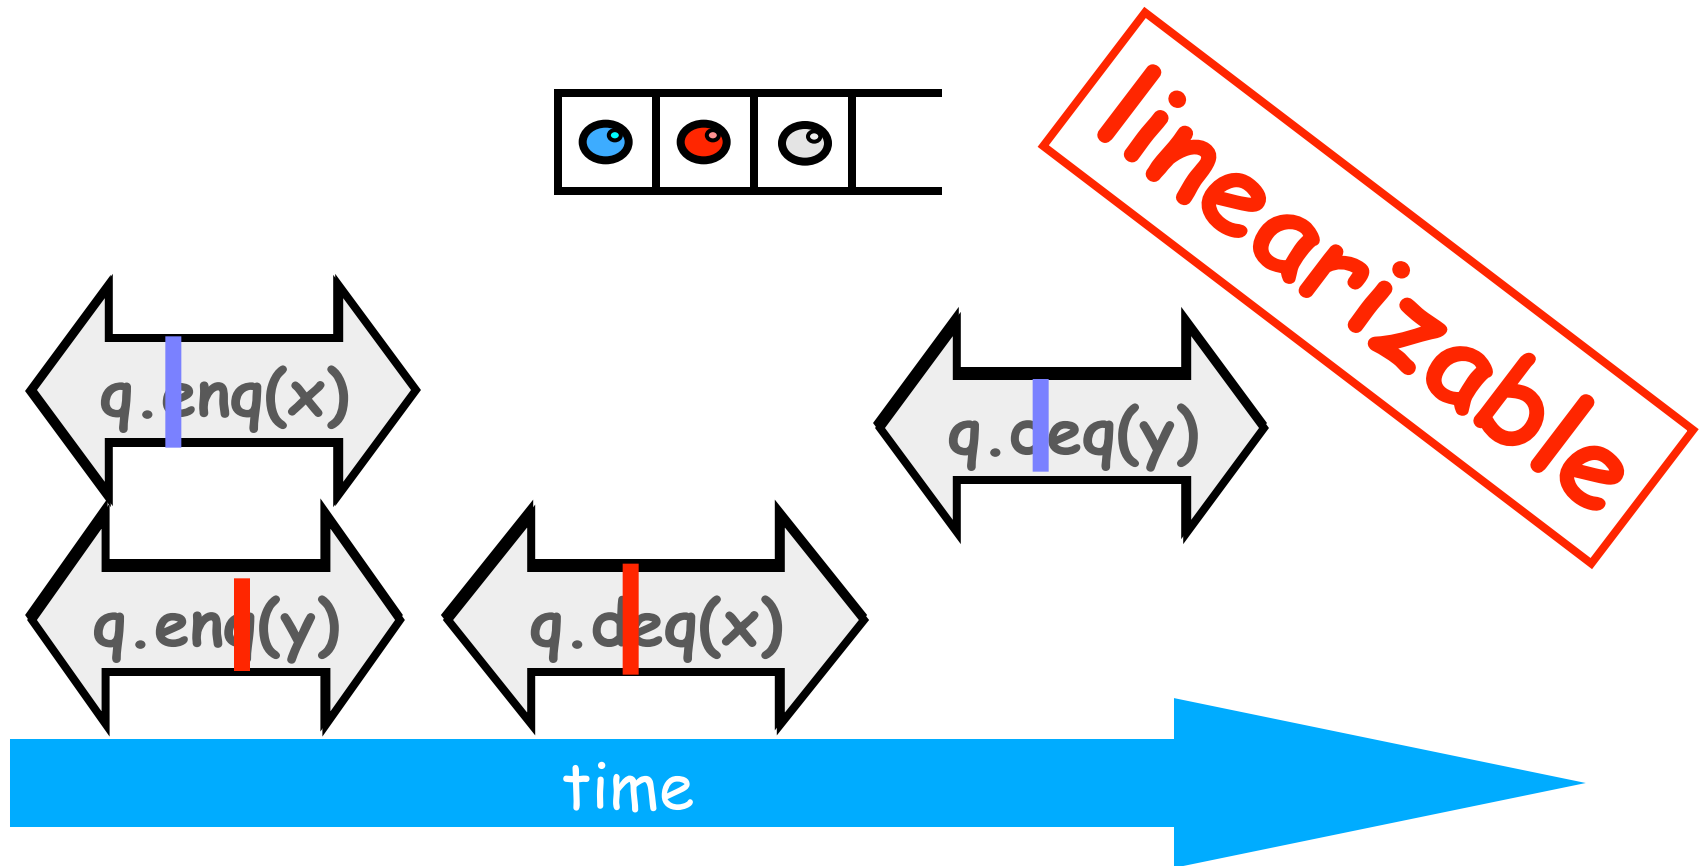


## Example 1

- Assume the following execution history.
- Is it linearizable?

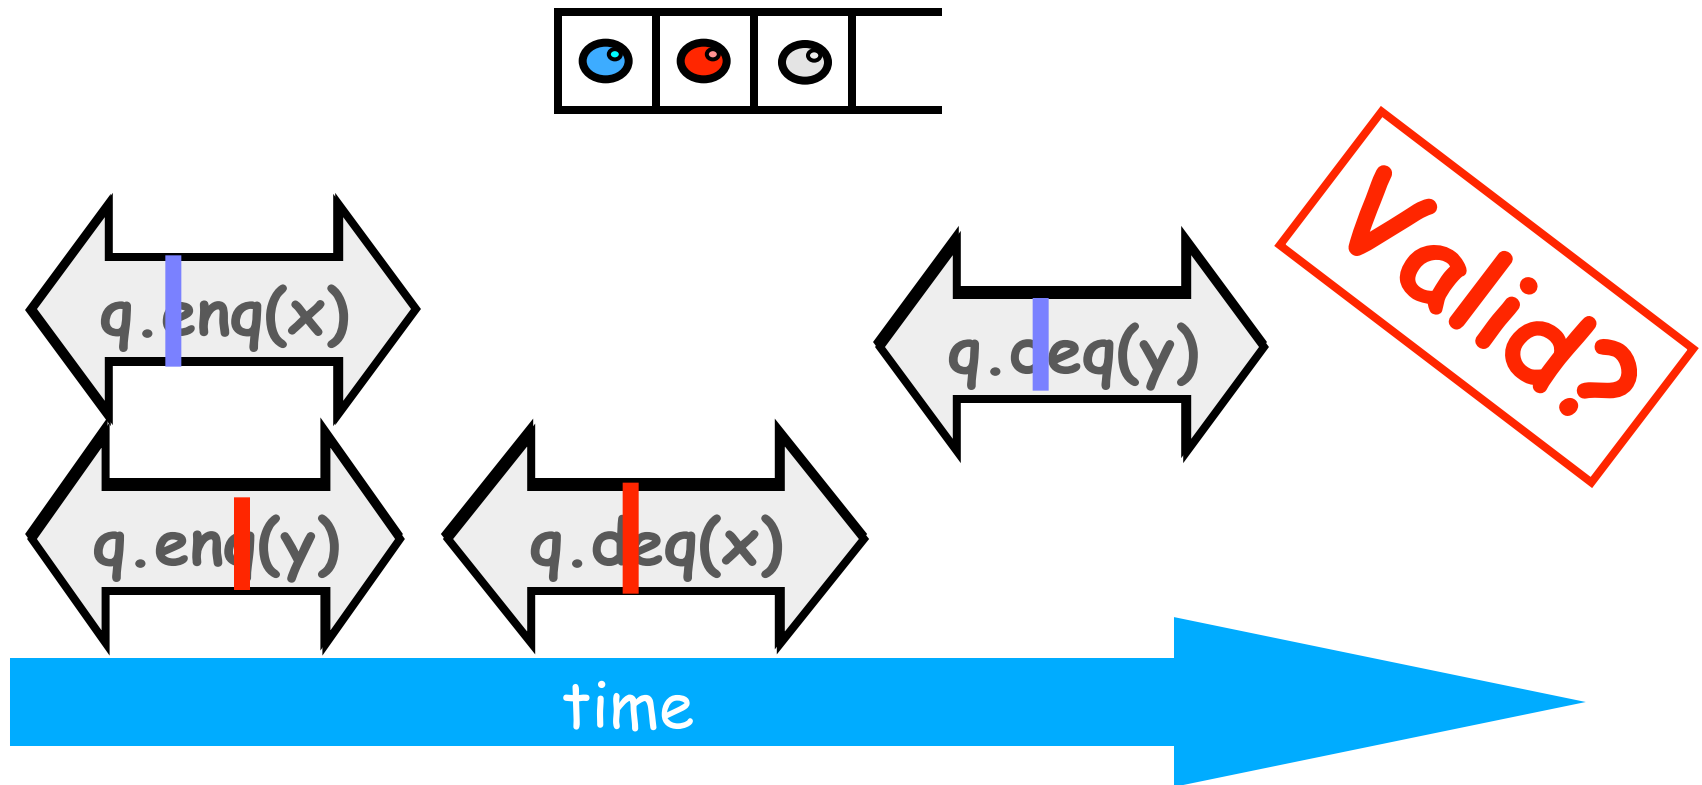


# Example 1



# Example 1



- What if we choose other enq linearization points?





## Example 2

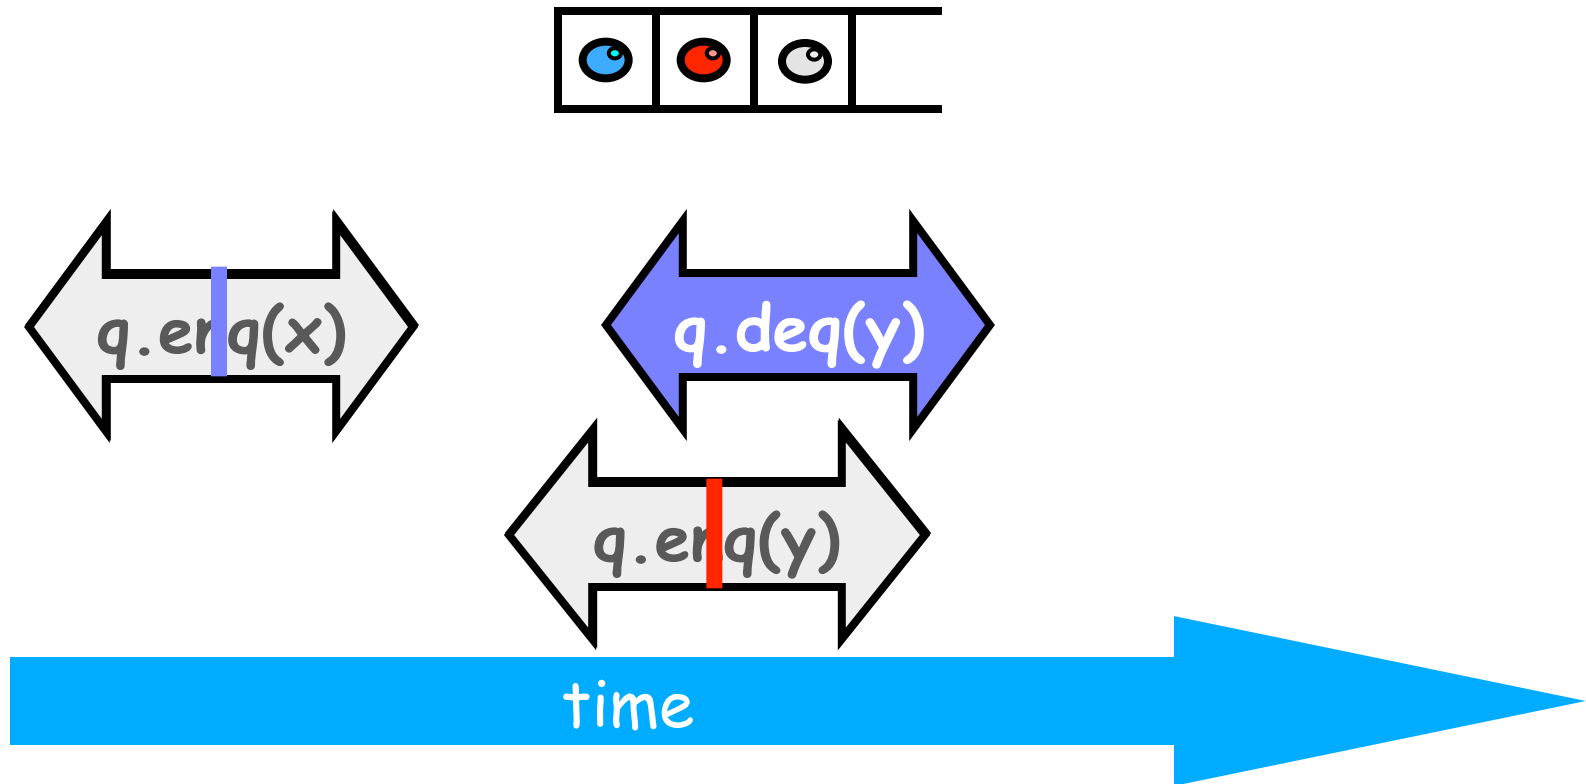
- Assume the following execution history.
- Is it linearizable?








## Example 2





## Example 2

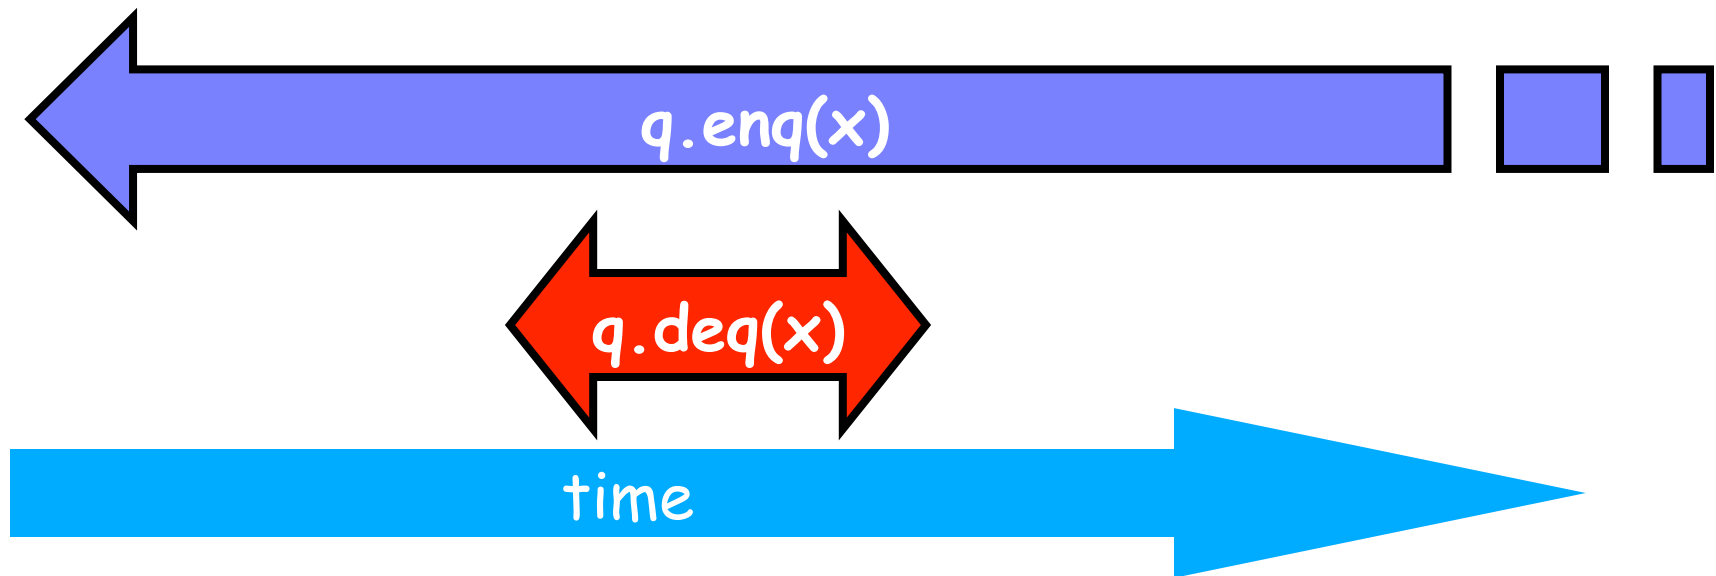

**not linearizable**





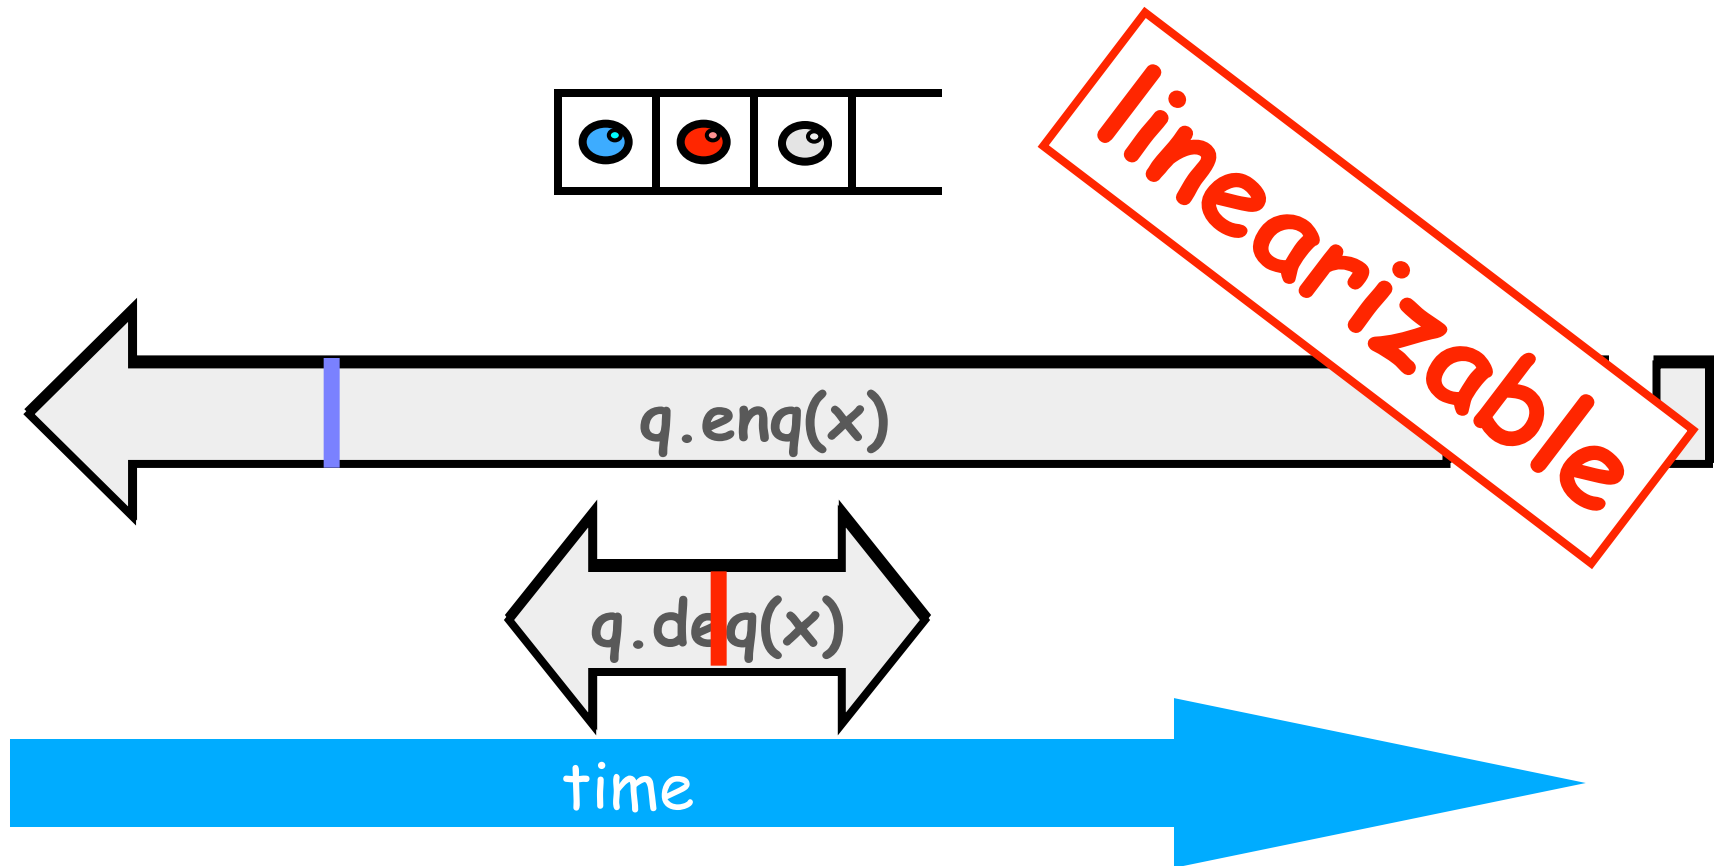
## Example 3

- Is this execution history linearizable?





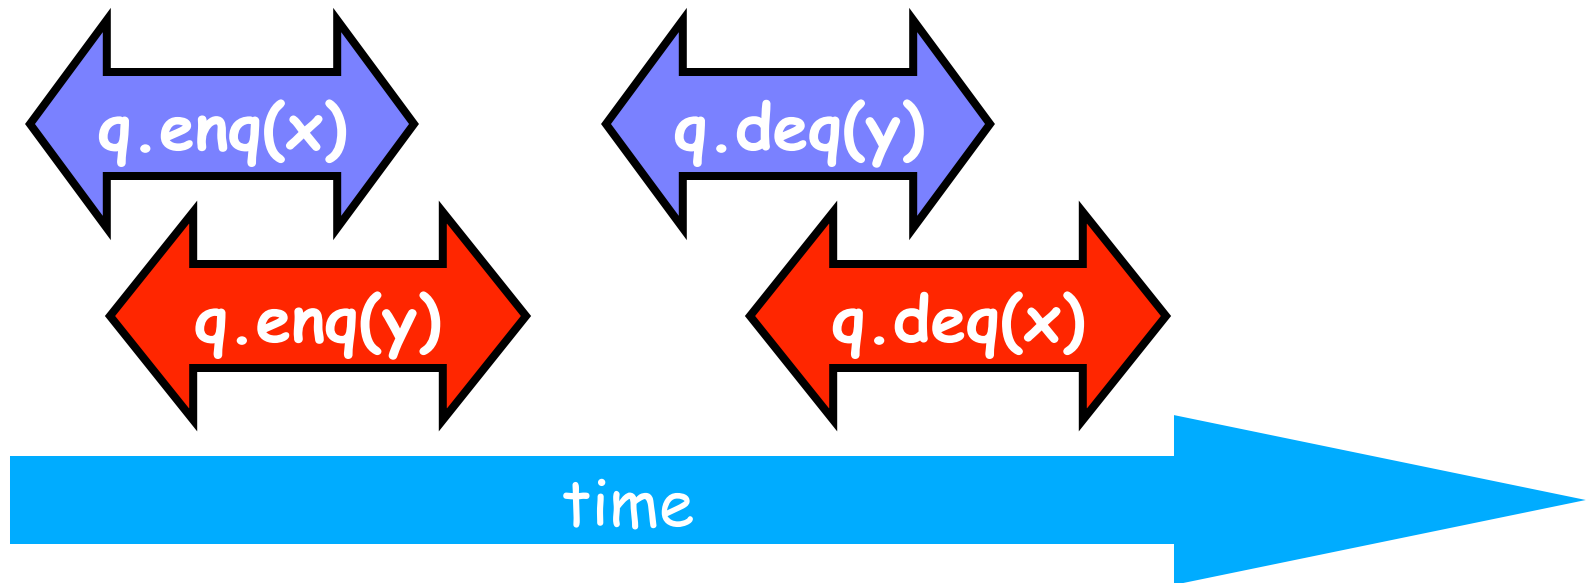

## Example 3





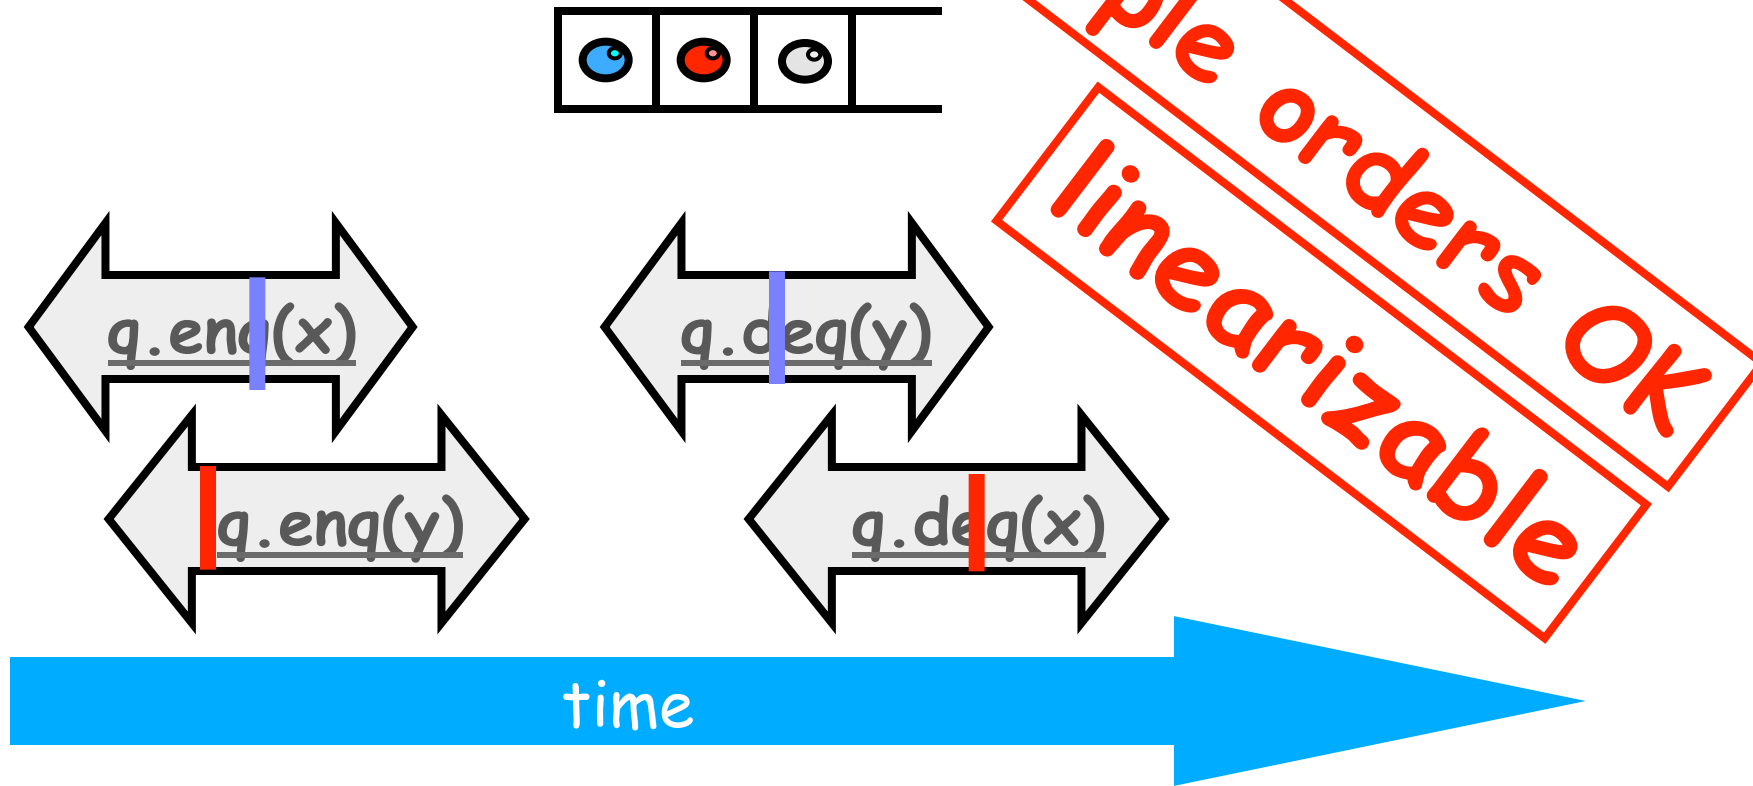
## Example 4

- Is this history linearizable?




Comme ci  
Comme ça

Example 4



# Reasoning About Linearizability: Locking (Lock-based Queue Example)

```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```



# Reasoning About Linearizability: Locking

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Linearization points  
are when locks are  
released




# More Reasoning: Wait-free

```
public class waitFreeQueue {
```

```
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];
```

```
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }
```

```
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }  
}
```



## More Reasoning: Wait-free

```
public class WaitFreeQueue {
```

```
    int head, tail = 0;  
    Item[] items = new Item[capacity];
```

Linearization order is  
order head and tail  
fields modified

Remember that there  
is only one enqueuer  
and only one dequeuer

```
    void enq(Item x) {  
        if (tail == capacity) throw  
            new FullException();  
        items[tail % capacity] = x;  
    }
```

tail++;

```
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity];  
        return item;
```

head++;

```
}}
```

# Talking About Executions

- Why?
  - Can't we specify the linearization point of each operation without describing an execution?
- Not Always
  - In some cases, linearization point **depends on the execution**

## Strategy

- Identify one atomic step where method “happens”
  - Critical section
  - Machine instruction
- Doesn't always work
  - Might need to define several different steps for a given method

# Linearizability: Summary

- Linearization points: Operation takes effect instantaneously between invocation and response
- Powerful specification tool for shared objects
- Allows us to capture the notion of objects being “atomic”
- **Helps to describe, develop and use a concurrent object, and reason about its properties using a sequential specification**
- **Linearization is a local property**
- **Locality implies composability: a composition of linearizable objects is linearizable**
- Good for high level objects

# Why Does Composability Matter?

- Modularity
- Can prove linearizability of objects in isolation
- Can compose independently-implemented objects

# Correctness via Linearizability

- Reason about correctness using *linearizability*
  - Find linearization points
  - Linearize the concurrent execution to the sequential one and check the latter for correctness
- **Linearizability is a correctness condition for concurrent objects.**
- **Object is correct if its “sequential” behavior (i.e. linearized execution) is correct**

## Progress

- There might be implementation whose methods are lock-based (deadlock-free)
- There might be implementation whose methods did not use locks (lock-free)
- How do they relate?



# Progress Conditions

- Informal definitions
- ***Deadlock-free***: some thread trying to acquire the lock eventually succeeds.
- ***Starvation-free***: every thread trying to acquire the lock eventually succeeds.
- ***Lock-free***: some thread calling a method eventually returns.
- ***Wait-free***: every thread calling a method eventually returns.

# Progress Conditions

	<b>Non-Blocking</b>	<b>Blocking</b>
<b>Everyone makes progress</b>	<b>Wait-free</b>	<b>Starvation-free</b>
<b>Someone makes progress</b>	<b>Lock-free</b>	<b>Deadlock-free</b>

# Concurrent Object Case Study: A Concurrent Set based on Linked List

# Concurrent Objects

- Shared objects are potential sequential bottlenecks
- Scalable concurrent data structures: how to?
- Adding threads should not lower throughput of object methods
  - Contention effects
  - Mostly fixed by using Queue locks
- Should increase throughput
  - Not possible if inherently sequential
  - Surprising things are parallelizable

# Concurrent Objects: Design Space

- **Lock-based** (synchronized) vs **Lock-free** (non-synchronized)
  - It's about mutual exclusion
- **Blocking** vs **Non-Blocking**
  - It's about condition synchronization
- **Wait-Free (lock-free and non-blocking)**
  - Each call takes a finite number of steps
- **Fine-grained** vs **coarse-grained** synchronization
  - It's about granularity of synchronized objects
  - Fine-grained vs coarse-grained locking

# Coarse-Grained Lock-Based Synchronization

- Synchronized class: Take a sequential implementation of a class, add a scalable lock field, and ensure that each method call acquires and releases that lock.
- Each method locks the object
  - Avoid contention using queue locks
  - Easy to reason about
    - In simple cases
  - Standard Java model: Monitors
    - **Synchronized** blocks and methods
- So, are we done with concurrent objects?

# Coarse-Grained Lock-Based Synchronization

## (cont'd)

- Sequential bottleneck
  - Threads “stand in line”
- Adding more threads
  - Does not improve throughput
  - Struggle to keep it from getting worse
- So why even use a multiprocessor?
  - Well, some apps inherently parallel ...

## “Traffic Jam”

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
  - Enters critical section
  - And “eats the big muffin” (stops running)
    - Cache miss, page fault, descheduled ...
  - Everyone else using that lock is stuck!
  - Need to trust the scheduler....



# Four Patterns for Concurrent Objects

1. Fine-Grained Synchronization
2. Optimistic Synchronization
3. Lazy Synchronization (not considered here)
4. Lock-Free Synchronization
  - For *highly-concurrent* objects
    - Concurrent access
    - More threads, more throughput
  - Course-Grained Synchronization (the Monitor model) was considered earlier

## First: Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
  - Independently-synchronized components
- Methods conflict when they access
  - The same component ...
  - At the same time

## Second:

# Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check (validate)...
  - OK: we are done
  - Oops: start over
- Evaluation
  - Usually cheaper than locking, but
  - Mistakes are expensive

## Third:

# Lazy Synchronization

- Rather object-specific
  - Might work on collections, e.g. list
- Postpone hard work, e.g. physical removals
- Removing components is tricky
  - *Logical* removal
    - Mark component to be deleted
  - *Physical* removal
    - Do what needs to be done

## Fourth:

# Lock-Free Synchronization

- Don't use locks at all
  - Use **CompareAndSet()** & relatives ...
  - Transactional style: load shared state, modify locally, then CAS: Compare And Set

```
do {  
    old = state.get();  
    new = modify(old);  
} until (state.compareAndSet(old, new))
```

- Advantages
  - No Scheduler Assumptions/Support
- Disadvantages
  - Complex
  - Sometimes high overhead

## Case Study: Concurrent Set Based on Linked List

- To illustrate some of the patterns...
- Using a **list**-based **Set**
  - Common application
  - Building block for other apps
- Consider three implementation patterns
  1. Course-grained locking (the Monitor model)
  2. Fine-grained locking
  3. Lock-free
  - For other patterns – optimistic and lazy synchronizations – see the book ***The Art of Multiprocessor Programming***, by Maurice Herlihy, and Nir Shavit

# The Set Interface

- Unordered collection of items
- No duplicates
- Methods
  - **add(x)** put **x** in set
  - **remove(x)** take **x** out of set
  - **contains(x)** tests if **x** in set

## List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```



## List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

**Add item to set**

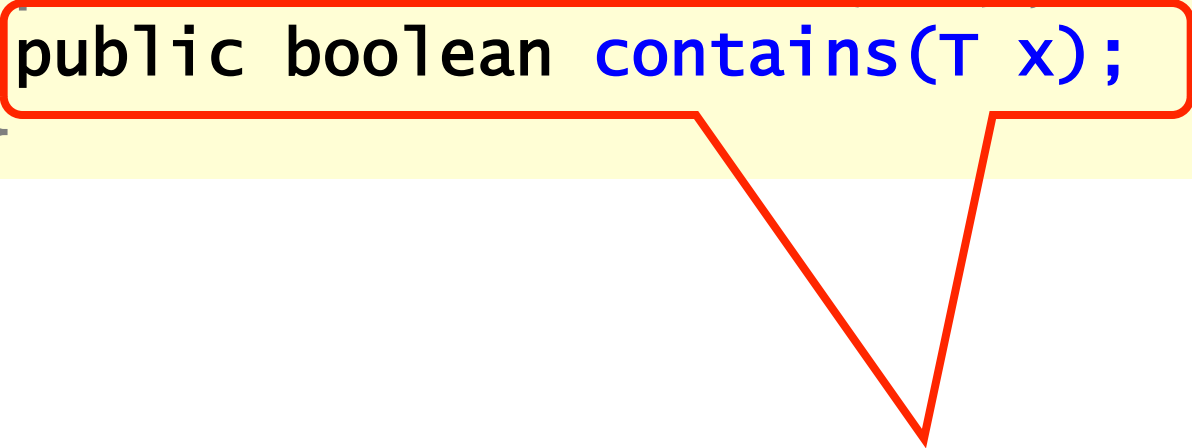
## List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

**Remove item from set**

## List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```



**Is item in set?**

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

**item of interest**

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```


Usually hash code

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

Reference to next node

## The List-Based Set



Sorted using keys with Sentinel nodes  
(min & max possible keys)



# Sequential List Based Set

Add()




Remove()




# Sequential List Based Set


Add()




Remove()



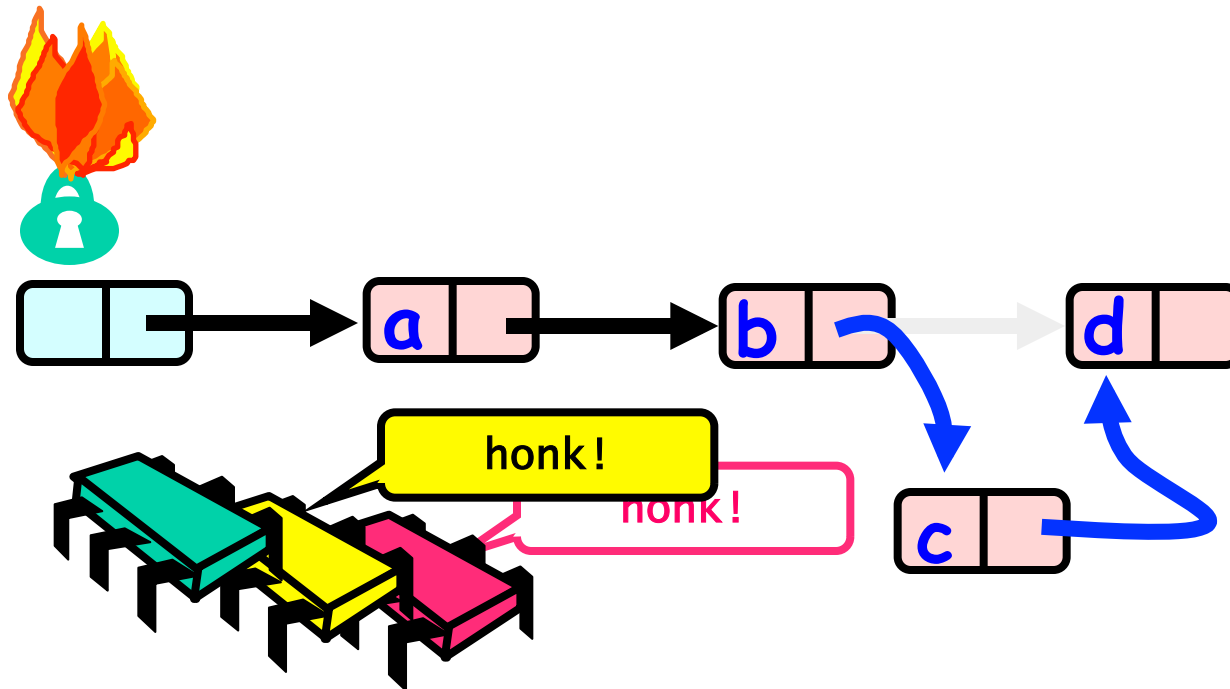
# Coarse-grained Locking: Coarse-Grained Synchronized Set



# Coarse-grained Locking



# Coarse-grained Locking



Simple but hotspot + bottleneck

# Coarse-grained Locking

- Easy, same as synchronized methods
  - **“One lock to rule them all ...”**
- Simple, clearly correct
  - Deserves respect!
- Works poorly with contention
  - Queue locks help
  - But bottleneck still an issue

# First Pattern: Fine-grained Locking



- Requires **careful** thought
  - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”
    - J.R.R. Tolkien, "The Fellowship of the Ring"
- Split object into pieces
  - Each piece has own lock
  - Methods that work on disjoint pieces need not exclude each other

## Fine-grained Locking Idea


- Improve concurrency by locking individual entries, rather than locking the entire list.
- Instead of a single lock on the entire list, add a lock to each entry.
- As a thread traverses the list, it locks each entry when it first visits, and some time later releases it.
- Such fine-grained locking allows concurrent threads to traverse the list together in a pipelined fashion.




# Hand-over-Hand Locking




# Hand-over-Hand Locking




# Hand-over-Hand Locking




# Hand-over-Hand Locking




# Hand-over-Hand Locking




# Removing a Node




remove(b)



# Removing a Node




# Removing a Node






# Removing a Node




# Removing a Node




## Concurrent Removes (each locks one node)

- Supposed we only locked one node and not two.
- Consider two concurrent threads that want to remove adjacent nodes.




# Concurrent Removes


- The red thread looks for the predecessor to c, while the green thread looks for the predecessor to b.



# Concurrent Removes




# Concurrent Removes




# Concurrent Removes

- The green thread looks for the predecessor to b.




# Concurrent Removes






# Concurrent Removes




## Uh, Oh

- But the final effect is to remove b but not c!



Uh, Oh



Bad news, C not removed




## Insight

- If a node is locked
  - No one can delete node's *successor*
- If a thread locks
  - Node to be deleted
  - And its predecessor
  - Then it works


## Concurrent Removes Again (each locks two nodes)




# Removing a Node



# Removing a Node




# Removing a Node







# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node



# Removing a Node




# Removing a Node







# Removing a Node




# Removing a Node



# Removing a Node



## Removing a Node

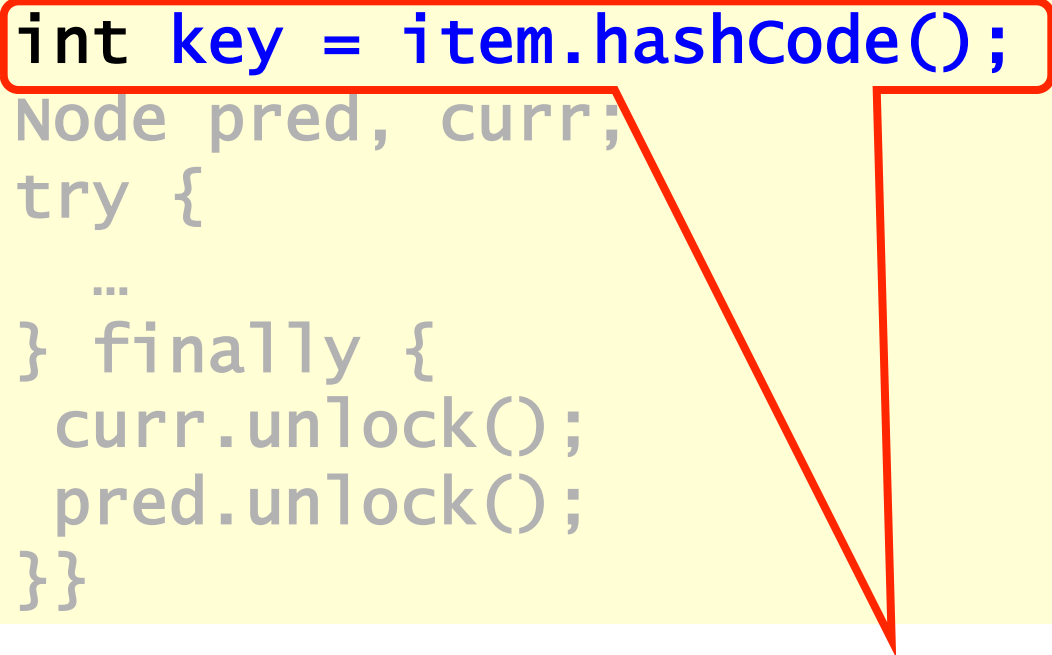


## Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

## Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```



**Key used to order node**

## Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        currNode.unlock();  
        predNode.unlock();  
    }  
}
```

**Predecessor and current nodes**

# Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;
```

```
    try {
```

```
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Make sure  
locks released**



## Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Everything else  
(on next slide)**


## Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```

## Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```


lock pred == head



## Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```


Lock current



# Remove method


```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```

Traversing list



## Remove: searching


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next; // remove  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



## Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```


Search key range



## Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**At start of each loop: curr  
and pred locked**






## Remove: searching

```
while (curr.key <= key) {
    if (item == curr.item) {
        pred.next = curr.next;
        return true;
    }
    pred.unlock();
    pred = curr;
    curr = curr.next;
    curr.lock();
}
```


## If item found, remove node



## Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```


**If node found, remove it**



## Remove: searching

## Unlock predecessor


```
while (curr.key <= key) {
    if (item == curr.item) {
        pred.next = curr.next;
        return true;
    }
    pred.unlock();
    pred = curr;
    curr = curr.next;
    curr.lock();
}
return false;
```



## Remove: searching

**Only one node locked!**


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



## Remove: searching


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

demote current



## Remove: searching


```
while (curr.key <= key) {  
    Find and lock new current  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = currNode;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



## Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = currNode;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

Lock invariant restored



## Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

**Otherwise, not present**



**return false;**



## Why does this work?

- To remove node e
  - Must lock e
  - Must lock e's predecessor
- Therefore, if you lock a node
  - It can't be removed
  - And neither can its successor

## Linearization points in remove()

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point if  
item is present**

## Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point  
if item not present  
(before return  
false)**

## Adding Nodes

- To add node e
  - Must lock predecessor
  - Must lock successor
- Neither can be deleted
  - (Is successor lock actually required?)
  - successor lock not actually required!

# Drawbacks of Fine-grained Locking

- Better than coarse-grained lock
  - Threads can traverse in parallel
- Still not ideal
  - Long chain of acquire/release
  - Inefficient

# Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
  - Enters critical section
  - And “eats the big muffin” (stops running)
    - Cache miss, page fault, descheduled ...
  - Everyone else using that lock is stuck!
  - Need to trust the scheduler....

## Fourth Pattern: Lock-Free

- Reminder: Lock-Free Data Structures
- No matter what ...
  - Guarantees minimal progress in any execution
  - i.e. Some thread will always complete a method call
  - Even if others halt at malicious times
  - Implies that implementation can't use locks




## A Lock-free List (Set)

- Next logical step
  - Wait-free **contains()**
    - Search by traversing the list from head to tail without locking
  - Lock-free **add()** and **remove()**
    - Search by traversing the list from head to tail without locking
    - Use CAS to complete the operation
- Use only **CAS: compareAndSet()**
  - What could go wrong?



## Lock-free List




Use CAS to verify pointer is correct

# Not enough!

# Physical Removal CAS pointer

# Problem...

Logical Removal =  
Set Mark Bit

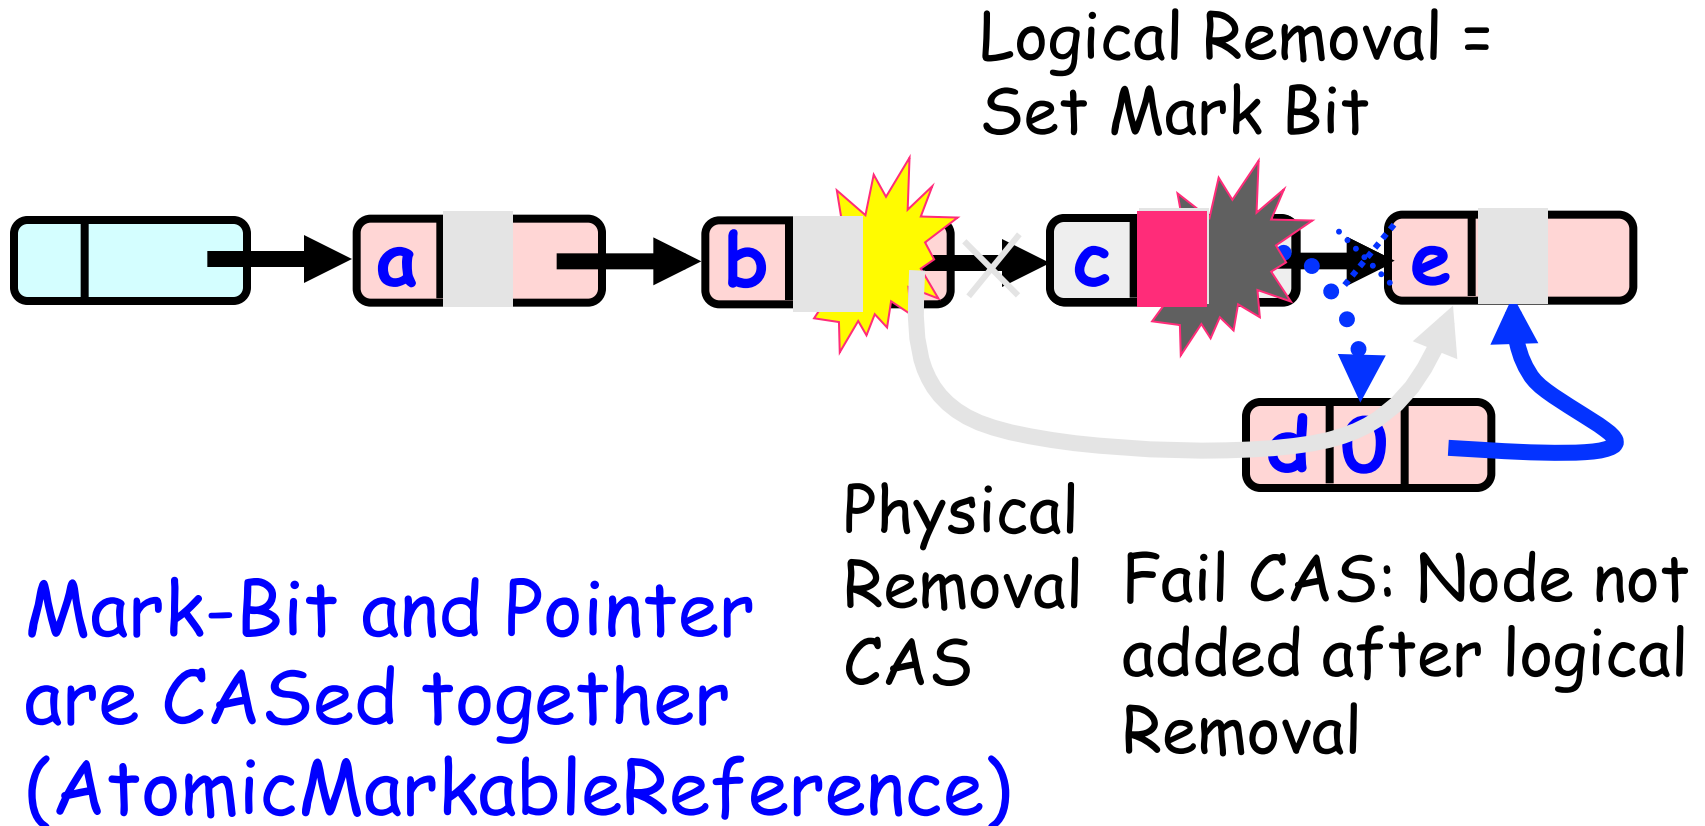


Problem:  
d not added to list...  
Must Prevent  
manipulation of  
removed node's pointer

Physical  
Removal  
CAS

Node added  
Before  
Physical  
Removal CAS

## The Solution: Combine Bit and Pointer




## Solution

- Use **AtomicMarkableReference**
- Atomically
  - Swing reference and
  - Update flag
- Remove in two steps
  - Set **mark** bit in **next** field
  - Redirect predecessor's pointer

# Marking a Node

- **AtomicMarkableReference** class
  - `java.util.concurrent.atomic` package



## Getting Reference (& Mark)

- Getting Reference & Mark

```
Public Object get(boolean[] marked);
```

```
Public Object get(boolean[] marked);
```

Returns reference

Returns mark at array index 0!

- Getting Mark only

```
public boolean isMarked();
```

Value of mark

# Changing State

```
Public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark);
```

# Changing State

If this is the current  
reference ...

```
Public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark);
```

And this is the  
current mark ...



# Changing State

...then change to this  
new reference ...

```
Public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark);
```

... and this new  
mark

## Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark);
```


## Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark);
```

If this is the current  
reference ...


## Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark);
```




.. then change to  
this new mark.


# Removing a Node





# Removing a Node



# Removing a Node



# Removing a Node






## Traversing the List

- **Q:** what do you do when you find a “logically” deleted node in your path?
- **A:** finish the job.
  - CAS the predecessor’s **next** field
  - Proceed (repeat as needed)

# Lock-Free Traversal (only Add and Remove)

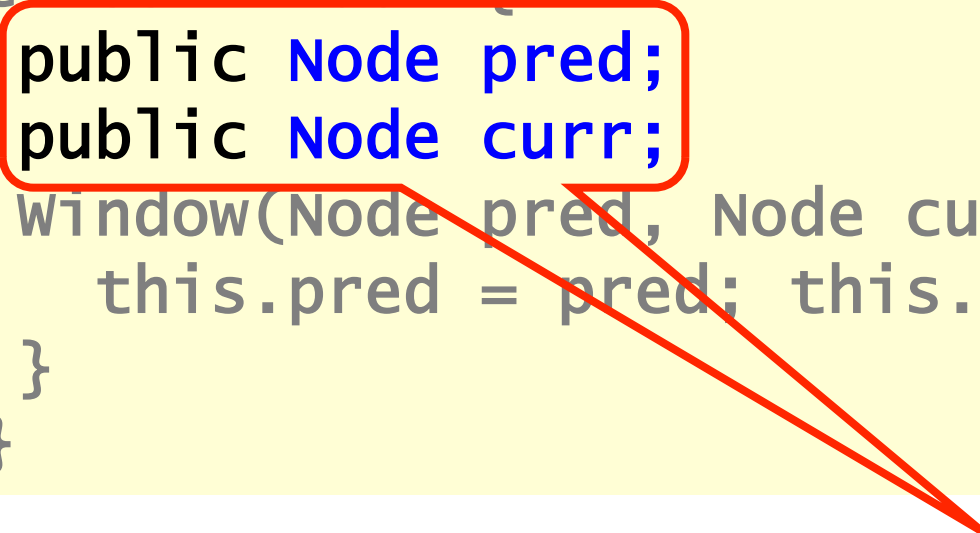


# The Window Class

```
class window {  
    public Node pred;  
    public Node curr;  
    window(Node pred, Node curr) {  
        this.pred = pred; this.curr = curr;  
    }  
}
```

# The Window Class

```
class window {  
    public Node pred;  
    public Node curr;  
    window(Node pred, Node curr) {  
        this.pred = pred; this.curr = curr;  
    }  
}
```



**A container for pred  
and current values**

## Using the Find Method

```
window window = find(head, key);  
Node pred = window.pred;  
curr = window.curr;
```

- The **find()** method returns a structure containing the nodes on either side of the key. It removed marked nodes when it encounters them

## Using the Find Method

```
Window window = find(head, key);
```

```
Node pred = window.pred;  
curr = window.curr;
```

**Find returns window**

## Using the Find Method

```
window window = find(head, key);
```


```
Node pred = window.pred;  
curr = window.curr;
```

**Extract pred and curr**

# The Find Method

```
window window = find(item);
```

At some instant,



or ...




# The Find Method

```
window window = find(item);
```

At some instant,

item

not in list



# Remove

```
public boolean remove(T item) {
    Boolean snip;
    while (true) {
        Window window = find(head, key);
        Node pred = window.pred, curr = window.curr;
        if (curr.key != key) {
            return false;
        } else {
            Node succ = curr.next.getReference();
            snip = curr.next.compareAndSet(succ, succ, false,
true);
            if (!snip) continue;
            pred.next.compareAndSet(curr, succ, false,
false);
            return true;
        }
    }
}
```

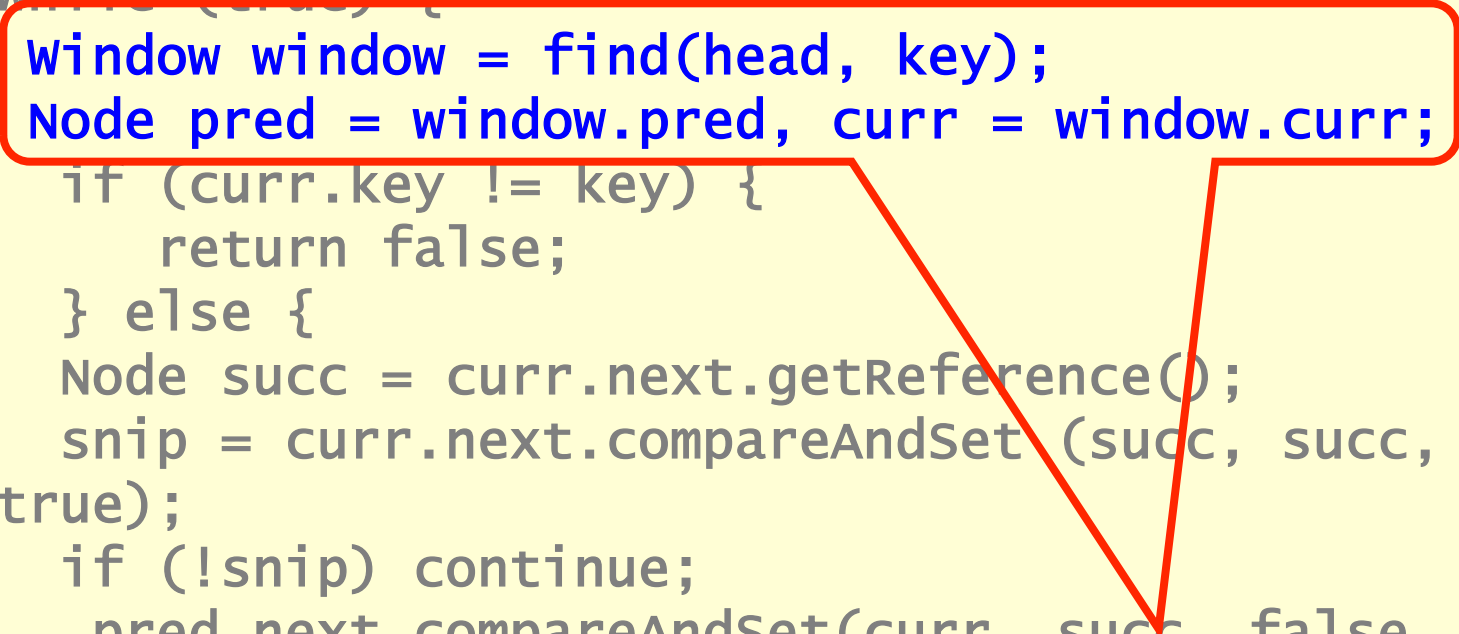
# Remove

```
public boolean remove(T item) {  
    Boolean snip;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key != key) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet (succ, succ, false,  
true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false,  
false);  
            return true;  
        }  
    }  
}
```

**Keep trying**

# Remove

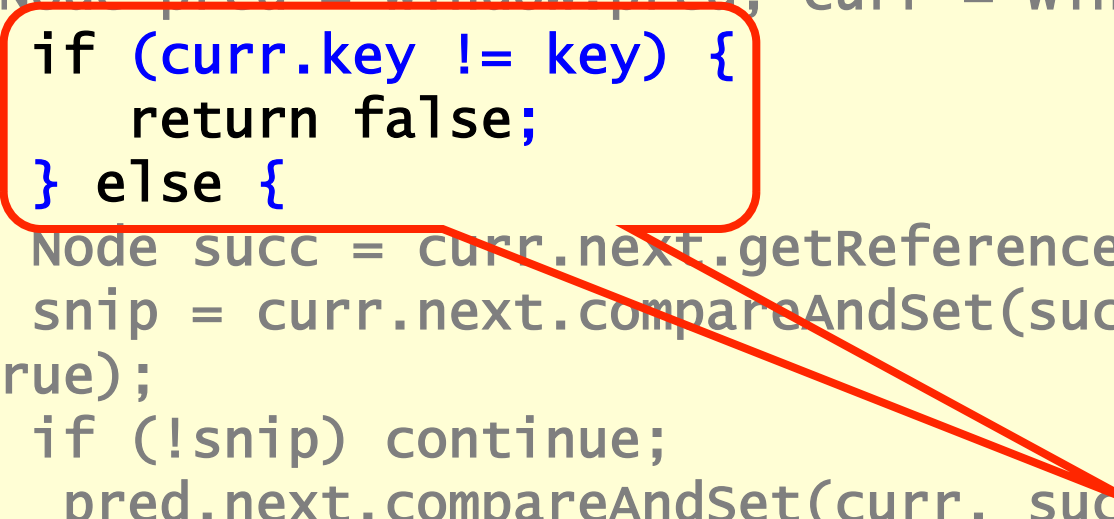
```
public boolean remove(T item) {  
    Boolean snip;  
    while (true) {  
        window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key != key) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet(succ, succ, false,  
true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false,  
false);  
            return true;  
        }  
    }  
}
```



**Find neighbors**

# Remove

```
public boolean remove(T item) {
    Boolean snip;
    while (true) {
        Window window = find(head, key);
        Node pred = window.pred, curr = window.curr;
        if (curr.key != key) {
            return false;
        } else {
            Node succ = curr.next.getReference();
            snip = curr.next.compareAndSet(succ, succ, false,
            true);
            if (!snip) continue;
            pred.next.compareAndSet(curr, succ, false,
            false);
            return true;
        }
    }
}
```



**She's not there ...**

# Remove

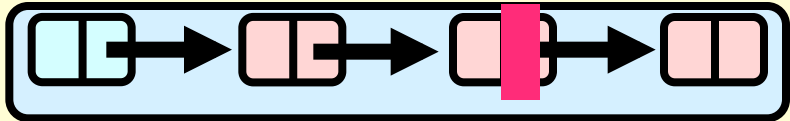
```
public boolean remove(T item) {  
    Boolean snip;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key != key) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet(succ, succ, false,  
true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false,  
false);  
            return true;  
        }  
    }  
}
```

**Try to mark node as deleted**

# Remove

```
public boolean remove(T item) {  
    Boolean success = false;  
    while (curr != null) {  
        window.window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr != null && curr.key.equals(key)) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet(succ, succ, false,  
            true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false,  
            false);  
            return true;  
        }  
    }  
}
```

If it doesn't work, just retry, if it does, job essentially done



# Remove

```
public boolean remove(T item) {
```

```
    Boolean snip;
```

```
    while (true) {
```

```
        window window = find(head,
```

```
        Node pred = window.pred, curr = window.curr;
```

```
        if (curr.key != key) {
```

**Try to advance reference**

**(if we don't succeed, someone else did or will).**

```
        snip = curr.next.compareAndSet(succ, succ, false,
```

```
        true);
```


```
        if (!snip) continue;
```

```
        pred.next.compareAndSet(curr, succ, false,
```

```
        false);
```

```
        return true;
```

```
    }  
}
```





# Add

```
public boolean add(T item) {
    boolean splice;
    while (true) {
        Window window = find(head, key);
        Node pred = window.pred, curr = window.curr;
        if (curr.key == key) {
            return false;
        } else {
            Node node = new Node(item);
            node.next = new AtomicMarkableRef(curr, false);
            if (pred.next.compareAndSet(curr, node, false,
false)) {return true;}
        }
    }
}
```


## Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key == key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
        }  
    }  
}
```

**Item already there.**

# Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key == key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
        }  
    }  
}
```




**create new node**

# Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        (curr, key)  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
        }  
    }  
}
```

**Install new node,  
else retry loop**



## Wait-free Contains


```
public boolean contains(T item) {  
    boolean marked;  
    int key = item.hashCode();  
    Node curr = this.head;  
    while (curr.key < key)  
        curr = curr.next;  
    Node succ = curr.next.get(marked);  
    return (curr.key == key && !marked[0])  
}
```

# Wait-free Contains

```
public boolean contains(T item) {  
    boolean marked;  
    int key = item.hashCode();  
    Node curr = this.head;  
    while (curr.key < key)  
        curr = curr.next;  
    Node succ = curr.next.get(marked);  
    return (curr.key == key && !marked[0])  
}
```

**get and check  
marked together  
with key**

## Summary: Wait-free Contains



Use Mark bit + list ordering

1. Not marked  $\rightarrow$  in the set
2. Marked or missing  $\rightarrow$  not in the set

# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false}; boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
            curr = succ;
        }
    }
}
```




# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false}; boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getreference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
            curr = succ;
        }
    }
}
```

**If list changes  
while  
traversed,  
start over  
Lock-Free  
because we  
start over only  
if someone else  
makes progress**

# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null;
    boolean[] marked = {false};
    boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
            curr = succ;
        }
    }
}
```



# Lock-free Find

```
public Window find(Node head, int key) {  
    Node pred = null, curr = null, succ = null;  
    boolean[] marked = {false}; boolean snip;  
    retry: while (true) { Move down the list  
        pred = head;  
        curr = pred.next.getReference();  
        while (true) {  
            succ = curr.next.get(marked);  
            while (marked[0]) {  
                ...  
            }  
            if (curr.key >= key)  
                return new Window(pred, curr);  
            pred = curr;  
            curr = succ;  
        }  
    }  
}
```

# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false}; boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
            curr = succ;
        }
    }
}
```

**Get ref to successor and  
current deleted bit**

# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false}; boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
        }
    }
}
```

**Try to remove deleted nodes in  
path...code details soon**

# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false}; boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        succ = curr.next.getReference();
        if (curr.key >= key)
            return new Window(pred, curr);
        pred = curr;
        curr = succ;
    }
}
```

**If curr key that is greater or equal, return pred and curr**

# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false};
    boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        while (true) {
            succ = curr.next.getReference();
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
            curr = succ;
        }
    }
}
```

**Otherwise advance window and loop again**

# Lock-free Find


```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
                                         succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked);  
    }  
    ...  
}
```



# Lock-free Find

Try to snip out node

```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
                                         succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked);  
    }  
    ...  
}
```




# Lock-free Find

if predecessor's next field  
changed must retry whole

traversal

```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked);  
    }  
    ...  
}
```



# Lock-free Find

Otherwise move on to  
check if next node deleted

```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
                                         succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked);  
    }  
    ...  
}
```

## “To Lock or Not to Lock”

- **Locking vs. Non-blocking:** Extremist views on both sides
- **The answer:** nobler to compromise, combine locking and non-blocking
  - Remember: Blocking/non-blocking is a property of a method
  - Individual methods of a given item implementation don't necessarily provide the same properties

# This work is licensed under a Creative Commons Attribution-ShareAlike 2.5 License.

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.