

ID1217 Concurrent Programming
Lecture 15



Distributed Programming with Message Passing

Vladimir Vlassov
KTH/ICT/EECS



ROYAL INSTITUTE
OF TECHNOLOGY

Outline

- Introduction
- Message passing
 - Channels and messages
 - Asynchronous message passing
 - Synchronous message passing
 - Synchronous vs asynchronous message passing

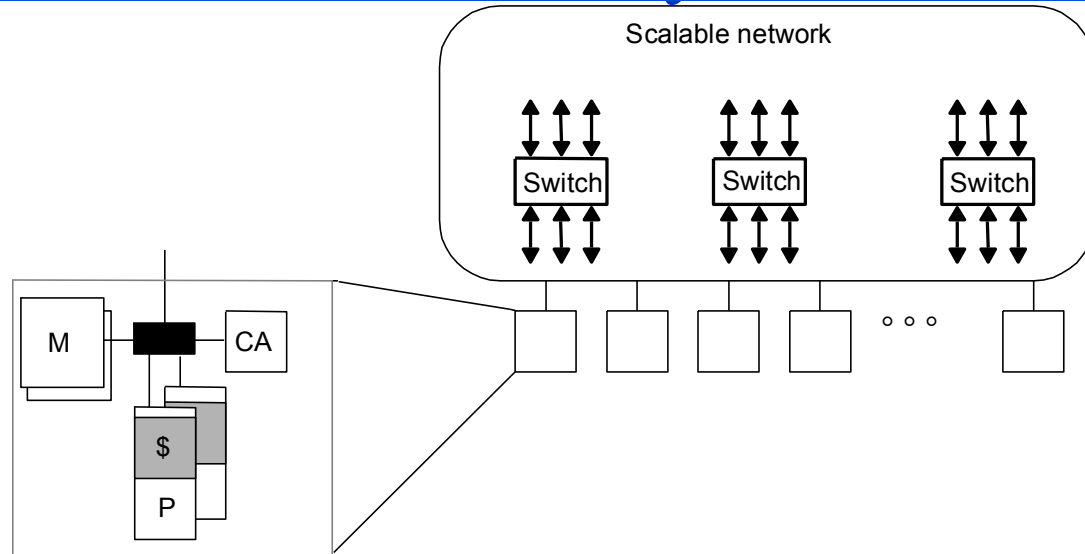


ROYAL INSTITUTE
OF TECHNOLOGY

Distributed Programming

- A **distributed program** is formed of processes distributed among nodes of a distributed memory platform
 - No shared variables, only communication channels are shared
 - Processes have to exchange messages in order to interact
- **Message passing** provide ability for moving data across process memory spaces
 - Message passing primitives: **send** and **receive**
 - Distributed programs use message passing for inter-process communication
- Programming options
 - Design a special programming language or extend an existing one
 - Provide an existing (sequential) language with an external library for distributed programming, e.g. socket API, MPICH

Distributed Memory Architecture



- A node is a single-processor or SMP (Symmetric MultiProcessor)
- No physical memory shared across nodes
- Scalable interconnection network for communication
 - Add nodes – add switches
 - Large number of independent communication paths between nodes
 - Allows a large number of concurrent network transactions
 - Bulk transfer

Some Existing Mechanisms and Environments

- Communication mechanisms:
 - Asynchronous message passing
 - Synchronous message passing
 - Remote procedure call (RPC) and rendezvous
 - Remote method invocation (RMI)
- Distributed programming environments and APIs:
 - Berkley Sockets API – substrate technology
 - RPC API, Java socket API, CORBA, Java RMI
 - PVM (Parallel Virtual Machine)
 - **MPI (Message Passing Interface)**
 - Pro-Active component programming environment
 - Web-service programming environments in Java EE
 - ...




Programming with Distributed Processes

- Processes of a distributed applications are heavy-weight processes started asynchronously on different (or the same) computers
 - One proc can spawn a child by **fork**
 - The child can execute a different program by **exec** (or **reexec**)
- **Single Program Multiple Data (SPMD) model**
 - One program is written for several processes (tasks)
 - Control (conditional) statements are used to select a task for each process
 - To select a task, process needs to know: the total number of processes in the application (or in a group), its own number (rank), and a problem size.
 - For example: MPICH uses SPDM
- **Multiple Program Multiple Data (MPMD) model**
 - Different programs are written for different processes (tasks)
 - For example, servers and clients
 - For example, PVM allows MPMD

Message Passing

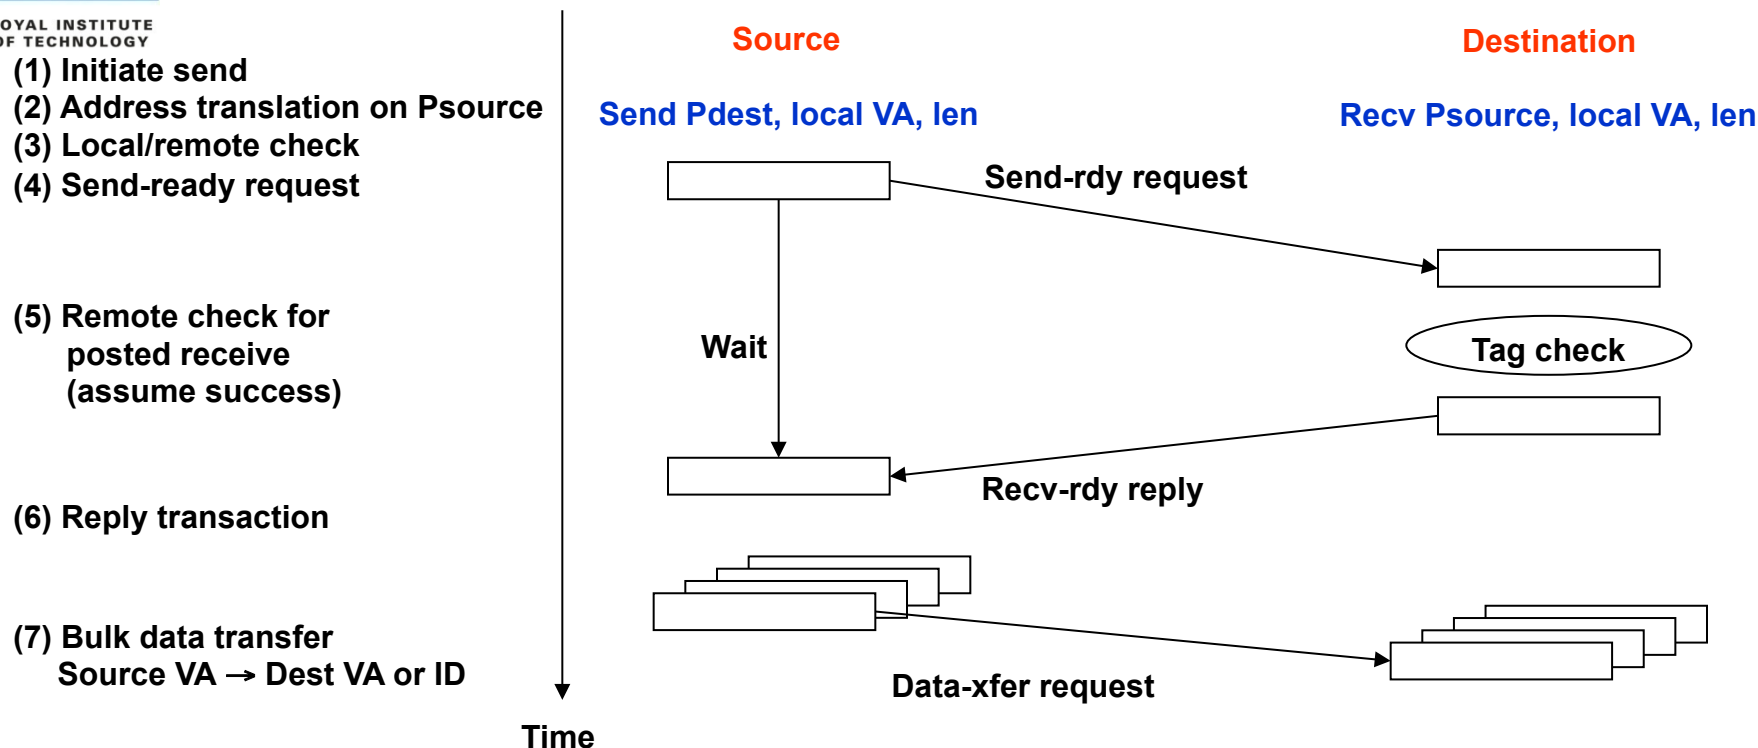
- **Message passing** is sending and receiving messages via communication channels shared by processes
 - Send/receive is one-way data transfer from the memory of a source (sending) proc to the memory of destination (receiving) proc
- Message passing involves synchronization - a message cannot be received until it has been sent.
- A channel can be implemented as:
 - a shared buffer – on a single processor or a shared memory multiprocessor
 - a communication link (communication HW assist) – on a distributed memory multiprocessor or network of workstations



Message Passing Models

- **Synchronous message passing** – blocking semantics of both, send and receive
 - Send completes after matching receive is posted and data are sent
 - Receive completes after data has been received from the matching send
 - Channel can be implemented without buffering
- **Asynchronous message passing** – non-blocking send and blocking receive – send is asynchronous w.r.t. receive
 - Send completes after the send buffer may be reused
 - Channel is an unbounded FIFO queue of messages

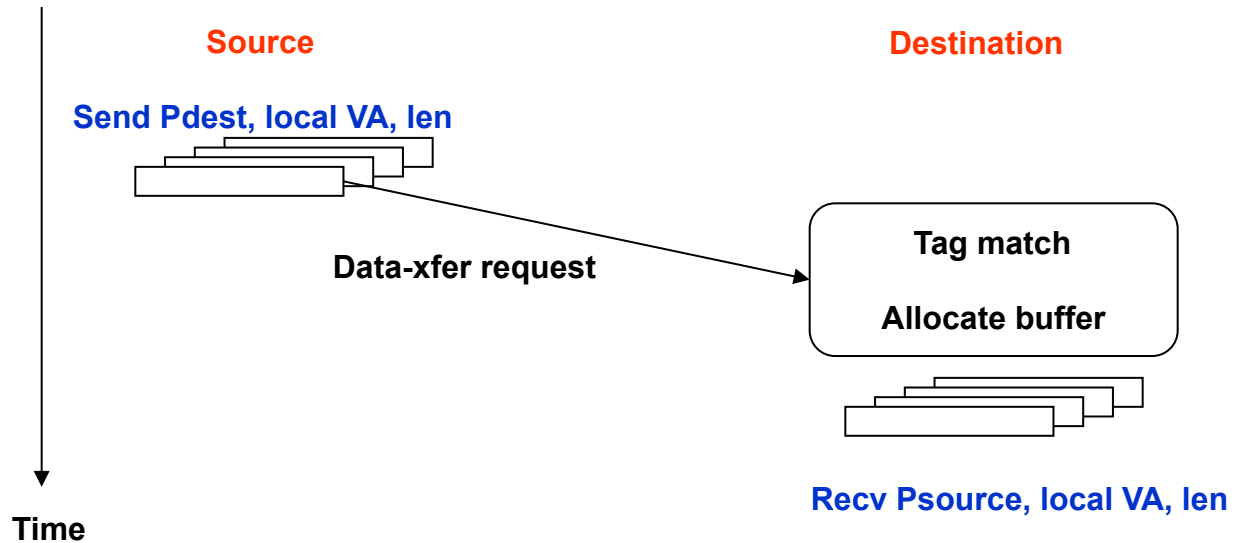
Synchronous Message Passing



- **3-phase protocol** (match table at receiver): send-rdy, recv-rdy and data-xfer
 - 2-phase protocol (receiver-initiated, match table at sender): recv-rdy and data-xfer. Receiver must know who will send
- Constrained programming model. Deterministic

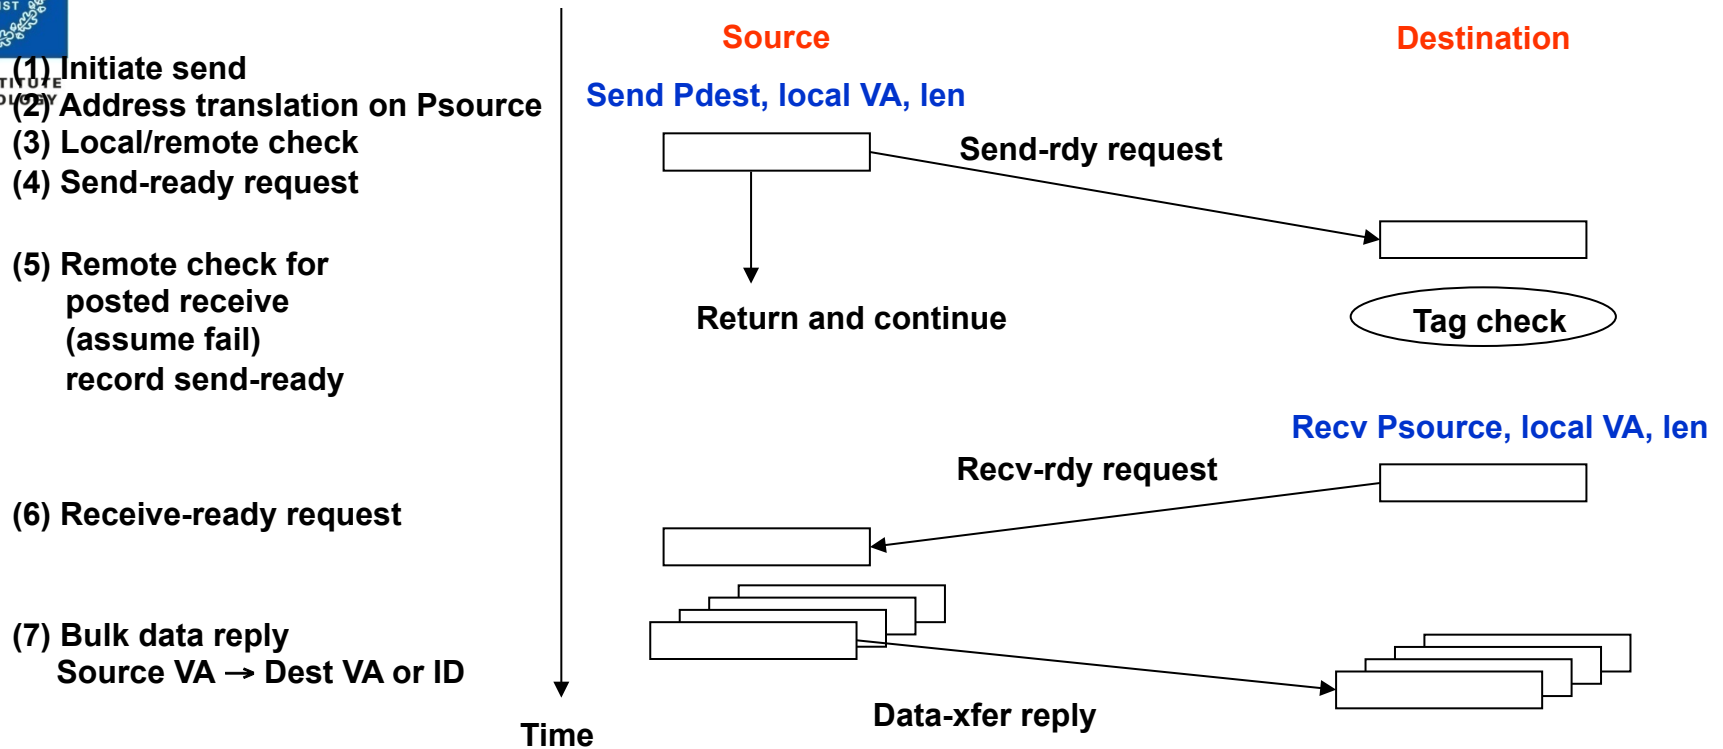
Asynchronous Message Passing: Optimistic

- (1) Initiate send
- (2) Address translation
- (3) Local/remote check
- (4) Send data
- (5) Remote check for posted receive; on fail, allocated buffer



- **Assumes an unbounded amount of storage** outside usual process memory
 - Storage for buffering required within msg layer.
 - Problem of input buffer overflow
- **No ready-to-receive handshaking before the data transfer**
- More powerful programming model
 - Allows to avoid deadlocks
 - Wildcard receive => **non-deterministic**

Asynchronous Message Passing: Conservative



- **3-hase protocol with hand-shake.**
 - Send returns as soon as data are buffered out
 - Where is the buffering? First at sender then can move to receiver (if it will)
 - Short message optimizations - send without hand-shake, use destination credit (reserved space), check the credit locally before sending

Basic Interaction Patterns in Distributed Programs



- **Producer-Consumer** (pipelines) – one way
 - Output is a function of the input and the local state
 - Receive then send
- **Client/server** – two way as master/slave
 - Client initiates and sends first, server receives first and then replies
 - Should interact with a pure request/response protocol to avoid deadlocks
 - Roles may change in particular interaction
- **Interacting peers** – two way as equals
 - Typically interaction is initiated by one of the peers
 - All should be ready to receive

Asynchronous Message Passing

- First we consider asynchronous message passing
 - Assumes arbitrary storage “outside the local address spaces”
 - **Shared channels** – communication abstraction – unidirectional links with unbounded buffering
- Communication primitives
 - **Non-blocking send** – sends data from local memory on a specified channel
 - **Blocking receive** – delays until a message has arrived from a specified channel, receives a message and store it to a specified location

Channels

- **Channel** is a unidirectional communication link from a sending process(es) to a receiving process(es) with an unbounded buffer for a FIFO queue of messages
 - Can be defined as bi-directional links
- Operations:
 - **Send** – append a message to a channel queue
 - **Receive** – remove a message from the head of the channel queue



Channels (cont'd)

- Channels are known to all processes – global variables.
- A single channel:
chan name(type₁ field₁; ...; type_n field_n);
 - List of pairs **type_i field_i** defines a structure of messages transferred over the channel;
 - Optional field names can be omitted:

chan name(type₁; ...; type_n);

- Array of channels:
chan name[n](type₁, ..., type_m);
 - Defines an array of **n** identical channels

- Examples:

chan input(char);

chan disk_access(int cylinder, int block, int count, char* buf);

chan result[n](int)

Communication Primitives

- **Send** a message to a channel **ch**
send ch(expr₁, ..., expr_n);
 - Non-blocking (asynchronous)
 - Evaluates expressions **expr₁, ..., expr_n**
 - Collects results into a message (types must match to those in the channel declaration)
 - Appends the message to the end of the channel queue **ch**
- **Receive** a message from a channel **ch**
receive ch(var₁, ..., var_n);
 - Blocking: Delays the receiver until there is a message in **ch**
 - Extract the message at the head of the channel queue
 - Store message fields to **var₁, ..., var_n**
- Send and Receive are atomic
- FIFO ordering: messages are received in the same order as they were sent



ROYAL INSTITUTE
OF TECHNOLOGY

Communication Primitives (cont'd)


- Additional operation:

boolean **empty(ch);**

- Returns true if channel ch contains no messages, otherwise returns false
- Exposes race condition

Clients and Servers

- **Client** process is proactive
 - Sends a request to a server process over a request channel
 - Waits until the request is serviced and a reply is sent by the server
 - Receives the reply from a reply channel
- **Server** process is reactive
 - Waits for a client request on the request channel
 - Receives and handles the request, generates a reply
 - Sends the reply to the client over the reply channel associated with a client
- In general, server's reply is a function of client's request and the server state



Server and Clients

- The single-threaded server behaves as an active remote monitor.
- Each request type is mapped to a procedure (operation) at the server side.

```
type op_kind = enum(op1, ..., opn);
type arg_type = union(arg1, ..., argn);
type result_type = union(res1, ..., resn);
chan request(int clientID, op_kind, arg_type);
chan reply[n](res_type);

process Server {
    int clientID; op_kind kind; arg_type args;
    res_type results; declarations of other variables;
    initialization code;
    while (true) {      ## loop invariant MI
        receive request(clientID, kind, args);
        if (kind == op1)
            { body of op1; }
        ...
        else if (kind == opn)
            { body of opn; }
        send reply[clientID](results);
    }
}

process Client[i = 0 to n-1] {
    arg_type myargs; result_type myresults;
    place value arguments in myargs;
    send request(i, opj, myargs);      # "call" opj
    receive reply[i](myresults);      # wait for reply
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

Order of Servicing Client Requests

- Order of servicing may depend on the state of the server and on a client request (e.g. type of request, request parameters)
- **FCFS** – first come first served – a simple server architecture without scheduling.
- **“Best” request is serviced first (next)**
 - Best: shortest, largest, system, etc.
 - To choose the current best, server must receive and examine all requests pending in a channel
 - The server maintains a list of received requests – a request list
 - Requests are scheduled from the list in desired order
 - Two threads can do this:
 - One (receiving thread) receives and places requests on the request list
 - Another (servicing/replying thread) fetches a request from the list and services it

Example: A File Server

- Illustrates a multithreaded server and **session** communication – **conversational continuity**
- File server is a part of DFS (distributed file system)
 - Provides for remote clients access to external files on the server's disk
- Client requests:
 - OPEN access to a file – also opens a session
 - READ/WRITE the file
 - CLOSE the file – also closes the session
- Multi-processed (multithreaded) server – to run several concurrent sessions
- A server keeps a session with a client open until it receives CLOSE request from the client

File Server Processes and Channels

- **n** server processes – a scalable session-oriented file server
process File_Server[n]
 - Allow servicing **n** clients (access **n** files) simultaneously
- Request channels
 - One channel for OPEN requests
chan open(fname, clientID)
 - **n** channels for access (READ/WRITE/CLOSE) requests
chan access[n](kind, args)
- Replies channels
 - **m** channels for OPEN replies (one per client)
chan open_reply[m](serverID)
 - **m** channels for access replies (one per client)
chan access_reply[m](results)



Client – Server Interaction

- Rendezvous part – opens file and opens session
 - Client with **myid**:
send open(fname, myid);
receive open_reply[myid](serverID);
 - Server with **myid** which services a client with **clientID**
receive open(fname, clientID);
open the file;
if (successful)
send open_reply[clientID](myid);
- Conversation part – **session** – request-response interaction
 - one client to one server, using **clientID** and **serverID**

File Server and Clients

- Channels and file server processes

```
type kind = enum(READ, WRITE, CLOSE);
chan open(string fname; int clientID);
chan access[n](int kind, types of other arguments);
chan open_reply[m](int serverID); # server id or error
chan access_reply[m](types of results); # data, error, ...

process File_Server[i = 0 to n-1] {
  string fname; int clientID;
  kind k; variables for other arguments;
  bool more = false;
  variables for local buffer, cache, etc.;
  while (true) {
    receive open(fname, clientID);
    open file fname; if successful then:
    send open_reply[clientID](i); more = true;
    while (more) {
      receive access[i](k, other arguments);
      if (k == READ)
        process read request;
      else if (k == WRITE)
        process write request;
      else # k == CLOSE
        { close the file; more = false; }
      send access_reply[clientID](results);
    }
  }
}
```


File Server and Clients (cont'd)

- Client processes:

```
process Client[j = 0 to m-1] {  
    int serverID;  declarations of other variables;  
    send open("foo", j);      # open file "foo"  
    receive open_reply[j](serverID); # get back server id  
    # use file then close it by executing the following  
    send access[serverID](access arguments);  
    receive access_reply[j](results);  
    ...  
}
```

Use of Global Channels

- In the File Server example, **n** server processes use one global channel **open** for OPEN requests
- What if cannot have a global channel with multiple receivers?
 - usually each channel has exactly one receiver
- Use a separate manager process that allocates a server process to a client

receive open(...);

Pick a free server process (or create a new one);
forward the open request to it

- Need “done” messages from servers to keep track of free servers
- This implies that the manager receives two kinds of messages:
(1) opens from clients and (2) dones from servers



ROYAL INSTITUTE
OF TECHNOLOGY

Interacting Peers

- Two way as equals: peer-to-peer
 - Server-to-server
 - Business-to-business
 - Client-to-client
- Environments and applications:
 - (Un)Structured overlay networks, DHT (distributed hash tables)
 - File sharing,
 - P2P distributed services (e.g. naming service, etc.)
 - Scalable P2P-based services in distributed systems
 - Distributed games




ROYAL INSTITUTE
OF TECHNOLOGY

Example: Exchange of Values


- A typical problem in peer-to-peer computing.
- Assume, n processes; each has a value; want every process to learn every value
- Assume a collective reduction operation: every process needs to know the smallest and the largest of the values distributed among processes
- Trade-off: Number of messages

Solution Structures


- a) Centralized – star configuration, with coordinator in center
- b) Symmetric – fully-connected (complete) graph
- c) Ring – closed (circular) pipeline



(a) Centralized solution



(b) Symmetric solution



(c) Ring solution

(a) Centralized Solution

- One process collects all values, does all the “work” and broadcasts result to others
- In total $2(n-1)$ messages over $n+1$ channels
 - With a “broadcast” message – n distinct messages

```
chan values(int), results[n](int smallest, int largest);
process P[0] {    # coordinator process
    int v;    # assume v has been initialized
    int new, smallest = v, largest = v; # initial state
    # gather values and save the smallest and largest
    for [i = 1 to n-1] {
        receive values(new);
        if (new < smallest)
            smallest = new;
        if (new > largest)
            largest = new;
    }
    # send the results to the other processes
    for [i = 1 to n-1]
        send results[i](smallest, largest)
}
process P[i = 1 to n-1] {
    int v;    # assume v has been initialized
    int smallest, largest;
    send values(v);
    receive results[i](smallest, largest);
}
```

(b) Symmetric Solution

- Each process sends its value to others, collects all values from others and does the “work” individually
 - In total $n(n-1)$ messages over n channels
 - With a “broadcast” message – n distinct messages

```
chan values[n] (int);
process P[i = 0 to n-1] {
    int v;    # assume v has been initialized
    int new, smallest = v, largest = v; # initial state
    # send my value to the other processes
    for [j = 0 to n-1 st j != i]
        send values[j] (v);
    # gather values and save the smallest and largest
    for [j = 1 to n-1] {
        receive values[i] (new);
        if (new < smallest)
            smallest = new;
        if (new > largest)
            largest = new;
    }
}
```

c) Circular Pipeline


- Two consecutive rounds:
 - (1) Compute current min and max in a pipeline fashion;
 - (2) Distribute results along the ring
- In the first round, each process $P[i]$
 - Receives two values (current min and max) from its predecessor $P[i-1]$
 - Compare the received values with its own value, and
 - Sends result of comparison, min and max, to its successor $P[(i+1) \bmod n]$
 - To avoid deadlock in the first round – $P[0]$ first sends and then receives
- In the second round, each process $P[i]$
 - Receives two values (final min and max) from its predecessor $P[i-1]$
 - Stores the values locally and sends them to its successor $P[(i+1) \bmod n]$

c) Circular Pipeline (cont'd)


- Employs $2n-1$ messages over n channels
- No broadcast

```
chan values[n](int smallest, int largest);
process P[0] { # initiates the exchanges
    int v;    # assume v has been initialized
    int smallest = v, largest = v; # initial state
    # send v to next process, P[1]
    send values[1](smallest, largest);
    # get global smallest and largest from P[n-1] and
    # pass them on to P[1]
    receive values[0](smallest, largest);
    send values[1](smallest, largest);
}
process P[i = 1 to n-1] {
    int v;    # assume v has been initialized
    int smallest, largest;
    # receive smallest and largest so far, then update
    # them by comparing their values to v
    receive values[i](smallest, largest)
    if (v < smallest)
        smallest = v;
    if (v > largest)
        largest = v;
    # send the result to the next processes, then wait
    # to get the global result
    send values[(i+1) mod n](smallest, largest);
    receive values[i](smallest, largest);
    if (i < n-1)
        send values[i+1](smallest, largest);
}
```


Summary of Peer-to-Peer Communication Patterns



(a) Centralized solution



(b) Symmetric solution



(c) Ring solution

- All-to-all communication – $O(n^2)$ messages, but a symmetric solution is easy to program when all know all
- Star and ring configurations – $O(n)$ messages
- The circular pipeline (ring) configuration – can be inefficient
 - In the example, messages circulate around the pipeline two full circles before every process has learned the global result

Synchronous Message Passing

- Blocking semantics of both, send and receive
 - Send completes after a matching receive has been posted and the source has sent data
 - Receive completes after data transfer from the matching send completes
 - Synchronization via a handshake before bulk data transfer
 - 3 phase protocol: send_rdy, receive-rdy, transfer.
- Synchronous send over a channel:
`synch_send ch(expr1, ..., exprn);`
- Channels can be implemented without buffering
- Synchronous communication was introduced and formalized by Tony Hoare in 1978 via CSP (Communicating Sequential Processes).

Synchronous Message Passing Limits

Concurrency

- Assume, Producer and Consumer communicate with synchronous message passing:
 - Producer and Consumer arrive at communication stage at different times;
 - This causes each pair of send/receive to delay, and, as consequence, the total execution time to increase
- Another example: Client-Server with synch. message passing:
 - Unnecessary delays:
 - When Client wants to releases resource managed by Server
 - When Client issues a write request to an output device controlled by Server that does not need to be acknowledged

```
channel values(int);

process Producer {
    int data[n];
    for [i = 0 to n-1] {
        do some computation;
        synch_send values(data[i]);
    }
}

process Consumer {
    int results[n];
    for [i = 0 to n-1] {
        receive values(results[i]);
        do some computation;
    }
}
```



Synchronous Message Passing is Deadlock Prone

- Example: two processes exchange values
- A symmetric solution with synchronous send leads to a deadlock
 - The deadlock would not occur with asynchronous sends

```
channel in1(int), in2(int);  
process P1 {  
    int value1 = 1, value2;  
    synch_send in2(value1);  
    receive in1(value2);  
}  
  
process P2 {  
    int value1, value2 = 2;  
    synch_send in1(value2);  
    receive in2(value1);  
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

Asynchronous vs Synchronous Message Passing

- **Asynchronous message passing:**
 - Pros
 - More convenient for programming;
 - More concurrency; less deadlock situations
 - Cons:
 - Assumes an unbounded amount of storage outside process memory
 - More complicated to implement: the amount of memory space for channels is unpredictable (and so must be allocated dynamically)
- **Synchronous message passing:**
 - Pros
 - Channels can be implemented without intermediate buffering
 - Cons:
 - Limits concurrency; deadlock prone