

# ID1217 Concurrent Programming

## Lecture 17



## Remote Procedure Call (RPC); Rendezvous; Remote Method Invocation (RMI). Java RMI

Vladimir Vlassov  
KTH/ICT/EECS



ROYAL INSTITUTE  
OF TECHNOLOGY

# Outline

- Remote Procedure Call (RPC)
  - Syntax of RPC
  - Executing a remote procedure call
- Rendezvous
  - Input statements (accept, select)
  - Typical rendezvous
- Examples of RPC and rendezvous
- RPC versus Rendezvous
- Multiple primitives notation
- Remote Method Invocation (RMI)
- Case study: Java RMI



ROYAL INSTITUTE  
OF TECHNOLOGY

# Java RMI References

- Trail: RMI  
<http://docs.oracle.com/javase/tutorial/rmi/index.html>
- Java RMI  
<http://docs.oracle.com/javase/8/docs/technotes/guides/rmi/index.html>




# Motivation for RPC and RMI

- **Message passing** is convenient for consumers-producers (filters) and P2P, but it is somewhat low level for client-server applications
  - Client/server interactions are based on a request/response protocol;
  - Client requests are typically mapped to procedures on server;
  - A client waits for a response from the server.
- Need for more convenient (easier to use) communication mechanisms for developing client/server applications
- **Remote Procedure Call (RPC) and rendezvous**
  - Procedure interface; message passing implementation
- **Remote Method Invocation (RMI)**
  - RMI is an object-oriented analog of RPC
- RPC, rendezvous and RMI are implemented on top of message passing.


# RPC: Remote Procedure Call

- **RPC** is a mechanism that allows a program running on one computer (VM) to cause a procedure to be executed on another computer (VM) without the programmer needing to explicitly code for this.
- Two processes involved: **caller** and **callee**.
- **Caller (RPC client)** is a **calling process** that initiates an RPC to a server.
- **Callee (RPC server)** is a **called process** that accepts the call.



# RPC: Remote Procedure Call (cont'd)

- Each RPC is executed in a **separate process (thread)** on the server side
- **An RPC is a synchronous operation.**
  - The caller is suspended until the results of the remote procedure are returned.
  - Like a regular or local procedure call.
  - Guess why?





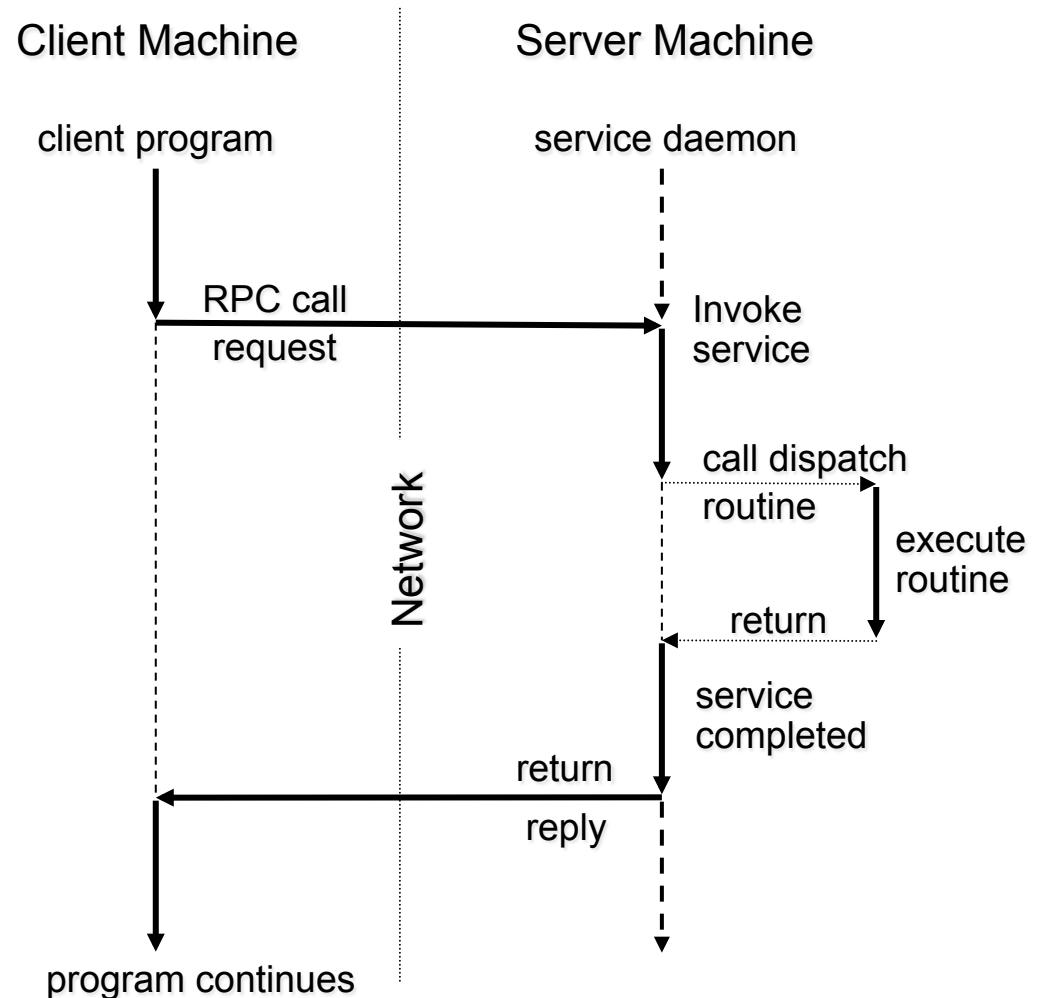
ROYAL INSTITUTE  
OF TECHNOLOGY

# Identifying a Remote Procedure

- Each RPC procedure is **uniquely identified** by
  - A program number
    - identifies a group of related remote procedures
  - A version number
  - A procedure number
- An RPC call message has **three unsigned fields**:
  - Remote program number
  - Remote program version number
  - Remote procedure number
  - The three fields uniquely identify the procedure to be called.

# Executing RPC

- On each RPC the server starts a **new process** to execute the call.
- The new process terminates when the procedure returns and results are sent to the caller.
- Calls from the same caller and calls from different callers are serviced by **different concurrent processes** on server.
- Concurrent invocations might interfere with each other when accessing shared objects – **might need synchronization**





# RPC Syntax

## Modules (Servers)

```
module mname
    interface, i.e. headers of exported operations;
body
    variable declarations;
    initialization code;
    procedures for exported operations;
    local procedures and processes;
end mname
```

- Exported operation (method of a remote interface)

```
op opname(formal identifiers) [returns result]
```

- Procedure – operation implementation


```
proc opname(formal identifiers) returns result identifier
    declarations of local variables;
    statements
end
```

- Client makes a remote call to a module (server):

```
call mname.opname(arguments)
```


# Rendezvous

- **Rendezvous** – remote invocation on a running process
- The server proc is not always available, and execution of a call might be deferred on the server until the server **accepts** the call
- Central inter-process communication mechanism in ADA



# Properties of Rendezvous

- Each remote call is serviced by **the same server process**
- As only one servicing process is active, so avoiding interference.
- In either case, RPC or Rendezvous, **client blocks until remote procedure completes** and result is returned (received).



# Rendezvous Syntax (similar to ADA)

- Declaration of operations (**server interface**):

```
op opname (types of formals);
```

- **Input statement:**

- The simplest form (**accept**) defines one operation that can be called

```
in opname (formal identifiers) [returns result] -> S; ni
```

- The general form (**select**) defines a list of guarded operations

```
in opname1 (formals1) and B1 by e1 -> S1;  
[] ...  
[] opnamen (formalsn) and B2 by en -> Sn;  
ni
```

- **Call** of a remote operation by client – like in RPC

```
call mname.opname (arguments)
```

# The Input Statement

```
in opname1 (formals1) and B1 by e1 -> S1;  
[] ...  
[] opnamen (formalsn) and B2 by en -> Sn;  
ni
```


- The part before **->** is a **guard** of **S<sub>i</sub>**
  - **B<sub>i</sub>** – **optional synchronization expression**;
  - **e<sub>i</sub>** – **optional scheduling expression**;
  - Synchronization and scheduling expressions may reference formals.
- Execution rules:
  - Execution of the input statement delays until some guard succeeds;
  - A guard of **op** succeeds when the **op** has been called and **B** is true (or omitted);
  - If more then one guard succeeds, the oldest one is selected if **e** is omitted;
  - **e** controls selection: an **op** with minimum value of **e** is selected for execution
  - The statement terminates, when execution of the selected **op** completes and results are sent to the caller.

# Some Typical Rendezvous


- Basic input statement defines a remote subroutine (similar to **accept** in ADA)  
`in opname(formals) -> body ni`
- Multiple entries (similar to **select** in ADA)  
`in readblk(...) -> ...  
[] writeblk(...) -> ...  
ni`
- Delay based on local state using synchronization expression:  
`in request() and avail > 0 -> avail--; ...  
[] release() -> avail++; ...  
ni`
- Delay based on arguments and local state:  
`in getforks(i) and not eating[left(i)] and not eating[right(i)] -> eating[i]  
= true;  
[] relforks(i) -> eating[i] = false;  
ni`
- Scheduling based on arguments using scheduling expression (by)  
`in request(time) and free by time -> free := false;  
[] release() -> free := true;  
ni`

# Example: A Distributed File System Using RPC

- Modules (servers):
  - **FileCache**
    - A write-back allocate-on-write cache of file blocks
    - Exports **read** and **write** operations
    - On a miss calls remote **FileServer**
  - **FileServer**
    - Provides access to file blocks stored on a disk
    - Exports **readblk** and **writeblk**
    - Uses the local **DiskDriver** process to access the disk.



# File Cache



```


module FileCache    # located on each diskless workstation
  op read(int count; result char buffer[*]);
  op write(int count; char buffer[]);
body
  cache of file blocks;
  variables to record file descriptor information;
  semaphores for synchronization of cache access (if needed);

  proc read(count,buffer) {
    if (needed data is not in cache) {
      select cache block to use;
      if (need to write out the cache block)
        FileServer.writeblk(...);
      FileServer.readblk(...);
    }
    buffer = appropriate count bytes from cache block;
  }

  proc write(count,buffer) {
    if (appropriate block not in cache) {
      select cache block to use;
      if (need to write out the cache block)
        FileServer.writeblk(...);
    }
    cache block = count bytes from buffer;
  }
end FileCache
  
```



# The File Server and The Disk Driver



```

module FileServer    # located on a file server
  op readblk(int fileid, offset; result char blk[1024]);
  op writeblk(int fileid, offset; char blk[1024]);
body
  cache of disk blocks;
  queue of pending disk access requests;
  semaphores to synchronize access to the cache and queue;
  # N.B. synchronization code not shown below

  proc readblk(fileid, offset, blk) {
    if (needed block not in the cache) {
      store read request in disk queue;
      wait for read operation to be processed;
    }
    blk = appropriate disk block;
  }

  proc writeblk(fileid, offset, blk) {
    select block from cache;
    if (need to write out the selected block) {
      store write request in disk queue;
      wait for block to be written to disk;
    }
    cache block = blk;
  }

  process DiskDriver {
    while (true) {
      wait for a disk access request;
      start a disk operation; wait for interrupt;
      awaken process waiting for this request to complete;
    }
  }
end FileServer
  
```



ROYAL INSTITUTE  
OF TECHNOLOGY

# RPC versus Rendezvous

- In both, RPC and rendezvous
  - A client issues remote calls on a server.
  - Each call is **synchronous**: the client blocks until remote procedure completes and result is returned (received).
- RPC:
  - On each remote call, a server creates a **new process** to handle the call.
  - A server should be always available.
  - As the calls are executed in separate threads – **synchronization** might be needed to avoid interference between calls (due to shared resources).
- Rendezvous:
  - Calls are accepted, selected and served by an **existing process** – no interference.
  - Inter-process communication mechanism in ADA.
- Both, RPC and rendezvous are ideally suited for programming client/server applications (interactions).

# Multiple Primitives Environment

- Combines RPC, rendezvous and message passing in one coherent programming environment.
- A server exports operations that can be executed on a client request

- In a new process:

**proc** **op(formals)** { ... }

- In the server process that executes an input statement:

**in** **op1(formals1)** -> **S1**; [] **op2(formals2)** -> **S2**;...**ni**

- Client interacts with a server's module using:

- Either RPC: **call** **Mname.opname(arguments)**;


- Or message passing: **send** **Mname.opname(arguments)**;

- Four possible combinations:
 

	<i>invocation</i>	<i>service</i>	<i>effect</i>
	<b>call</b>	<b>proc</b>	procedure call
	<b>call</b>	<b>in</b>	rendezvous
	<b>send</b>	<b>proc</b>	dynamic process creation
	<b>send</b>	<b>in</b>	asynchronous message passing

# Example: Replicated Files

- Suppose:
  - $n$  distributed file servers, each exports **open**, **close**, **read** and **write** operations.
  - A file is replicated on each server.
  - A client knows and communicates with one (nearest) server.
  - Servers use a write-update protocol to maintain coherency of copies.
- Invariant:
 
$$(nr == 0 \vee nw == 0) \wedge nw \leq 1$$
  - Concurrent reads, exclusive writes.
  - Assume one lock per copy.
  - When a file is open for reading, a local copy of the file must be locked.
  - When the file is open for writing, all replicas of the file must be locked.





ROYAL INSTITUTE  
OF TECHNOLOGY

## The File Server

```
module FileServer[myid = 1 to n]
  type mode = (READ, WRITE);
  op open(mode), close(),      # client operations
    read(result result types), write(value types);
  op startwrite(), endwrite(), # server operations
    remote_write(value types);
body
  op startread(), endread();  # local operations
  mode use; declarations for file buffers;

  proc open(m) {
    if (m == READ) {
      call startread();      # get local read lock
      use = READ;
    } else {                # mode assumed to be WRITE
      # get write locks for all copies
      for [i = 1 to n]
        call FileServer[i].startwrite();
      use = WRITE;
    }
  }

  proc close() {
    if (use == READ)      # release local read lock
      send endread();
    else                  # use == WRITE, so release all write locks
      for [i = 1 to n]
        send FileServer[i].endwrite()
  }
}
```

## File Server (cont'd)

```
proc read(results) {
    read from local copy of file and return results;
}

proc write(values) {
    if (use == READ)
        return with error: file was not opened for writing;
    write values into local copy of file;
    # concurrently update all remote copies
    co [i = 1 to n st i != myid]
        call FileServer[i].remote_write(values);
}

proc remote_write(values) { # called by other servers
    write values into local copy of file;
}

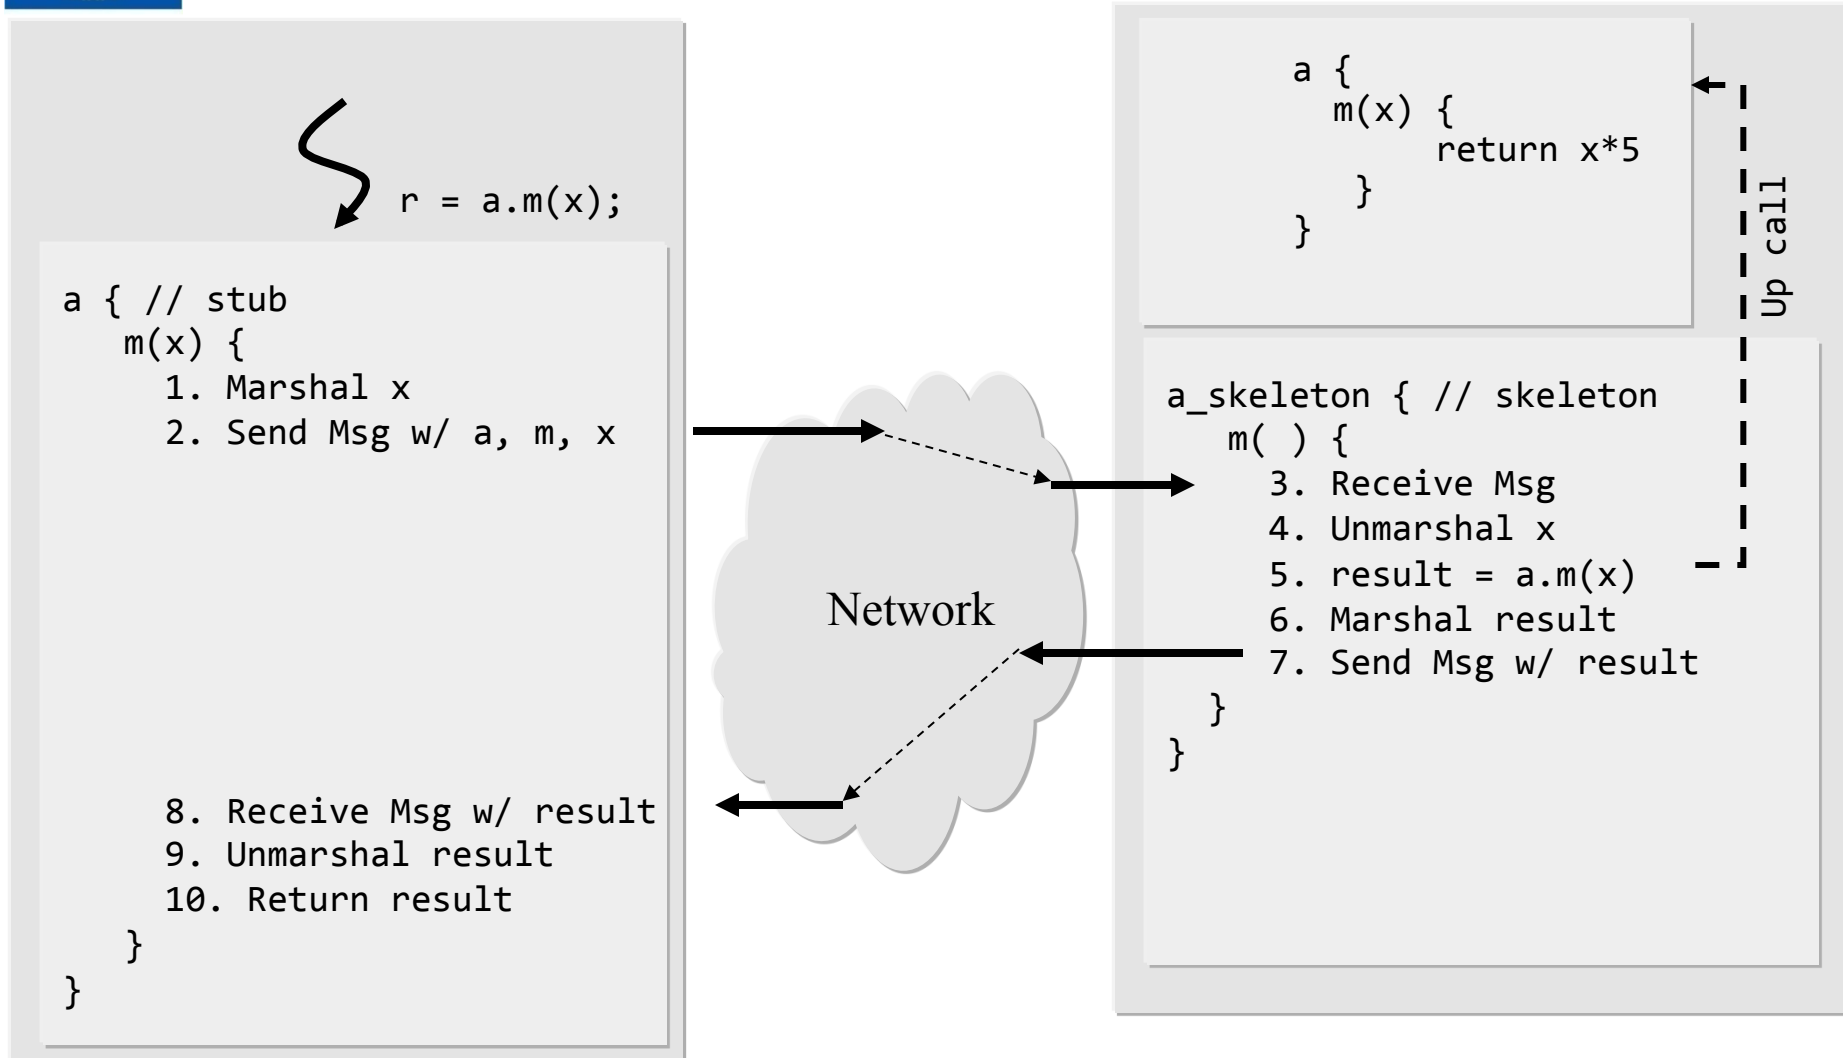
process Lock {
    int nr = 0, nw = 0;
    while (true) {
        ## RW: (nr == 0 ∨ nw == 0) ∧ nw ≤ 1
        in startread() and nw == 0 -> nr = nr+1;
        [] endread() -> nr = nr-1;
        [] startwrite() and nr == 0 and nw == 0 ->
            nw = nw+1;
        [] endwrite() -> nw = nw-1;
        ni
    }
}

end FileServer
```

# Remote Method Invocation (RMI)

- **Remote method invocation (RMI)** is a mechanism to invoke a method on remote object, i.e. object in another computer or virtual machine.
- RMI is the object-oriented analog of RPC in an distributed OO environment, e.g. OMG CORBA, Java RMI, .NET
  - RPC allows calling procedures over a network
  - RMI invokes objects' methods over a network
- **Location transparency**: invoke a method on a stub like on a local object
- **Location awareness**: the stub makes remote call across a network and returns a results via stack

# Remote Method Invocation

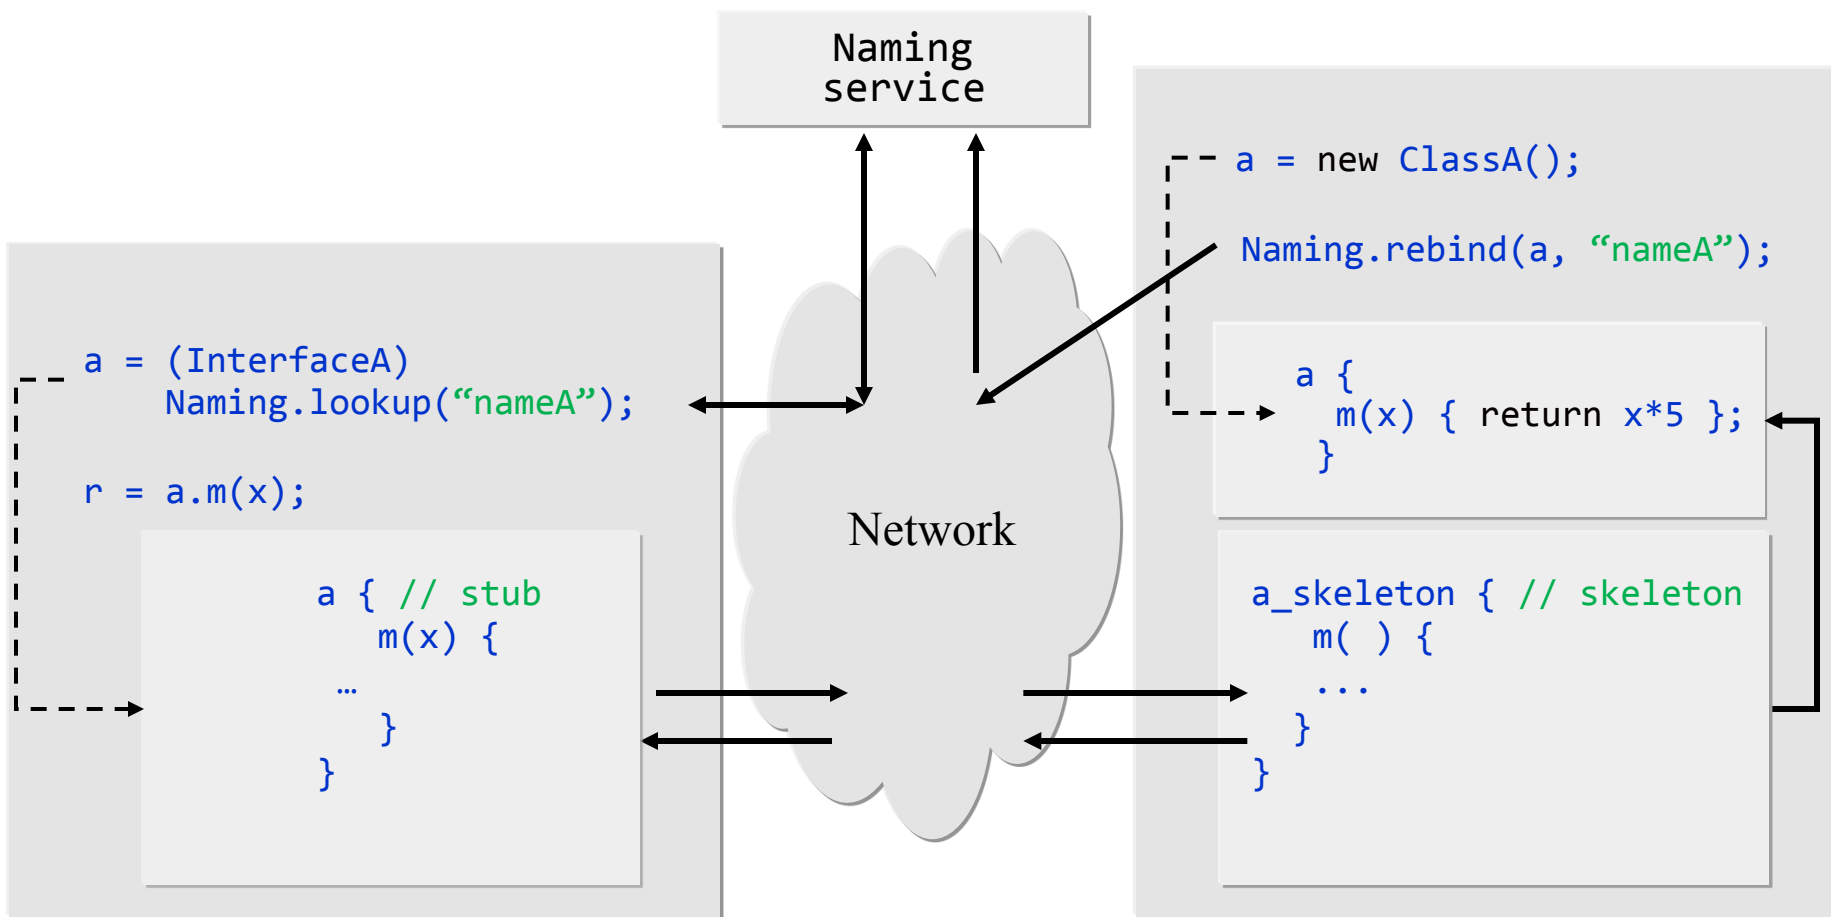




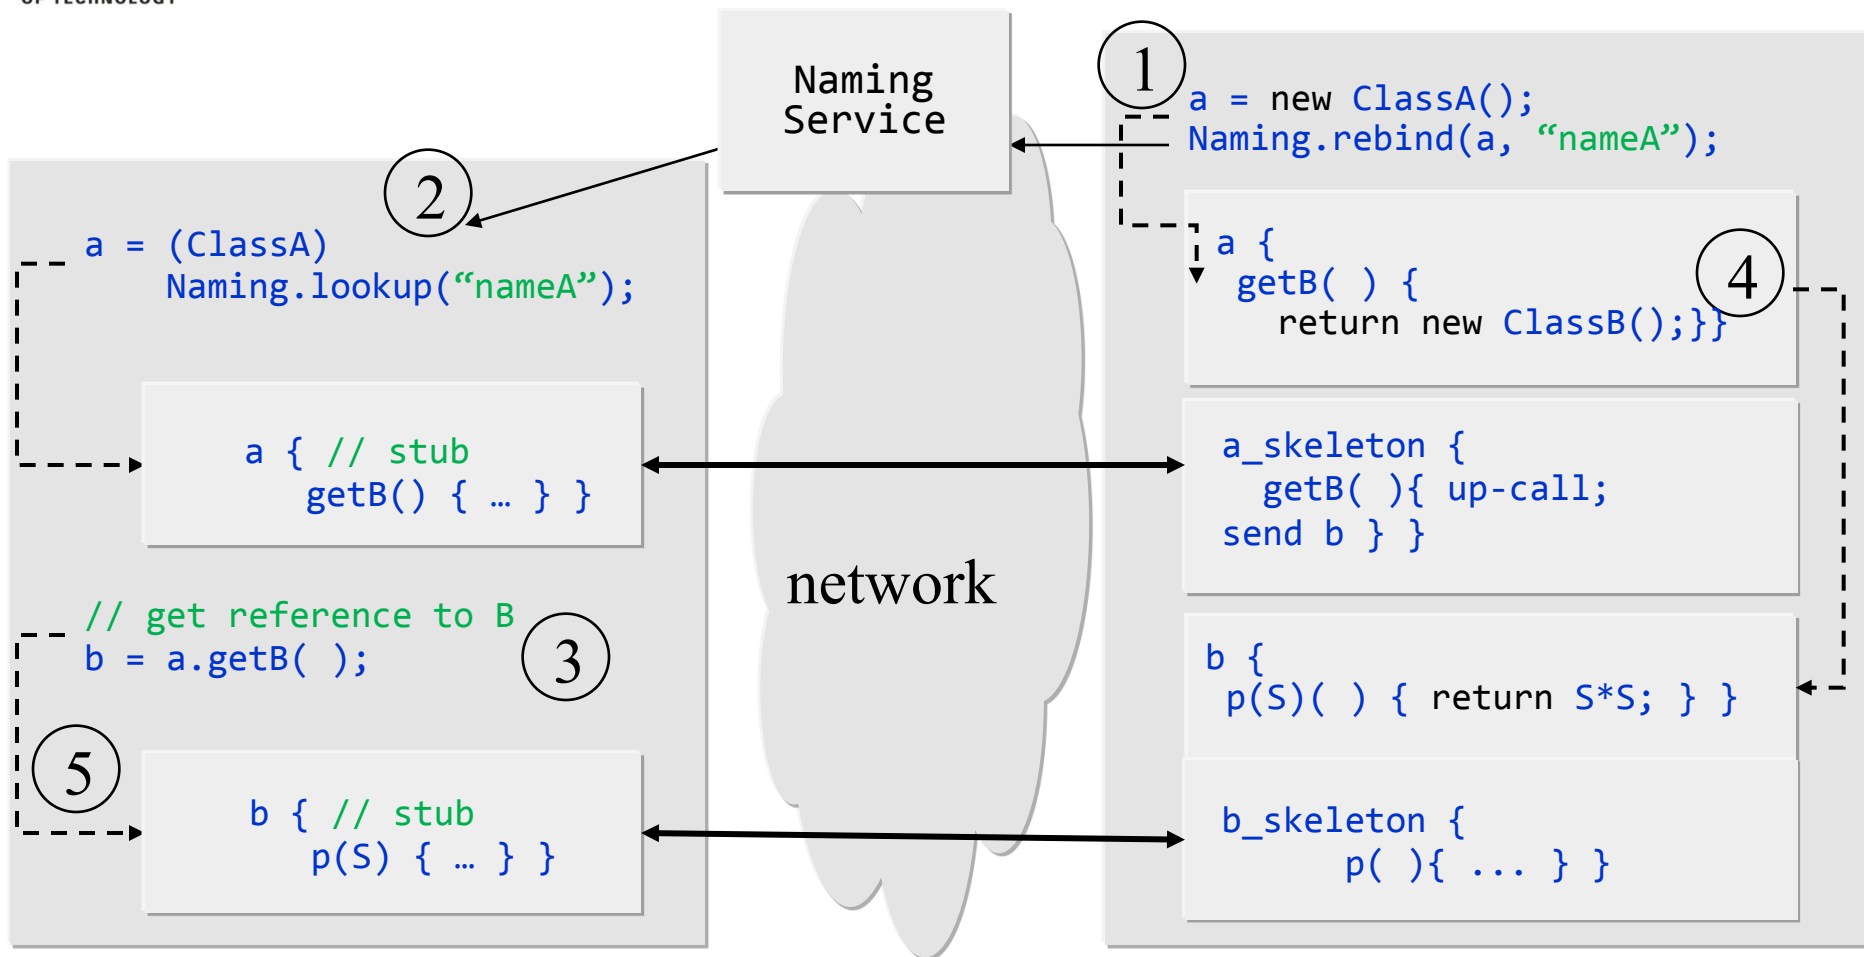
# Locating Objects

- How does a caller get a reference to a remote object, i.e. stub?
- One approach is to use a distributed **Naming Service**:
  - Associate a unique name with an object.
  - Bind the name to the object at the Naming Service.
    - The record typically includes name, class name, object reference (i.e. location information) and other information to create a stub.
  - The client looks up the object by name in the Naming Service.
- The primary reference problem: How to locate the Naming Service?
  - Configuration problem: URL of the naming service

# Use of Naming Service



# Remote Reference in Return





ROYAL INSTITUTE  
OF TECHNOLOGY

# Case Study: Java RMI



# Java RMI (Remote Method Invocation)

- **Java RMI** is a mechanism that allows a thread in one JVM to invoke a method on a object located in another JVM.
  - Provides Java native ORB (Object Request Broker)
- The Java RMI facility allows applications or applets running on different JVMs, to interact with each other by invoking remote methods:
  - Remote reference (stub) is treated as local object.
  - Method invocation on the reference causes the method to be executed on the remote JVM.
  - Serialized arguments and return values are passed over network connections.
  - Uses Object streams to pass objects “by value”.



ROYAL INSTITUTE  
OF TECHNOLOGY

# RMI Classes and Interfaces

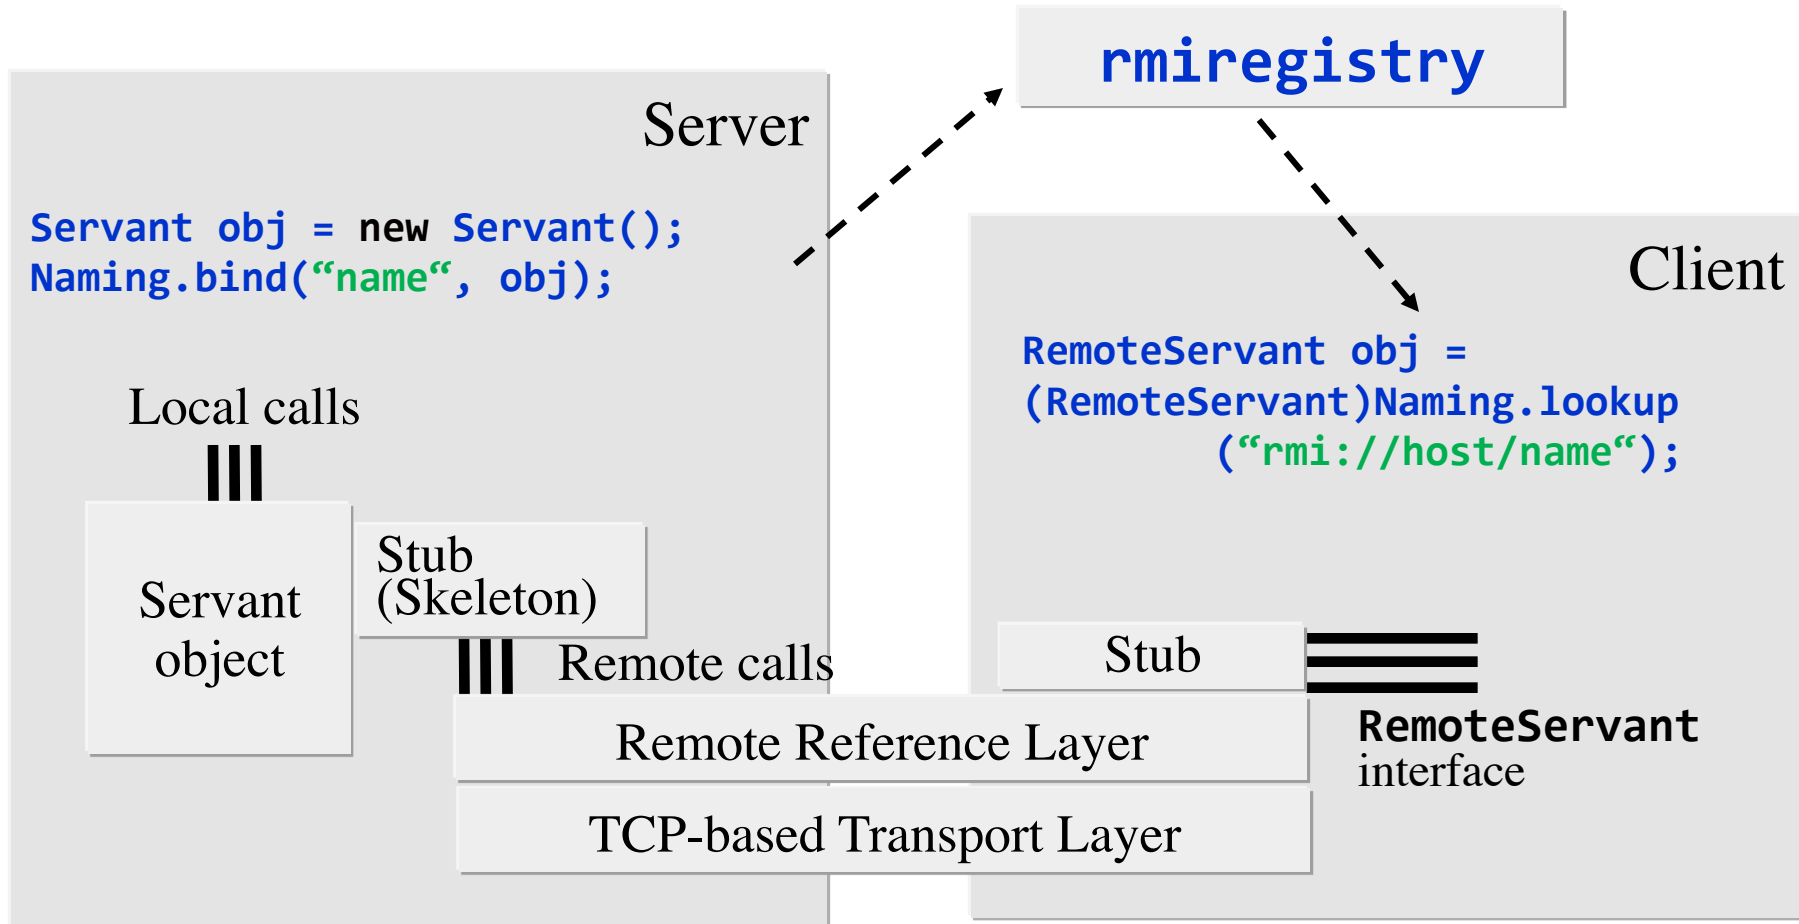
- **java.rmi.Remote**
  - Interface that indicates interfaces whose methods may be invoked from a non-local JVM -- remote interfaces.
- **java.rmi.Naming**
  - The RMI Naming Service client that is used to bind a name to an object and to lookup an object by name.
- **java.rmi.RemoteException**
  - The common superclass for a number of communication-related RMI exceptions.
- **java.rmi.server.UnicastRemoteObject**
  - A class that indicates a non-replicated remote object.

# Developing and Executing a Distributed Application with Java RMI

Typical development and execution steps

1. Define a remote interface(s) that extends **java.rmi.Remote**.
2. Develop a class (a.k.a. servant class) that implements the interface.
3. Develop a server class that provides a container for servants, i.e. creates the servants and registers them at the Naming Service.
4. Develop a client class that gets a reference to a remote object(s) and calls its remote methods.
5. Compile all classes and interfaces using **javac**.
6. (Optional) Generate stubs for the servant classes by **rmic**.
7. Start the Naming service **rmiregistry**
8. Start the server on a server host, and run the client on a client host.

# Architecture of a Client-Server Application with Java RMI





# Declaring and Implementing a Remote Interface

- A remote interface must extend **java.rmi.Remote**
  - Each method must throw **java.rmi.RemoteException**
- A class may implement one or several remote interface
  - The class either extends the **UnicastRemoteObject** class or its instance is exported via the static call  
**UnicastRemoteObject.exportObject(Remote obj)**
- An object of the class that implements the remote interface is called a **servant**.
  - A servant is created by a server and lives until the server dies.
  - The servant and the server can be encapsulated into one class (typically, a primary class).
- A **stub** and a **skeleton** are generated from a servant class automatically

# The Naming Service **rmiregistry**.

## The Naming Client **Naming**

- A Remote object can be registered with a specified name at the Naming service, **rmiregistry**, provided in J2SE.
  - A name can be specified as a URL of the form **rmi://host:port/objectName**
  - The URL indicates host/port of rmiregistry (defaults to localhost:1099).
- The **Naming** class provides a static client of the RMI registry.

- A server binds a name to an object:

```
try {  
    Bank bank = new BankImpl("CityBank");  
    Naming.rebind("rmi://" + host + ":" + port +  
        "/CityBank", bank);  
    System.out.println(bank + " is ready.");  
} catch (Exception e) {  
    e.printStackTrace();  
}
```

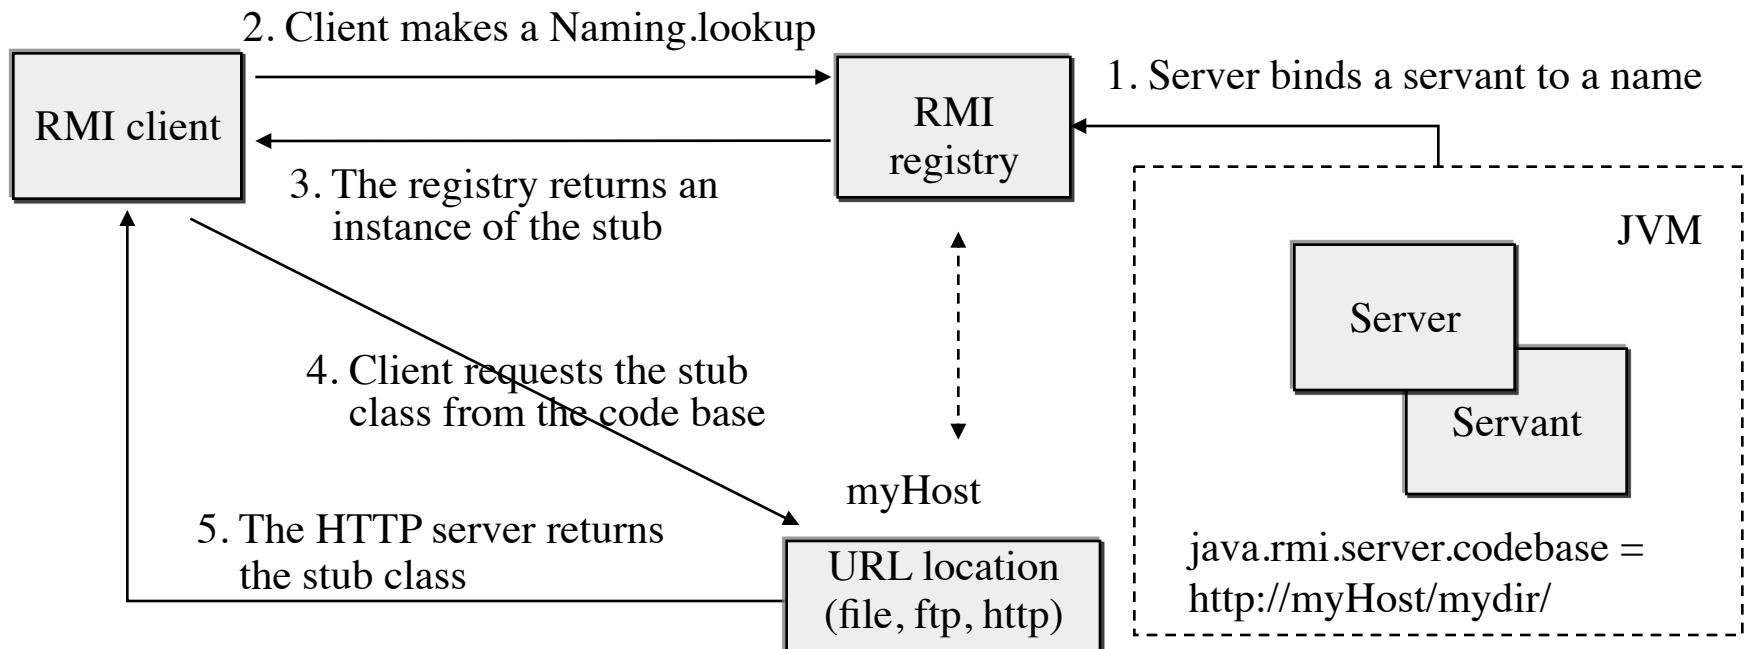
- A client looks up a remote object by name:

```
String rmiregURL = "rmi://" + host;  
try {  
    bank = (Bank)Naming.lookup(rmiregURL + "/CityBank");  
} catch (Exception e) {  
    System.out.println("The runtime failed: " + e);  
    System.exit(0);  
}
```

# Loading Stub Classes

**Stubs are dynamically generated and loaded when needed** either from the local file system or from the network using the URL specified on server side using the `java.rmi.server.codebase` property.

- The property can be set in a command line of an application, for example:  
`-Djava.rmi.server.codebase=http://myHost/mydir/`
- See: <http://java.sun.com/j2se/1.4.2/docs/guide/rmi/spec/rmi-properties.html>





# Parameters and Returns in Java RMI

- Primitive data types and non-remote **Serializable** objects are passed by values.
  - If an object is passed by value, it is cloned at the receiving JVM, and its copy is no longer consistent with the original object.
  - The class name collision problem. Versioning.
- Remote objects are passed by references.
  - A remote reference can be returned from a remote method. For example:

```
try {  
    // lookup for the bank at rmiregistry  
    Bank bankobj = (Bank)Naming.lookup(bankname);  
    // create a new account in the bank  
    Account account = bankobj.newAccount(clientname);  
    account.deposit(value);  
} catch (Rejected e) { process the exception }  
...
```
  - A remote object reference can be passed as a parameter to a remote method.



ROYAL INSTITUTE  
OF TECHNOLOGY

# Example: A Bank Manager

- An application that controls accounts.
- Remote interfaces:
  - **Account** – deposit, withdraw, balance;
  - **Bank** – create a new account, delete an account, get an account;
- Classes that implement the interfaces:
  - **BankImpl** – a bank servant class the implements the **Bank** interface used to create, delete accounts;
  - **AccountImpl** – a account servant class implements the **Account** interface to access accounts.



ROYAL INSTITUTE  
OF TECHNOLOGY

# Bank and Account Remote Interfaces

- The **Bank** interface

```
package bankrmi;
import java.rmi.*;
import bankrmi.Account;
import bankrmi.Rejected;
public interface Bank extends Remote {
    public Account newAccount(String name) throws RemoteException, Rejected;
    public Account getAccount (String name) throws RemoteException;
    public boolean deleteAccount(String name) throws RemoteException, Rejected;
}
```

- The **Account** interface

```
package bankrmi;
import java.rmi.*;
import bankrmi.Rejected;
public interface Account extends Remote {
    public float balance() throws RemoteException;
    public void deposit(float value) throws RemoteException, Rejected;
    public void withdraw(float value) throws RemoteException, Rejected;
}
```



# A Sketch of a Bank Implementation

```
package bankrmi;
import java.rmi.server.UnicastRemoteObject;
import java.util.Hashtable;
import java.rmi.*;
import bankrmi.*;
public class BankImpl extends UnicastRemoteObject implements Bank {
    private String _bankname = "Noname";
    private Hashtable _accounts = new Hashtable(); // accounts
    public BankImpl(String name) throws RemoteException {
        super(); _bankname = name;
    }
    public BankImpl() throws RemoteException {
        super();
    }
    public synchronized Account newAccount(String name) throws RemoteException, Rejected {
        AccountImpl account = (AccountImpl) _accounts.get(name);
        if (account != null) {
            System.out.println("Account [" + name + "] exists!!!");
            throw new Rejected("Rejected: Bank: " + _bankname + " Account for: " + name + " already exists: "
+ account);
        }
        account = new AccountImpl(name);
        _accounts.put(name, account);
        System.out.println("Bank: " + _bankname + " Account: " + name + " Created for " + name);
        return (Account)account;
    }
}
```

. . .



# A Sketch of an Account Implementation

```
package bankrmi;
import java.rmi.server.UnicastRemoteObject;
import java.rmi.*;
import bankrmi.*;

public class AccountImpl extends UnicastRemoteObject implements Account {
    private float _balance = 0;
    private String _name = "noname";
    public AccountImpl(String name) throws RemoteException {
        super();
        this.name = name;
    }
    public AccountImpl() throws RemoteException {
        super();
    }
    public synchronized void deposit(float value) throws RemoteException, Rejected {
        if (value < 0) throw new Rejected("Rejected: Account " + name + ": Illegal value: " + value);
        _balance += value;
        System.out.println("Transaction: Account "+name+": deposit: $" + value + ", balance: $" + _balance);
    }
    public synchronized void withdraw(float value) throws RemoteException, Rejected {
        ...
    }
    public synchronized float balance() throws RemoteException { return _balance; }
}
```





# The Server Application

```
package bankrmi;
import java.rmi.*;
import bankrmi.*;
public class Server {
    static final String USAGE = "java bankrmi.Server <bank_url>";
    static final String BANK = "NordBanken";
    public Server(String[] args) {
        String bankname = (args.length > 0)? args[0] : BANK;
        if (args.length > 1 || bankname.equalsIgnoreCase("-h")) {
            System.out.println(USAGE);
            System.exit(1);
        }
        try {
            Bank bankobj = (Bank)(new BankImpl(bankname));
            Naming.rebind(bankname, bankobj);
            System.out.println(bankobj + " is ready.");
        } catch (Exception e) { System.out.println(e); }
        Object sync = new Object();
        synchronized(sync) { try { sync.wait();} catch (Exception ie) {} }
    }
    public static void main(String[] args) {
        new Server(args).start();
    }
}
```



# A Sketch of a Client Application

```
package bankrmi;
import bankrmi.*;
import java.rmi.*;
public class SClient {
    static final String USAGE = "java Client <bank_url> <client> <value>";
    String bankname = "Noname", clientname = "Noname"; // defaults
    float value = 100;
    public SClient(String[] args) {
        // Read and parse command line arguments (see Usage above)
        ...
        try {
            Bank bankobj = (Bank)Naming.lookup(bankname);
            Account account = bankobj.newAccount(clientname);
            account.deposit(value);
            System.out.println (clientname + "'s account: $" + account.balance());
        } catch (Exception e) {
            System.out.println("The runtime failed: " + e);
            System.exit(0);
        }
    }
    public static void main(String[] args) {
        new SClient(args);
    }
}
```