

ID1217 Concurrent Programming
Lecture 18



Paradigms for Process Interaction
in Distributed Computing.
Parallelism in Scientific Computing

Vladimir Vlassov
KTH/ICT/EECS



ROYAL INSTITUTE
OF TECHNOLOGY

Paradigms for Process Interactions

- Basic interaction paradigms in distributed programs:
 1. Producer/consumer (filters) – one way
 2. Client/Server – two ways as master/slave
 3. Interacting peers – two ways as equals
- Larger paradigms (models) for process interactions
 1. Manager/Workers (a.k.a. “bag of tasks” or “work farm”)
 2. Heartbeat algorithms
 3. Pipeline algorithms
 4. Probe/echo algorithms; gossip algorithms
 5. Broadcast algorithms
 6. Token-passing algorithms
- The first three paradigms are commonly used in parallel computing; the others arise in distributed systems.



ROYAL INSTITUTE
OF TECHNOLOGY

1. Manager-Workers (Work Farm)

- Based on a **distributed bag-of-tasks** model
 - Manager maintains a bag of independent tasks and collects results
 - Each worker gets a task from the bag, executes it, and possibly generates new tasks that it puts to the bag.
 - Manager acts as a server; workers – as clients.
 - How can the manager detect termination? every worker is waiting to get a new task (i.e. idle) and the bag is empty.
- Advantages:
 - **Scalability**: Easy to vary the number of workers, and granularity of tasks.
 - **Good load balancing**: Easy to ensure that each worker does about the same amount of work.

Manager/Worker (cont'd)

- Worker:

```
call getTask(...)    // request-response  
compute;  
generate a new task;  
send putResult(...) // asynchronous send
```

- Manager exports operations “get task” and “put result”

```
in getTask(...) -> ... // returns a task or end  
[] putResult(...) -> ... // receives result  
ni
```



ROYAL INSTITUTE
OF TECHNOLOGY

Example: Adaptive Quadrature Using Manager/Workers

```
module Manager
  op getTask(result double left, right);
  op putResult(double area);
body Manager
  process manager {
    double a, b;          # interval to integrate
    int numIntervals;     # number of intervals to use
    double width = (b-a)/numIntervals;
    double x = a, totalArea = 0.0;
    int tasksDone = 0;
    while (tasksDone < numIntervals) {
      in getTask(left, right) st x < b ->
        left = x; x += width; right = x;
      [] putResult(area) ->
        totalArea += area;
        tasksDone++;
    }
    ni
  }
  print the result totalArea;
}
end Manager

double f() { ... }          # function to integrate
double quad(...) { ... }   # adaptive quad function

process worker[w = 1 to numWorkers] {
  double left, right, area = 0.0;
  double fleft, fright, lrarea;
  while (true) {
    call getTask(left, right);
    fleft = f(left); fright = f(right);
    lrarea = (fleft + fright) * (right - left) / 2;
    # calculate area recursively as shown in Section 1.5
    area = quad(left, right, fleft, fright, lrarea);
    send putResult(area);
  }
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

2. Heartbeat Algorithms

- Divide work (evenly).
- Processes periodically exchange information using a **send (expand)** and then **receive (contract)** interaction.
 - Typically each proc exchanges data with its neighbors
 - The exchanges provide a “fuzzy” barrier among the workers
- Used in many iterative applications with data parallelism
 - E.g. exchange edges in Jacobi iteration


A Typical Heartbeat Algorithm

```
process Worker[i = 1 to numWorkers] {  
    declarations of local variables;  
    initialize local variables;  
    while (not done) {  
        send values to neighbors;  
        receive values from neighbors;  
        update local values;  
    }  
}
```


- Examples: Jacobi iteration, region labeling, game of life

3. Pipeline Algorithms

- What: divide work evenly, compute and circulate data among workers
- When: used when workers need all the data, not just edges from neighbors.
- Pipeline structures –circular or open (or closed)



(a) open



(b) closed



(c) circular

Example: Sieve of Eratosthenes

- A closed pipeline that generates prime numbers

```
chan sieve[n](int);
process Sieve[1] { // coordinator
    int p = 2;
    for (i = 3 to n by 2)
        send sieve[2](i); // send odd numbers to Sieve[2]
    while (p != EOF) {
        println(p);
        receive sieve[1](p);
    }
}
process Sieve[i = 2 to L] {
    int p, next;
    bool sent = false;
    receive sieve[i-1](p); // receive prime p
    send sieve[1](p);      // send to print out
    while (true) {
        receive sieve[i-1](next); // receive next candidate
        if (next mod p) != 0 {    // if it might be prime,
            send sieve[i+1](next); // pass it on
            sent = true;
        }
        if (!sent) { send sieve[1](EOF); break; }
    }
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

4. Probe/Echo Algorithms

- Used to disseminate and/or to gather information
 - **Probes** – to disseminate request or information
 - **Echoes** – to collect information or to acknowledge
 - Each probe should be echoed
- For example:
 - Broadcast a message
 - Using a spanning tree – when knows a global topology
 - Using neighbor sets – when knows neighbors
 - The network topology problem: collect all local topologies and build their union , i.e. the entire network topology
 - Web "crawlers"



ROYAL INSTITUTE
OF TECHNOLOGY

Example: The Network Topology Problem

- Assume each node knows its local topology – neighbors
- Problem: collect all local topologies and build their union – **a network topology**
- Two phases:
 1. Each node sends a probe to its neighbors.
 2. Later each node sends an echo with local topology back to the node from which it received the first probe.
 - Eventually the initiating node **S** collects all echoes and computes a global topology.
 - **S** may distribute the global topology to other processes.



ROYAL INSTITUTE
OF TECHNOLOGY

Probe/echo Algorithm for Gathering the Topology of a Tree.

```
type graph = bool [n,n];
chan probe[n] (int sender);
chan echo[n] (graph topology)    # parts of the topology
chan finalecho(graph topology)  # final topology

process Node[p = 0 to n-1] {
  bool links[n] = neighbors of node p;
  graph newtop, localtop = ([n*n] false);
  int parent;    # node from whom probe is received
  localtop[p,0:n-1] = links;    # initially my links

  receive probe[p] (parent);
  # send probe to other neighbors, who are p's children
  for [q = 0 to n-1 st (links[q] and q != parent)]
    send probe[q] (p);

  # receive echoes and union them into localtop
  for [q = 0 to n-1 st (links[q] and q != parent)] {
    receive echo[p] (newtop);
    localtop = localtop or newtop;  # logical or
  }
  if (p == S)
    send finalecho(localtop);    # node S is root
  else
    send echo[parent] (localtop);
}

process Initiator {
  graph topology;
  send probe[S] (S)    # start probe at local node
  receive finalecho(topology);
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

5. Broadcast Algorithms

- For dissemination of information
- For making decentralized decisions (consensus)
 - Each process must participate in every decision
- For solving many distributed synchronization problems (e.g. distributed semaphores, distributed mutual exclusion)
- Use logical clocks to order communication events

Logical Clocks

- Many distributed algorithms use logical clocks to order communication events
- **A logical clock** (**lc**) is a private integer counter that a proc increases on every communication event.
- Proc attaches a **timestamp** (**ts**) to each message it sends:

```
ts = lc++; send(m, ts);
```

- Proc checks and corrects its **lc** when it receives a message with **ts**:

```
receive(m, ts); lc = max(lc, ts+1); lc++;
```




ROYAL INSTITUTE
OF TECHNOLOGY

Making Decisions (Achieving Consensus)

- Broadcast a message **m** with **ts** (and process id)
- Wait until have a **later message** from **every other** process
- Now know you won't ever see an earlier message, i.e. the message **m** is **fully acknowledged**.
- Hence if your broadcast is the oldest, you can act upon it:
 - Make decision based on **m** and messages in front of **m** in the message queue ordered by **ts** – **stable prefix** – all messages are fully acknowledged since they have smaller timestamps
 - Use logical clocks for ordering
 - Use something like process ids to break ties

6. Token Passing Algorithms

- Used to convey some permission (e.g. distributed mutual exclusion) or to detect termination
- Example: Mutual exclusion with a token ring
 - A lock is a token that is passed in the ring
 - A proc that needs the locks grabs it when the lock passes the proc.



Mutual Exclusion with a Token Ring

- There are **User** and **Helper** processes in the ring: one **Helper** per **User**.
- A token passed in the ring is a mutex lock.
- A **Helper** gets the token and checks if its **User** needs the token to enter its critical section.
- Invariant:

DMUTEX: **User**[*i*] is in its CS \Rightarrow
Helper[*i*] has the token \wedge
 there is exactly one token

```
chan token[1:n](), enter[1:n](), go[1:n](), exit[1:n]();

process Helper[i = 1 to n] {
  while (true) {    # loop invariant DMUTEX
    receive token[i]();      # wait for token
    if (not empty(enter[i])) { # does user want in?
      receive enter[i]();    # accept enter msg
      send go[i]();          # give permission
      receive exit[i]();     # wait for exit
    }
    send token[i%n + 1]();   # pass token on
  }
}

process User[i = 1 to n] {
  while (true) {
    send enter[i]();         # entry protocol
    receive go[i]();
    critical section;
    send exit[i]();          # exit protocol
    non-critical section;
  }
}
```



ROYAL INSTITUTE
OF TECHNOLOGY

Summary of Paradigms for Process Interaction

- Manager/Worker
 - A distributed bag of tasks or work farm model used in many parallel applications
- Heartbeat Algorithms
 - Exchange on each iteration – Used in many iterative applications with a fixed number of tasks and a fixed number of parallel processes –
- Pipeline Algorithms
 - Divide work evenly, compute and circulate data among workers
- Probe/Echo Algorithms
 - Used to disseminate and/or gather information in a distributed application
- Broadcast Algorithms
 - Used to disseminate information or to make decentralized decisions (consensus)
- Token-Passing Algorithms
 - Used to convey permission or to detect termination in a distr. application

Parallelism in Scientific Computing



ROYAL INSTITUTE
OF TECHNOLOGY

Outline

- Demand for parallel computing.
- Parallel scientific computing
 - Speedup, overhead
 - Steps in developing parallel programs
- Techniques in scientific computing
 - Grid computations – modeling of continuous processes and systems
 - Example: The heat problem (the Laplace equation)
 - Particle computations – modeling of discrete systems
 - Example: The N-body problem
 - Matrix computations – algebraic equations, image and signal processing, etc.



ROYAL INSTITUTE
OF TECHNOLOGY

Demand for Parallel Computing

- Why parallel computing?
- Application demands for performance
 - Our insatiable need for computing cycles in many areas
 - Scientific computing,
 - General-purpose computing: Engineering, commercial, financial, entertainment, etc.
- Ways to increase performance:
 - Faster clock – limited?
 - **Parallelism** – alternative to faster clock
 - Allows improving both, execution performance and utilization of CPUs



ROYAL INSTITUTE
OF TECHNOLOGY

Inevitability of Parallel Computing

- Application demands – major driving force
- Technology trends
 - Number of transistors on chip growing rapidly
 - Clock rates growing slowly; furthermore it declines
- Computer architecture trends
 - Instruction-level parallelism valuable but limited
 - Coarser-level parallelism, as in MPs and multi-core processors, the most viable approach
- Economics
 - Multiprocessors being pushed by software vendors (e.g. database) as well as hardware vendors (a line of products from PC to high-end servers)
- Current trends:
 - Today's microprocessors are multicore
 - Servers and workstations are multicore-based: HP, Dell, ...
 - Today's and tomorrow's microprocessors are multiprocessors.




ROYAL INSTITUTE
OF TECHNOLOGY

Scientific Computing

- Scientific computing is a way to examine physical phenomena by **computational modeling**
 - Weather forecast,
 - Modeling of ocean basin,
 - CFD (Computational Fluid Dynamics),
 - Modeling of nuclear reactions,
 - Modeling the evolution of galaxies, etc.
- **High Performance Computing (HPC)**
 - Targets to improving performance by any means – the traditional demand of scientific computing

Scientific Computing Demands. Grand Challenge problems (90s)



Parallel Scientific Computing

Goal: speedup on large problems (or solve an even larger problem)

- **Speedup:** $T_{\text{sequential execution}} / T_{\text{parallel execution on } n \text{ processors}}$.
- Assume N is a problem size; n is a number of processors
- **Amdahl's law** (for a fixed N): If fraction s of sequential execution is inherently serial, then speedup approaches $1/s$ as n approaches infinity

$$\text{speedup}(n) = 1/(s + (1 - s) / n) \leq 1/s$$

- **Gustafson's Law** (for a growing N): If the sequential fraction $s(N)$ diminishes with problem size N , then speedup approaches n as N approaches infinity

$$\text{speedup}(N, n) = n - s(N) (n - 1) \leq n$$



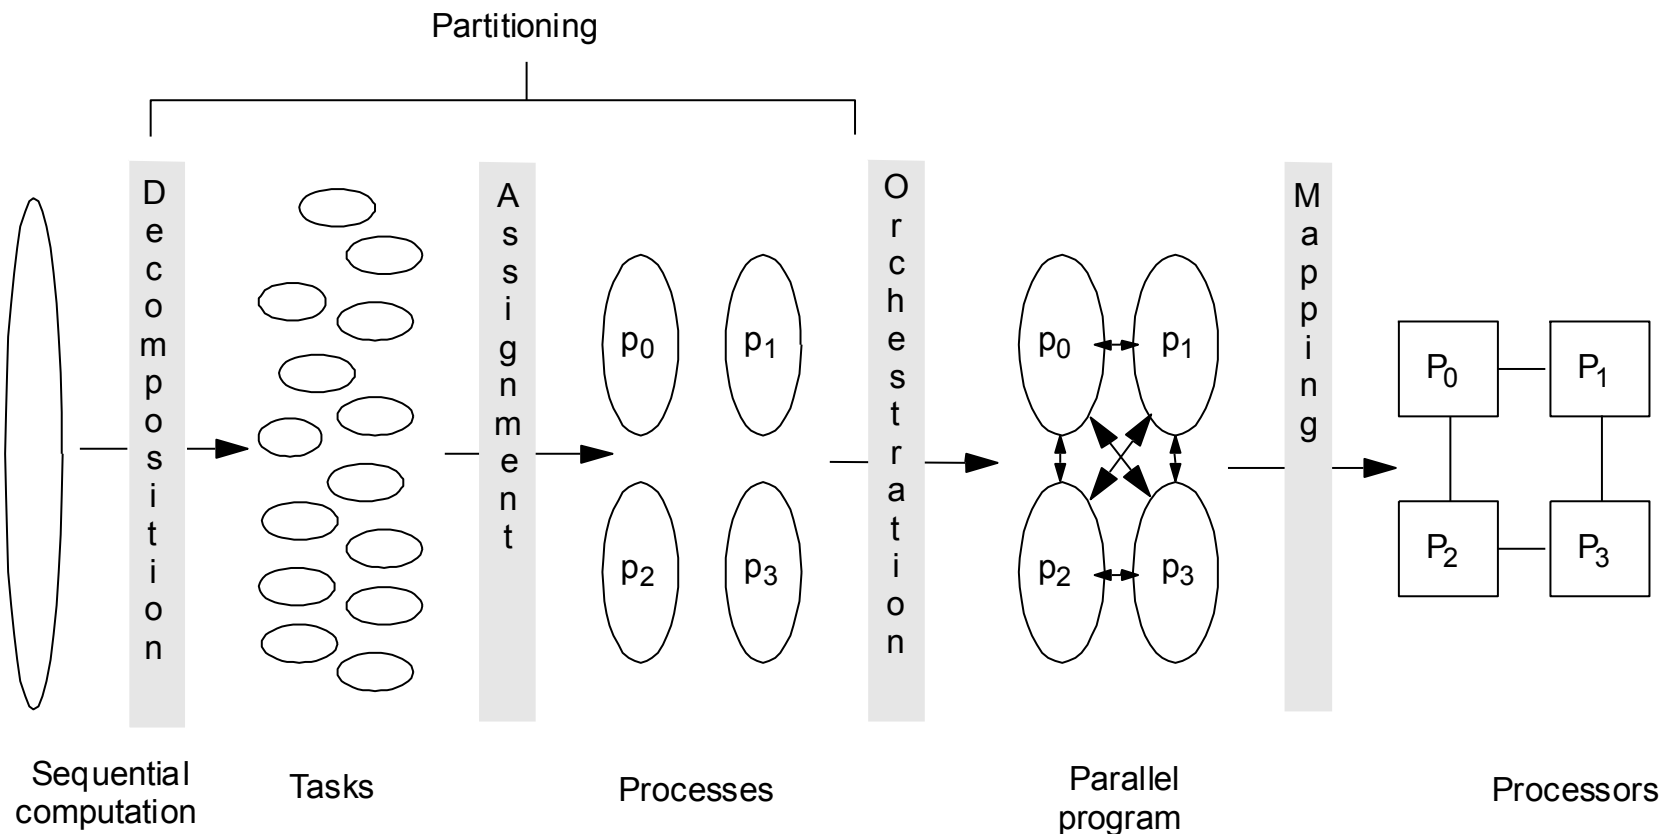
ROYAL INSTITUTE
OF TECHNOLOGY

Parallel Scientific Computing (cont'd)

- Parallelization:
 - Typically – **domain decomposition** (“divide and conquer”)
 - Simple way to identify parallelism is to look at loops (dependence analysis)
 - If not enough concurrency (e.g. loops are sequential), then look further and examine fundamental dependences, ignoring loop structure
- The challenge for a parallel program is to minimize overheads:
 - Create processes once (“work farm”, “thread pool”);
 - Achieve good load balancing;
 - Minimize the need for synchronization and use efficient algorithms for critical sections and barriers.

Recall: Four Steps in Creating a Parallel Program (see Lecture 1)

(1) Decomposition; (2) Assignment; (3) Orchestration; (3) Mapping



Parallel Simulation of a Phenomenon

- Parallelization based on the idea of **domain decomposition**: divide data into partitions (strips, stripes, blocks); assign a worker to each (or use a “bag of tasks”)
- A typical parallel simulation algorithm is a heartbeat algorithm:

– with shared memory:

start with a model;

// step through time

for [t = start to finish] {

Compute();

BARRIER();

Update();

BARRIER();

}

– with message passing:

start with a model;

// step through time

for [t = start to finish] {

Compute();

EXCHANGE();

Update();

EXCHANGE();

}



ROYAL INSTITUTE
OF TECHNOLOGY

Fundamental Techniques in Scientific Computing

- **Computations on a grid of points**
 - To solve partial differential equations (PDE) – to approximate the solution on a finite number of points using iterative numerical methods
 - PDEs are used to model continuous systems and processes (e.g. airflow over a wing)
- **Particle computations**
 - To model (discrete) systems of particles/bodies that exert influence on each other (e.g. stars)
- **Matrix computations**
 - Linear equations
 - Arise in many application domains like optimization problems, e.g., modeling the stock market or the economy, image processing, etc.

1. Computations on Grid of Points

- A continuous-time continuous-space system is modeled as a 3D (2D) grid of points in discrete time
 - The finer spatial and temporal steps are the greater accuracy can be achieved.
 - Many different computations per time step.
 - Applications: weather, fluid (air) flow, plasma physics, etc.
- PDE solvers
- Parallelization – the idea of **domain decomposition** (iterative data parallelism)
 - divide area into blocks or strips of points; assign a worker to each partition; iterations over time steps.

Example: Laplace's Equation

- The Laplace PDE with Dirichlet boundary conditions (2D):
$$\partial^2 \Phi / \partial^2 x + \partial^2 \Phi / \partial^2 y = 0$$
 - A.k.a., stationary heat equation
 - Here Φ is unknown potential such as stress or temperature (heat)
 - Values on boundaries are constant
 - Determine values in the interior area in the steady state (a spacial thermal gradient)
 - In practice: changing boundaries, multiple attributes (e.g., think about weather modeling)
- Solution:
 - Represent the region as an **evenly spaced grid of n by n points** and iterate using **a finite-difference method** until the difference between new values and previous values are within some tolerance ϵ

Iterative Techniques (From Slow to Fast)

- A Jacobi iteration (two arrays):

$$\text{newG}[i, j] = (\text{oldG}[i, j-1] + \text{oldG}[i-1, j] + \text{oldG}[i+1, j] + \text{oldG}[i, j+1]) / 4$$

- A Gauss-Seidel (GS) iteration:

$$G[i, j] = 0,25 \times (G[i, j-1] + G[i-1, j] + G[i+1, j] + G[i, j+1])$$


- A successive over-relaxation (SOR) iteration:

$$G[i, j] = \omega \times 0,25 \times (G[i, j-1] + G[i-1, j] + G[i+1, j] + G[i, j+1]) + (1 - \omega) \times G[i, j]$$

– here $0 < \omega < 2$

- Multigrid

– changing granularity of the grid





ROYAL INSTITUTE
OF TECHNOLOGY

Decomposition of the Dirichlet Problem

- Decomposition in general:
 - Simple way to identify concurrency is to look at loop iterations
 - dependence analysis; if not enough concurrency, then look further
 - If not much concurrency here at this level (most loops are sequential), examine fundamental dependences, ignoring loop structure
- For the Dirichlet problem we use the idea of **domain decomposition (iterative data parallelism)**
 - divide area into blocks or strips of points; assign a worker to each data partition, exchanges values on partitions' boundaries on each iteration (on each time step)

Assignment

- **Static assignments:** n/p rows are assigned to a process
 - Assignment by strips reduces communication by keeping adjacent rows together
 - Reduces concurrency from n to p
- **Dynamic assignment** using the bag-of-tasks technique:
 - get a row index, process the row, get a new row, and so on
- We use static assignment here: assignment of strips
 - Use the number of workers and the problem size to compute a strip size;
 - With shared variables: static assignment controlled by values of variables used as loop bounds
 - Use the strip size and a worker number to determine the bounds.

Jacobi Iteration Using Shared Variables

- For $n = 100$, $p = 4$, each worker is assigned a strip of 25 rows:
worker[1] – rows 1-25, worker[2] – rows 26-50, etc.
- Each worker performs a number of interactions defined by **MAXITERS**.
- To avoid swapping of arrays **grid** and **newgrid**, the loop is unrolled twice – two consecutive iterations in the loop body.
- Barrier after each iteration .

```

real grid[0:n+1,0:n+1], new[0:n+1,0:n+1];
int HEIGHT = n/PR;    # assume PR evenly divides n
real maxdiff[1:PR] = ([PR] 0.0);

procedure barrier(int id) {
    # efficient barrier algorithm from Section 3.4
}

process worker[w = 1 to PR] {
    int firstRow = (w-1)*HEIGHT + 1;
    int lastRow = firstRow + HEIGHT - 1;
    real mydiff = 0.0;
    initialize my strips of grid and new, including boundaries;
    barrier(w);
    for [iters = 1 to MAXITERS by 2] {
        # compute new values for my strip
        for [i = firstRow to lastRow, j = 1 to n]
            new[i,j] = (grid[i-1,j] + grid[i+1,j] +
                        grid[i,j-1] + grid[i,j+1]) * 0.25;

        barrier(w);
        # compute new values again for my strip
        for [i = firstRow to lastRow, j = 1 to n]
            grid[i,j] = (new[i-1,j] + new[i+1,j] +
                        new[i,j-1] + new[i,j+1]) * 0.25;

        barrier(w);
    }
    # compute maximum difference for my strip
    for [i = firstRow to lastRow, j = 1 to n]
        mydiff = max(mydiff, abs(grid[i,j]-new[i,j]));
    maxdiff[w] = mydiff;
    barrier(w);
    # maximum difference is the max of the maxdiff[*]
}

```

Jacobi Iteration Using Message Passing

- Each worker is assigned a partition – a strip of rows.
- Heartbeat algorithm: on each iteration, each worker exchanges edges with its neighbors (if any):


```
if (w > 1)
    send up[w-1](new[1,*]);
if (w < PR)
    send down[w+1](new[HEIGHT,*]);
if (w > PR)
    recv up[w](new[HEIGHT+1,*]);
if (w < 1)
    recv down[w](new[0,*]);
```

```
chan up[1:PR](real edge[0:n+1]);
chan down[1:PR](real edge[0:n+1]);
chan diff(real);

process worker[w = 1 to PR] {
    int HEIGHT = n/PR;    # assume PR evenly divides n
    real grid[0:HEIGHT+1,0:n+1], new[0:HEIGHT+1,0:n+1];
    real mydiff = 0.0, otherdiff = 0.0;
    initialize grid and new, including boundaries;
    for [iters = 1 to MAXITERS by 2] {
        # compute new values for my strip
        for [i = 1 to HEIGHT, j = 1 to n]
            new[i,j] = (grid[i-1,j] + grid[i+1,j] +
                        grid[i,j-1] + grid[i,j+1]) * 0.25;
        exchange edges of new -- see text;
        # compute new values again for my strip
        for [i = 1 to HEIGHT, j = 1 to n]
            grid[i,j] = (new[i-1,j] + new[i+1,j] +
                        new[i,j-1] + new[i,j+1]) * 0.25;
        exchange edges of grid -- see text;
    }
    # compute maximum difference for my strip
    for [i = 1 to HEIGHT, j = 1 to n]
        mydiff = max(mydiff, abs(grid[i,j]-new[i,j]));
    if (w > 1)
        send diff(mydiff);
    else    # worker 1 collects differences
        for [i = 1 to w-1] {
            receive diff(otherdiff);
            mydiff = max(mydiff, otherdiff);
        }
    # maximum difference is value of mydiff in worker 1
}
```

Red/Black Successive Over-Relaxation (SOR) Iteration


- To decompose, exploit application knowledge: reorder grid traversal: red-black ordering (R depends on B and visa versa)
 - Different ordering of updates: may converge quicker or slower
 - **Red sweep and black sweep are each fully parallel:**
 - Global synch between them (conservative but convenient)



See: MPT book Figure 11.6 Red/black Gauss-Seidel using shared variables.

Multigrid Methods

- To solve a large problem rapidly
- Idea: To use grids of different granularities and to switch between them to increase the rate of convergence of the finest grid



• *Fine grid point*

○ *Coarse and fine grid point*

Coarse Grid Correction

- Start with a fine grid, update points for a few iterations using any relaxation method (J, GS, SOR)
- Restrict the result to coarse-grained (twice coarser) by a **restriction operator**


0	1/8	0
1/8	1/2	1/8
0	1/8	0

$$\text{coarse}[x,y] = \text{fine}[i,j]*0.5 + (\text{fine}[i-1,j]+\text{fine}[i,j-1]+\text{fine}[i+1,j]+\text{fine}[i,j+1])*0.125$$

- Compute the coarse grid.
- Interpolate the coarse grid back to the fine grid by using an **interpolation operator**
- Update points on the fine grid for a few iterations.
- Change again to the coarse grid, etc.

1/4	1/2	1/4
1/2	1	1/2
1/4	1/2	1/4

Patterns for Multigrid



V cycle

W cycle



Full multigrid

2. Particle Computations

- To model a discrete system consisting of interacting particles/bodies
 - n particles requires $O(n^2)$ calculations on a time step
 - By using approximations, can be improved to $O(n \log_2 n)$
- Typically irregular iterative computations over time steps.
- Examples of applications:
 - Particle interactions due to chemical bonding,
 - Gravity (evolution of galaxies)
 - Electrical charge, fluid flow, etc.

Example: Gravitational N-Body Problem

- N bodies. Each body:
 - Position $\mathbf{p} (x, y, z)$
 - Velocity $\mathbf{v} (x, y, z)$
 - Mass m is constant
 - Force $\mathbf{F} (x, y, z)$
- Gravity causes the bodies to accelerate and to move.
- The motion of the bodies is simulated by stepping through discrete instants of time:

initialize bodies;

for time step {

calculate forces; // read \mathbf{p} , m ; compute \mathbf{F}

move bodies; // read \mathbf{F} , m ; compute \mathbf{v} and new \mathbf{p}

}

Calculation Forces. Moving Bodies

- **Force** on a body is the vector sum of the forces from all other bodies
 - Magnitude of the gravitational force between bodies i and j :
$$F[i,j] = (G * m[i] * m[j]) / r[i,j]**2$$
 - magnitude: symmetric ("equal and opposite")
 - direction: vector from one body to the other
- Changes in velocity and position (**moving a body**) – leapfrog scheme:
 - Acceleration: $a = F / m$
 - Change in velocity: $dv = a * DT$
 - Change in position:
$$dp = v * DT + (a/2)*DT**2 \approx (v + dv/2) * DT$$
 - Here DT – the length of a time step



ROYAL INSTITUTE
OF TECHNOLOGY

Parallelization of the N-Body Problem

- Assume, shared memory programming model
- **Parallelization – divide bodies among workers**, use a barrier after each phase:

```
initialize bodies;  
for time step {  
    calculate forces;  
    BARRIER;  
    move bodies;  
    BARRIER;  
}
```

Assignment

- Assume P worker processes.
- Three ways to assign work (distribute bodies among workers):
 - **Blocks** (assign each worker a continuous block of bodies):
 - $w[1]$ is assigned the first n/P bodies,
 - $w[2]$ is assigned the next n/P , etc.
 - **Stripes** (cyclic allocation of bodies to workers):
 - body 1 to $w[1]$, body 2 to $w[2]$, ..., body P to $w[P]$,
body $P+1$ to $w[1]$, body $P+2$ to $w[2]$, etc.
 - **Reverse stripes** (reverse cyclic allocation):
 - body 1 to $w[1]$, body 2 to $w[2]$, ..., body P to $w[P]$,
then reverse stripes:
body $P+1$ to $w[P]$, body $P+2$ to $w[P-1]$, etc.

Assignment of Bodies to Workers

- Assume, 2 workers (B – “Black” and W – “White”) and 8 bodies
- For each body i assigned to a worker, the worker computes forces between the body i and bodies $i + 1, \dots, n$

	1	2	3	4	5	6	7	8	
<i>Pattern</i>									<i>Workload</i>
blocks	B	B	B	B	W	W	W	W	B = 22, W = 6
stripes	B	W	B	W	B	W	B	W	B = 16, W = 12
reverse stripes	B	W	W	B	B	W	W	B	B = 14, W = 14

- The pattern of stripes leads to a fairly well-balanced workload that is easy to program.

Synchronization

- Barriers after each stage
- Critical sections in **calculate forces** phase:
 - Access to the same body by different workers (writers) must be synchronized (mutual exclusion)
 - Approaches:
 - One global lock – very inefficient
 - A lock per a body – too much synchronization
 - **Eliminate critical sections at all**
- To eliminate critical sections, replicate and split shared data:
 - Change the force vector into a force matrix: **a private force vector per worker**.
 - The result force vector: sums in columns – can be calculated in parallel in the moving phase



ROYAL INSTITUTE
OF TECHNOLOGY

N-Body Parallel Program with Shared Variables

- Procedures:
 - Calculate forces
 - Move bodies
 - Barrier
- Workers:

```
process Worker[w = 1 to PR] {  
  # run the simulation with time steps of DT  
  for [time = start to finish by DT] {  
    calculateForces(w);  
    barrier(w);  
    moveBodies(w);  
    barrier(w);  
  }  
}
```

```
type point = rec(double x, y); double G = 6.67e-11;  
point p[1:n], v[1:n], f[1:PR,1:n]; # position, velocity,  
double m[1:n]; # force and mass for each body  
initialize the positions, velocities, forces, and masses;  
  
procedure barrier(int w) { # efficient barrier from Section 3.4 }  
  
# calculate forces for bodies assigned to worker w  
procedure calculateForces(int w) {  
  double distance, magnitude; point direction;  
  for [i = w to n by PR, j = i+1 to n] {  
    distance = sqrt( (p[i].x - p[j].x)**2 +  
                    (p[i].y - p[j].y)**2 );  
    magnitude = (G*m[i]*m[j]) / distance**2;  
    direction = point(p[j].x-p[i].x, p[j].y-p[i].y);  
    f[w,i].x = f[w,i].x + magnitude*direction.x/distance;  
    f[w,j].x = f[w,j].x - magnitude*direction.x/distance;  
    f[w,i].y = f[w,i].y + magnitude*direction.y/distance;  
    f[w,j].y = f[w,j].y - magnitude*direction.y/distance;  
  }  
}  
  
# move the bodies assigned to worker w  
procedure moveBodies(int w) {  
  point deltav; # dv = f/m * DT  
  point deltap; # dp = (v + dv/2) * DT  
  point force = (0.0, 0.0);  
  for [i = w to n by PR] {  
    # sum the forces on body i and reset f[*,i]  
    for [k = 1 to PR] {  
      force.x += f[k,i].x; f[k,i].x = 0.0;  
      force.y += f[k,i].y; f[k,i].y = 0.0;  
    }  
    deltav = point(force.x/m[i] * DT, force.y/m[i] * DT);  
    deltap = point( (v[i].x + deltav.x/2) * DT,  
                  (v[i].y + deltav.y/2) * DT);  
    v[i].x = v[i].x + deltav.x;  
    v[i].y = v[i].y + deltav.y;  
    p[i].x = p[i].x + deltap.x;  
    p[i].y = p[i].y + deltap.y;  
    force.x = force.y = 0.0;  
  }  
}
```


Approximate Methods for the N-Body Problem

- Approximation (based on Newtonian mechanics):
 - If two bodies are far apart, the force between them is negligible
 - The force between a distant body and a group can be approximated by the force between the former and a single body that approximates the group (has the total mass of all the bodies in the group and is located in the center of mass of the group).
- Two methods that exploit above approximations:
 - Barnes-Hut algorithm
 - Fast Multipole Method
 - Both allows $O(n \log n)$ force calculation per time step
 - See section 11.2.4 in the MPD text.

Summary of Parallelization of the N-Body Problem

- Decomposition: distribute bodies among workers.
- $O(n^2)$ force calculations per time step
- Minimizing overheads:
 - create workers once
 - avoid critical sections: compute forces into "private" arrays; add them when ready to move bodies:
 - calculate phase: read positions and masses; compute forces
 - move phase: read forces and masses; compute new v and p
- Load balancing: assignment of stripes to workers
- Faster algorithms, $O(n \log n)$ use approximations



ROYAL INSTITUTE
OF TECHNOLOGY

3. Matrix Computations

- Linear equations
- Arise in many application domains
 - like optimization problems
 - e.g., modeling the stock market or the economy
 - image processing, etc.